

# ***Deep Learning Crash Course***



**Hui Xue**

**Fall 2022**


[www.deeplearningcrashcourse.org](http://www.deeplearningcrashcourse.org)

# Outline

- Loss function for classification
- Gradient descent
- Bias and variance


# *How do we know a model is good?*

Given a model and its parameters, how do we know this model gives correct prediction?



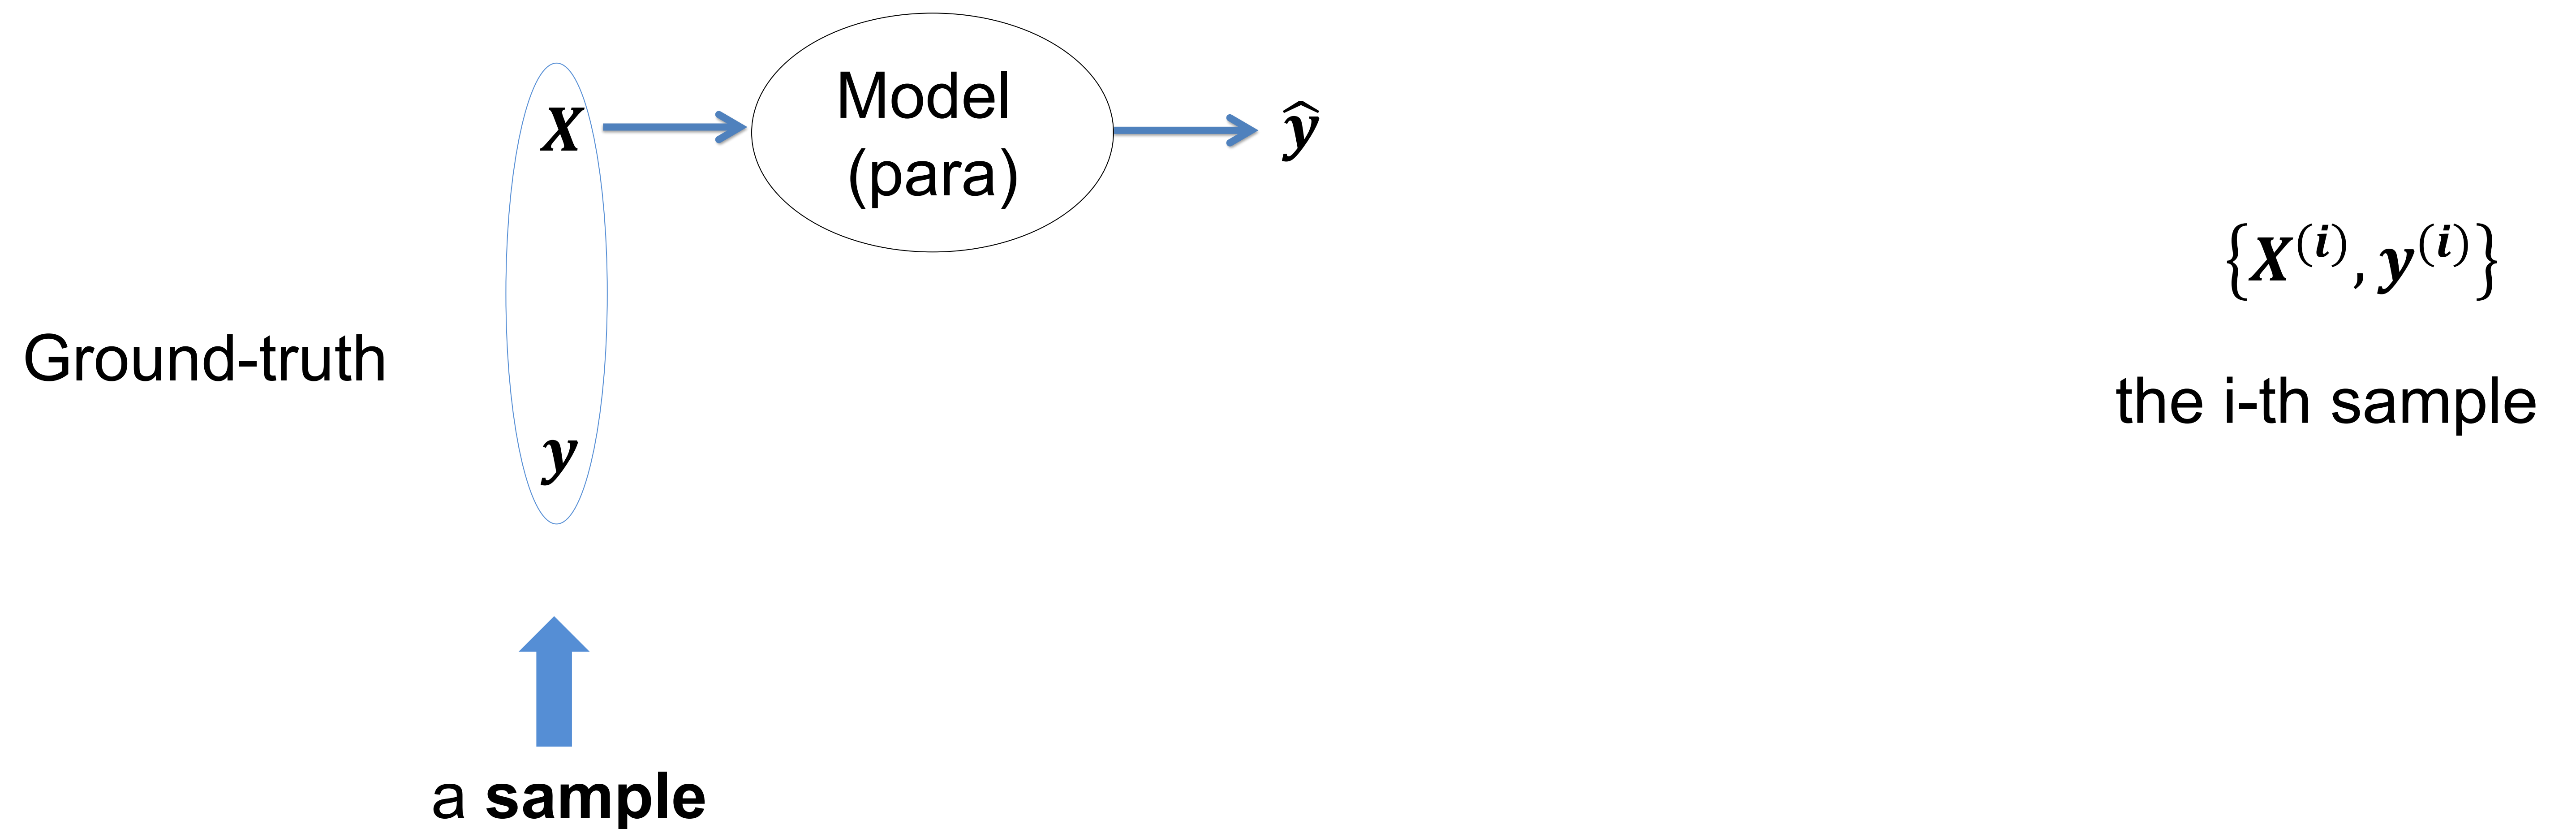
# *How do we know a model is good?*

Given a model and its parameters, how do we know this model gives correct prediction?



# *How do we know a model is good?*


Given a model and its parameters, how do we know this model gives correct prediction?





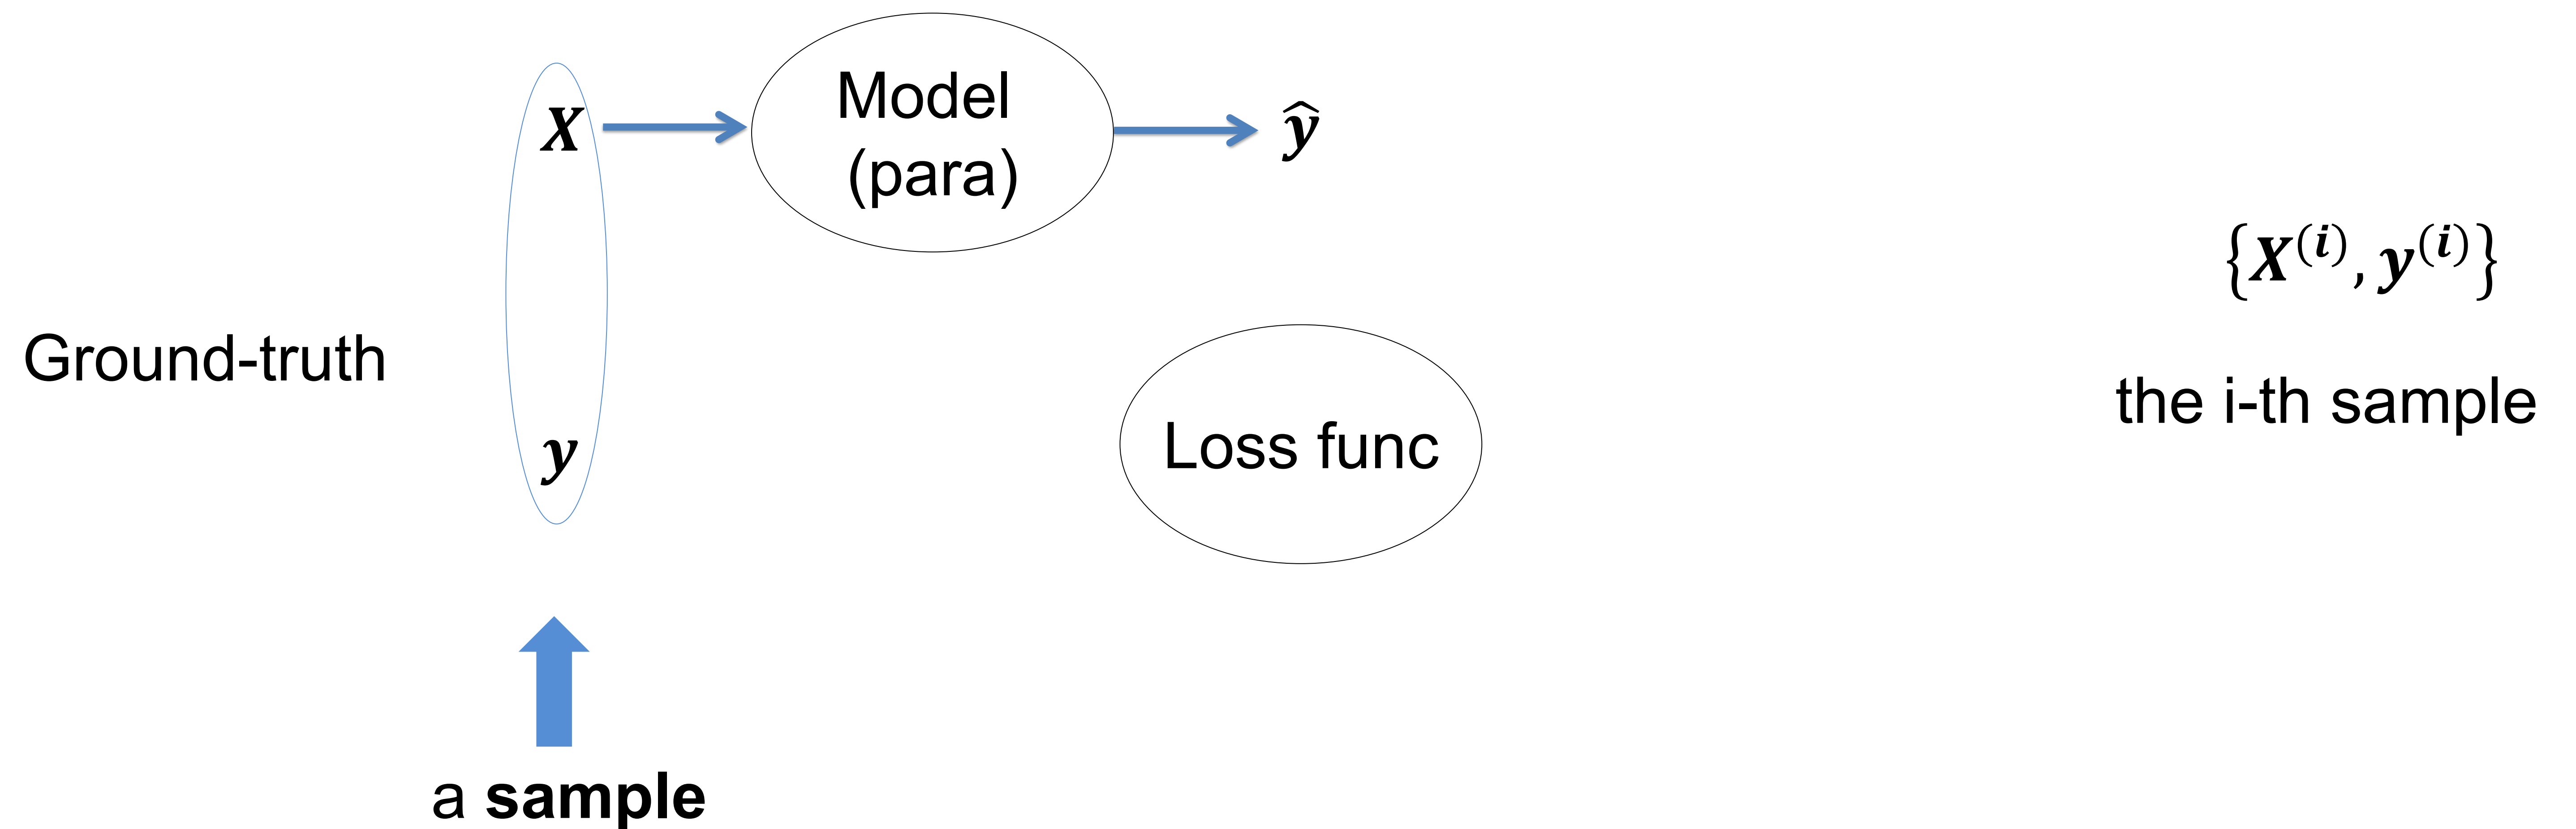
# How do we know a model is good?

Given a model and its parameters, how do we know this model gives correct prediction?



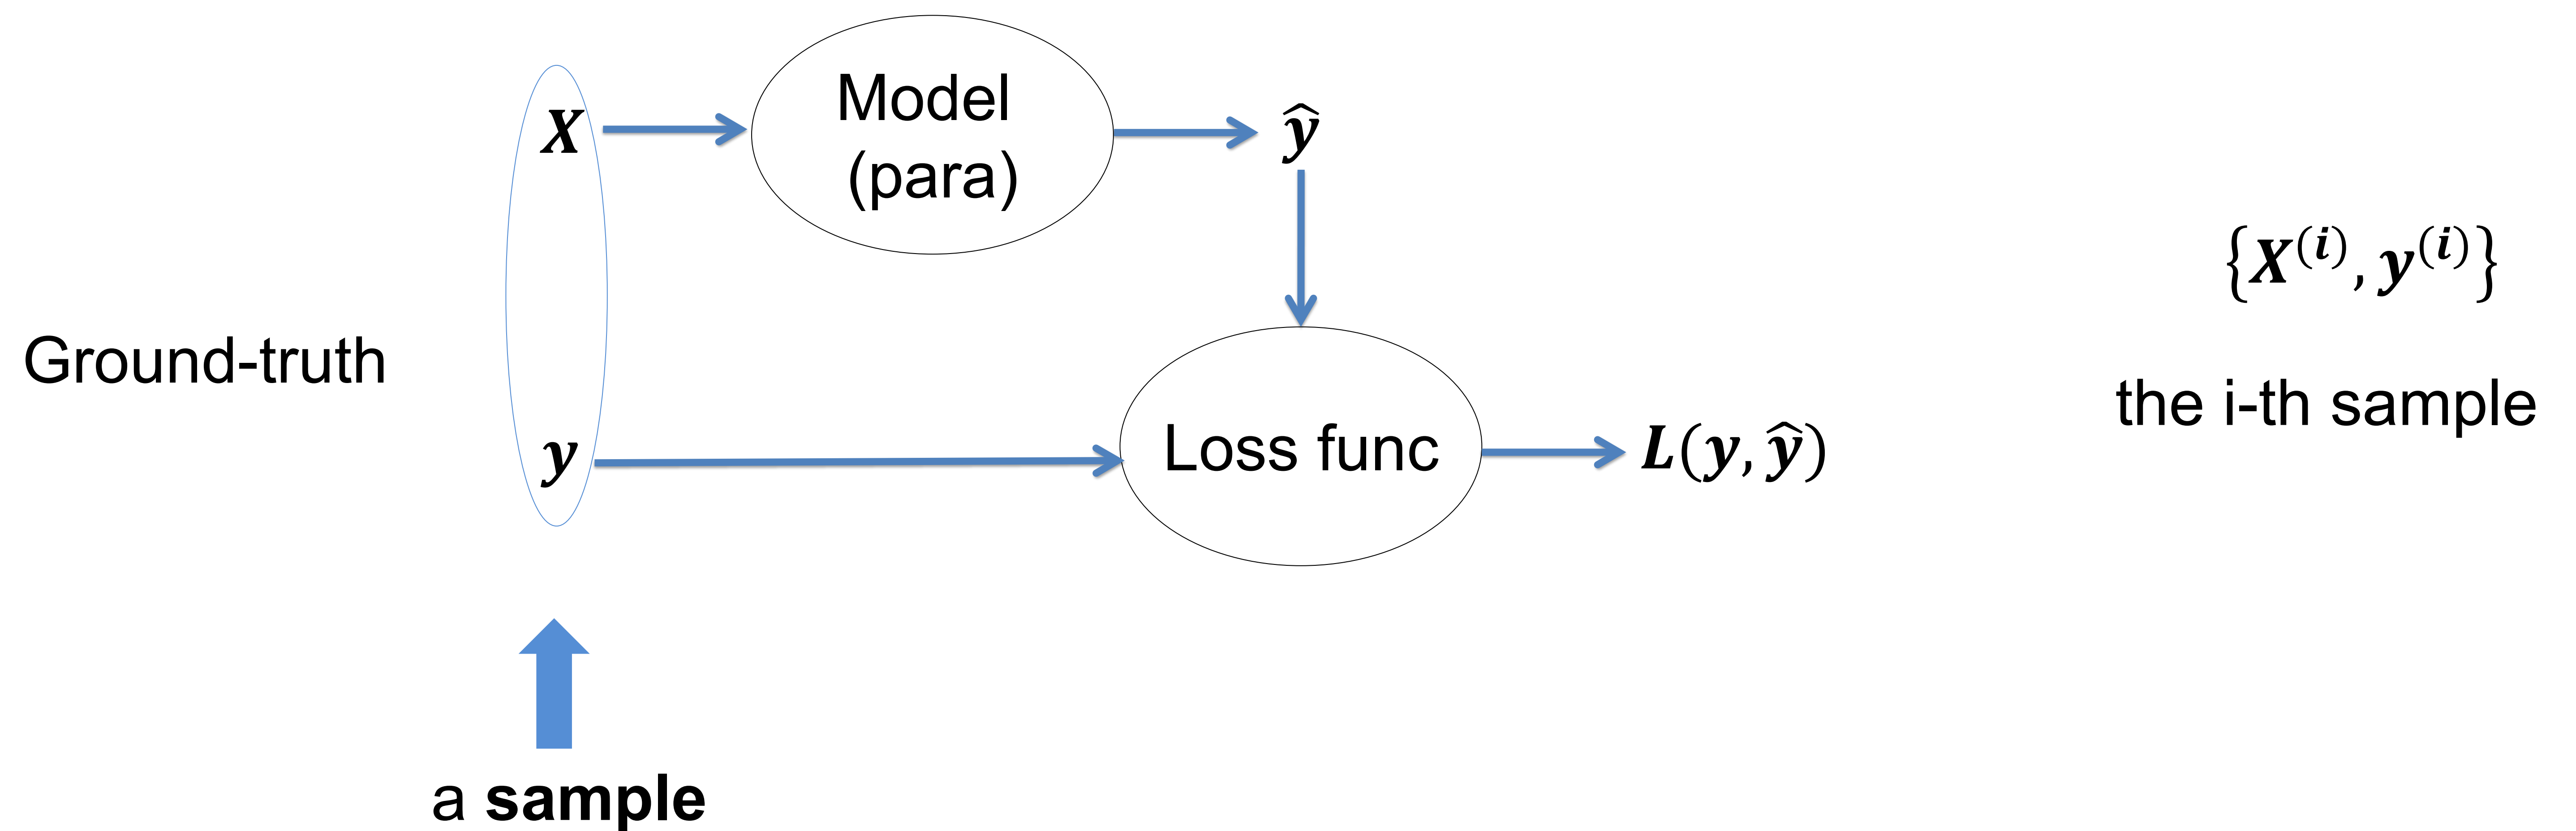
# How do we know a model is good?

Given a model and its parameters, how do we know this model gives correct prediction?



# How do we know a model is good?

Given a model and its parameters, how do we know this model gives correct prediction?





# Cross-entropy

Given two probability distribution  $p$  and  $q$ ,  $p$  is the target distribution and  $q$  is the model predicted distribution:

$$H(p, q) = - \sum_{\text{for all possible } x} p(x) \log[q(x; \theta)]$$

- Measure how close two distributions are (not 100% rigorous)
- Minimized if  $p=q$

$$H(p, q) = H(p) + D_{KL}(p \parallel q)$$

$D_{KL}(p \parallel q)$  is the KL divergence of  $p$  and  $q$ .

$$D_{KL}(p \parallel q) \geq 0$$

$D_{KL}(p \parallel q) = 0$  if and only if  $p=q$

$H(p)$  is the entropy of distribution  $p$

$$H(p, q) \neq H(q, p)$$

$$\begin{aligned} H(p, q) &= - \sum_x p(x) \log[q(x)] \\ &= - \sum_x p(x) \log \left[ \frac{p(x)}{p(x)} q(x) \right] \\ &= - \sum_x p(x) \left\{ \log[p(x)] + \log \left[ \frac{q(x)}{p(x)} \right] \right\} \\ &= - \sum_x p(x) \log[p(x)] - \sum_x p(x) \log \left[ \frac{q(x)}{p(x)} \right] \\ &= \underbrace{- \sum_x p(x) \log[p(x)]}_{H(p)} + \underbrace{\sum_x p(x) \log \left[ \frac{p(x)}{q(x)} \right]}_{D_{KL}(p \parallel q)} \end{aligned}$$

Entropy of distribution  $p$

# Binary Cross-entropy loss

For the binary classification,

$y = 1$  for the “Yes” class and 0 for the “No” class

$\hat{y} \in [0, 1]$ , probability belonging to the “Yes” class  
 $1 - \hat{y}$  for the “No” class

Then, the cross entropy (in this case, the BCE loss) is:

$$L_{BCE\_loss}(y, \hat{y}) = -[y \cdot \log(\hat{y}) + (1 - y) \cdot \log(1 - \hat{y})]$$

To minimize the BCE loss,

if  $y = 0$ ,  $\hat{y} = 0$  will reach minimum of  $\log(1)=0$

if  $y = 1$ ,  $\hat{y} = 1$  will reach minimum of  $\log(1)=0$

For all other cases,  $L_{BCE\_loss} > 0$

# Cross-entropy loss for multi-class

For the multi-class classification,

$$y = \begin{cases} CH3, & 1 & 0 & 0 & 0 & 0 \\ CH2, & 0 & 1 & 0 & 0 & 0 \\ CH4, & 0 & 0 & 1 & 0 & 0 \\ SAX, & 0 & 0 & 0 & 1 & 0 \\ Other, & 0 & 0 & 0 & 0 & 1 \end{cases}$$

$y$  is one-hot encoding for the correct class

$$y_k = \begin{cases} CH3, 0 \\ CH2, 1 \\ CH4, 2 \\ SAX, 3 \\ Other, 4 \end{cases}$$

$y_k$  is the index of correct class for sample  $y$

$\hat{y}$  is a  $[K, 1]$  vector for the probability of  $X$  belonging to each class

$\hat{y}[0]$  is a scalar to predict  $X$  being CH3

$\hat{y}[1]$  is a scalar to predict  $X$  being CH2

$\hat{y}[2]$  is a scalar to predict  $X$  being CH4

... ..

$\hat{y}[y_k]$  is a scalar to predict  $X$  being class  $y_k$

# Cross-entropy loss

For the multi-class classification,

$$L_{CE\_loss}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{k=0}^4 \mathbf{y}[k] \log(\hat{\mathbf{y}}[k])$$

$\mathbf{y}$  is one-hot encoding for the correct class

$\hat{\mathbf{y}}$  is a  $[K, 1]$  vector for the probability of  $X$  being each class

$\hat{\mathbf{y}}[y_k]$  is a scalar to predict  $X$  being class  $y_k$

$y_k$  is the index of correct class for sample  $\mathbf{y}$

$$L_{CE\_loss}(\mathbf{y}, \hat{\mathbf{y}}) = -\log(\hat{\mathbf{y}}[y_k])$$

Only one term left!

To minimize the cross-entropy loss,

We need  $\hat{\mathbf{y}}[y_k] = 1$ ; for all other cases,  $L_{CE\_loss} > 0$



# Cross-entropy loss

|       | $y$ | $\hat{y}$ |
|-------|-----|-----------|
| CH2   | 0   | 0.2       |
| CH3   | 0   | 0.3       |
| CH4   | 1   | 0.1       |
| SAX   | 0   | 0.15      |
| Other | 0   | 0.25      |

$$\begin{aligned} L_{CE\_loss}(\mathbf{y}, \hat{\mathbf{y}}) &= - \sum_{k=0}^4 y[k] \log(\hat{y}[k]) \\ &= -[0 \times \log(0.2) \\ &\quad + 0 \times \log(0.3) \\ &\quad + 1 \times \log(0.1) \\ &\quad + 0 \times \log(0.15) \\ &\quad + 0 \times \log(0.25)] \\ &= -\log(0.1) = 2.3 \end{aligned}$$




# Cross-entropy loss

|       | $y$ | $\hat{y}$ |
|-------|-----|-----------|
| CH2   | 0   | 0.02      |
| CH3   | 0   | 0.1       |
| CH4   | 1   | 0.8       |
| SAX   | 0   | 0.05      |
| Other | 0   | 0.03      |

$$\begin{aligned} L_{CE\_loss}(\mathbf{y}, \hat{\mathbf{y}}) &= - \sum_{k=0}^4 y[k] \log(\hat{y}[k]) \\ &= -[0 \times \log(0.02) \\ &\quad + 0 \times \log(0.1) \\ &\quad + 1 \times \log(0.8) \\ &\quad + 0 \times \log(0.05) \\ &\quad + 0 \times \log(0.03)] \\ &= -\log(0.8) = 0.22 \end{aligned}$$

# Where we are



$$\mathbf{Z}^{[1]} = \mathbf{W}^{[1]} \mathbf{X} + \mathbf{b}^{[1]}$$

$$\mathbf{a}^{[1]} = f(\mathbf{Z}^{[1]})$$

$$\mathbf{Z}^{[2]} = \mathbf{W}^{[2]} \mathbf{a}^{[1]} + \mathbf{b}^{[2]}$$

$$\mathbf{a}^{[2]} = f(\mathbf{Z}^{[2]})$$

$$\mathbf{Z}^{[3]} = \mathbf{W}^{[3]} \mathbf{a}^{[2]} + \mathbf{b}^{[3]}$$

$$\hat{\mathbf{y}} = \mathbf{a}^{[3]} = f(\mathbf{Z}^{[3]})$$


$$L_{BCE\_loss}(\mathbf{y}, \hat{\mathbf{y}})$$

$$L_{CE\_loss}(\mathbf{y}, \hat{\mathbf{y}})$$

Model : map input to output

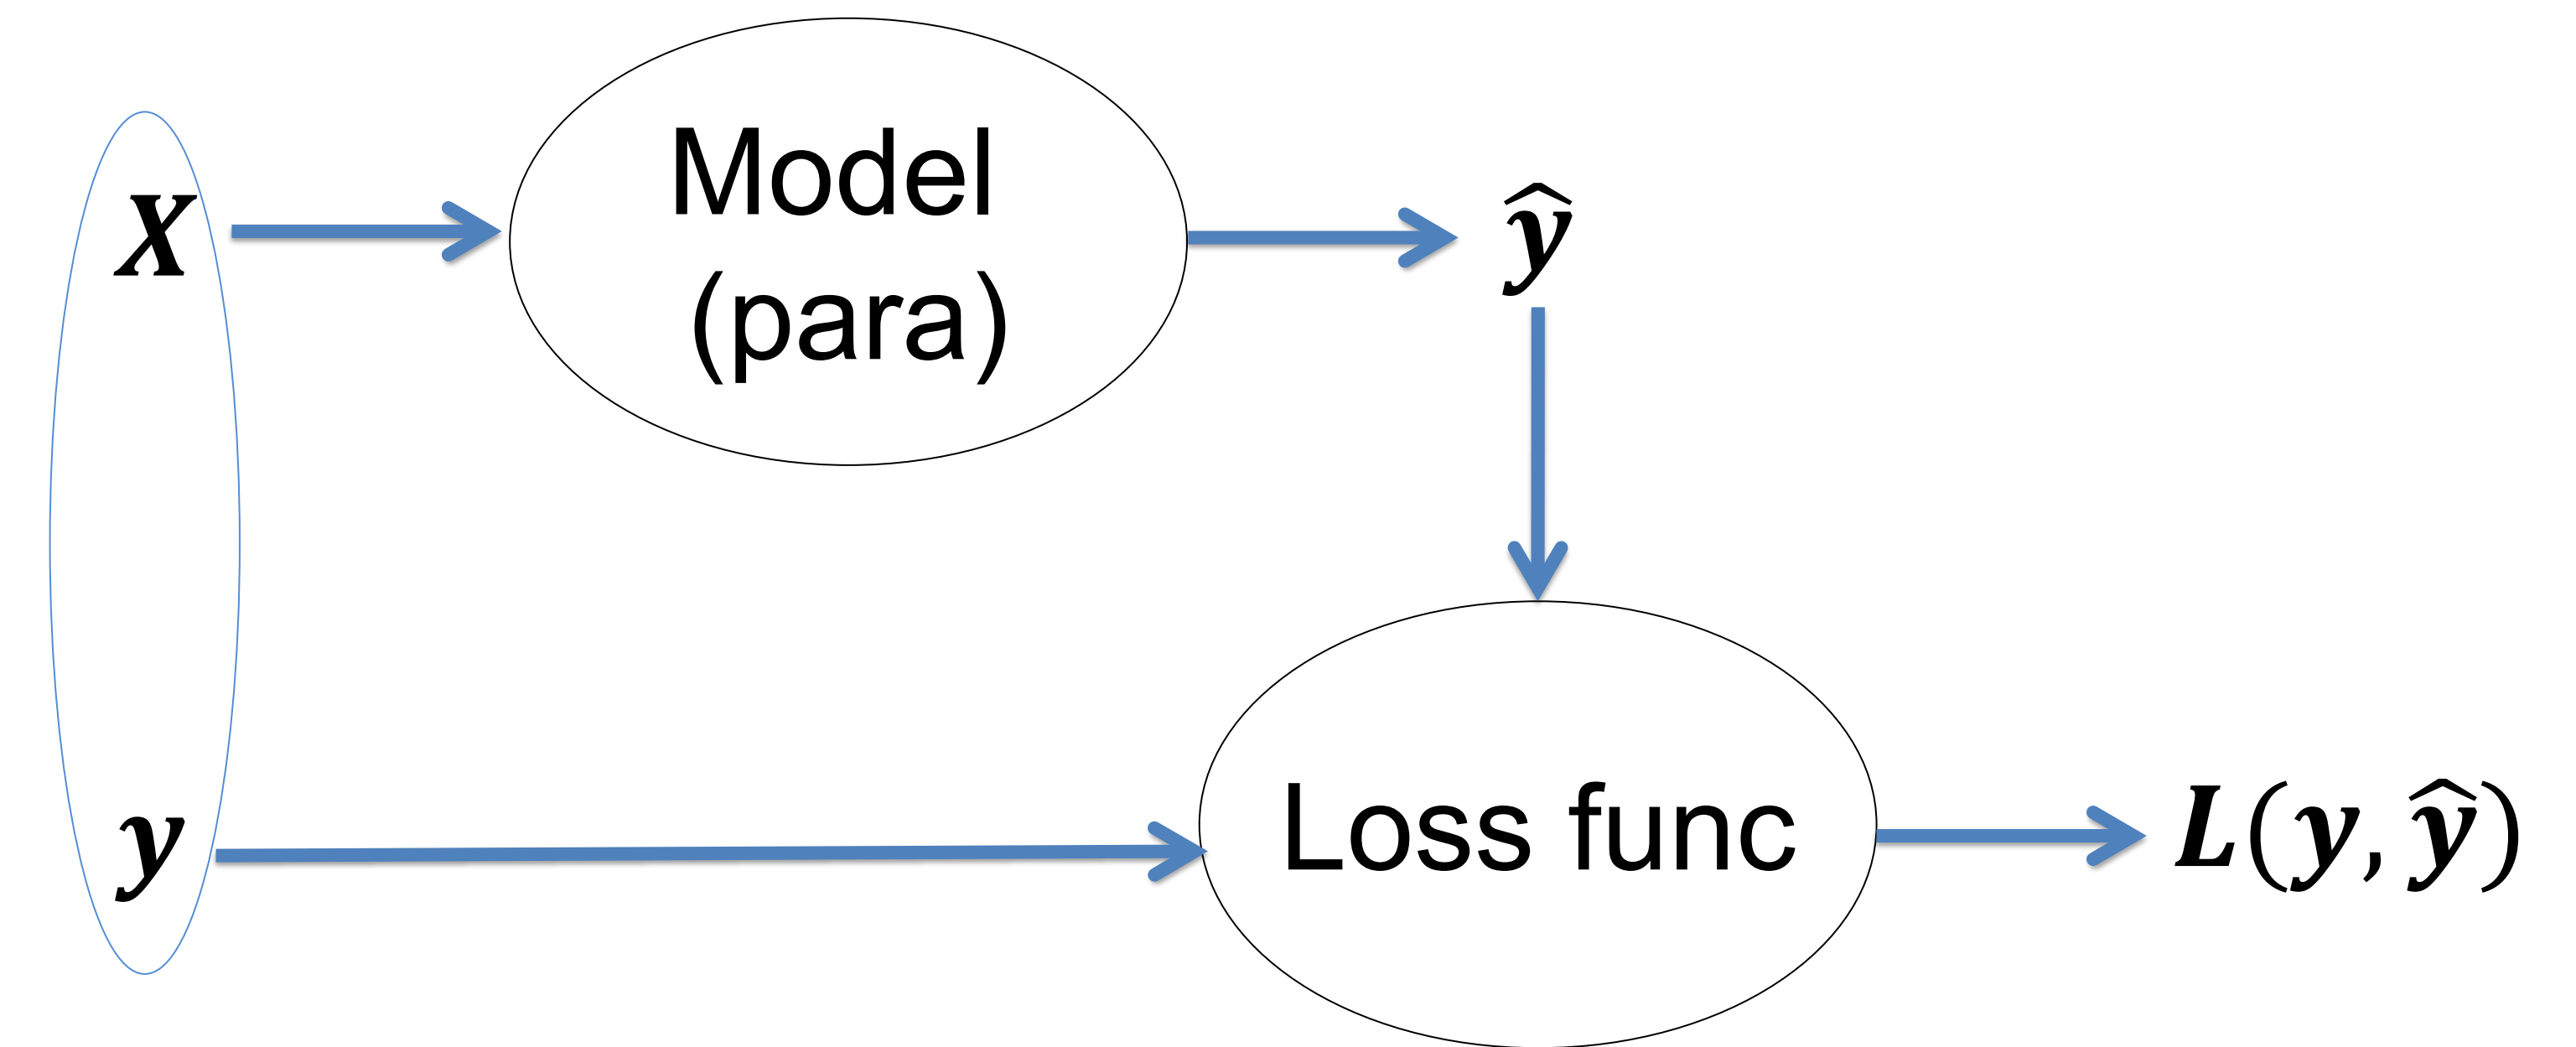
Loss : measure how good the output is, compared to the label

# How do we determine the model parameters?




$W^{[1]}, b^{[1]}$   
 $W^{[2]}, b^{[2]}$   
 $W^{[3]}, b^{[3]}$

Model : map input to output

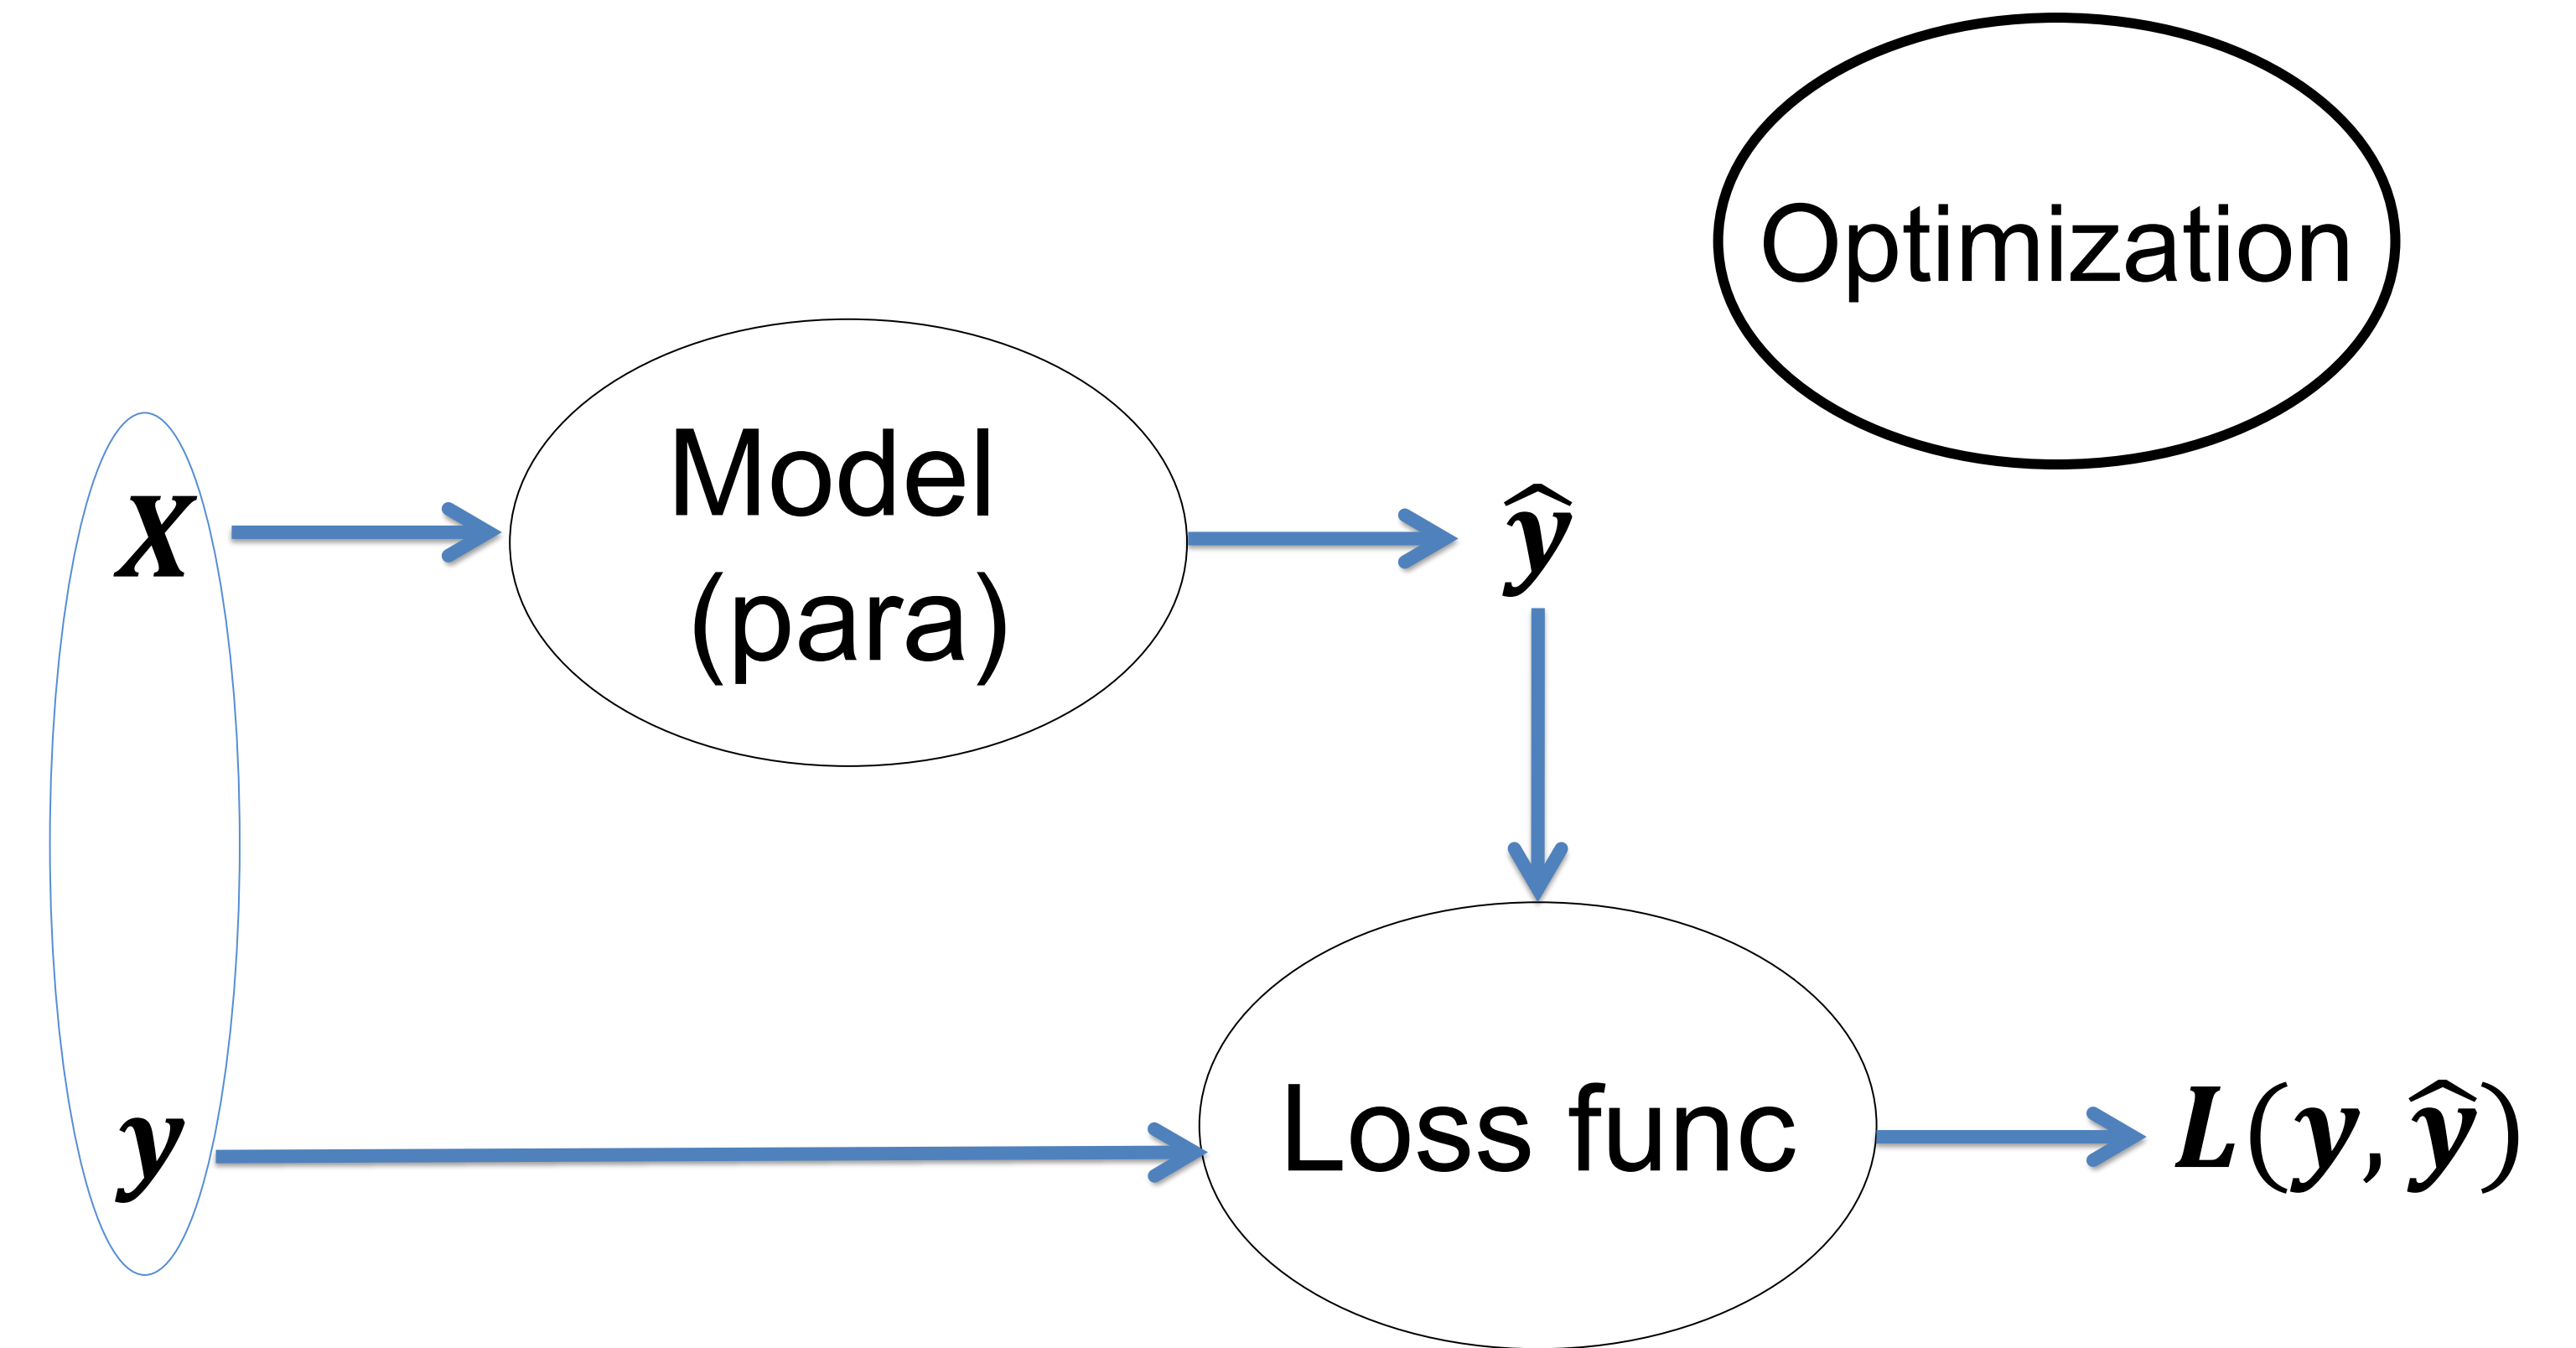


# How do we determine the model parameters?




$W^{[1]}, b^{[1]}$   
 $W^{[2]}, b^{[2]}$   
 $W^{[3]}, b^{[3]}$

Model : map input to output




# How do we determine the model parameters?




$W^{[1]}, b^{[1]}$   
 $W^{[2]}, b^{[2]}$   
 $W^{[3]}, b^{[3]}$

Model : map input to output



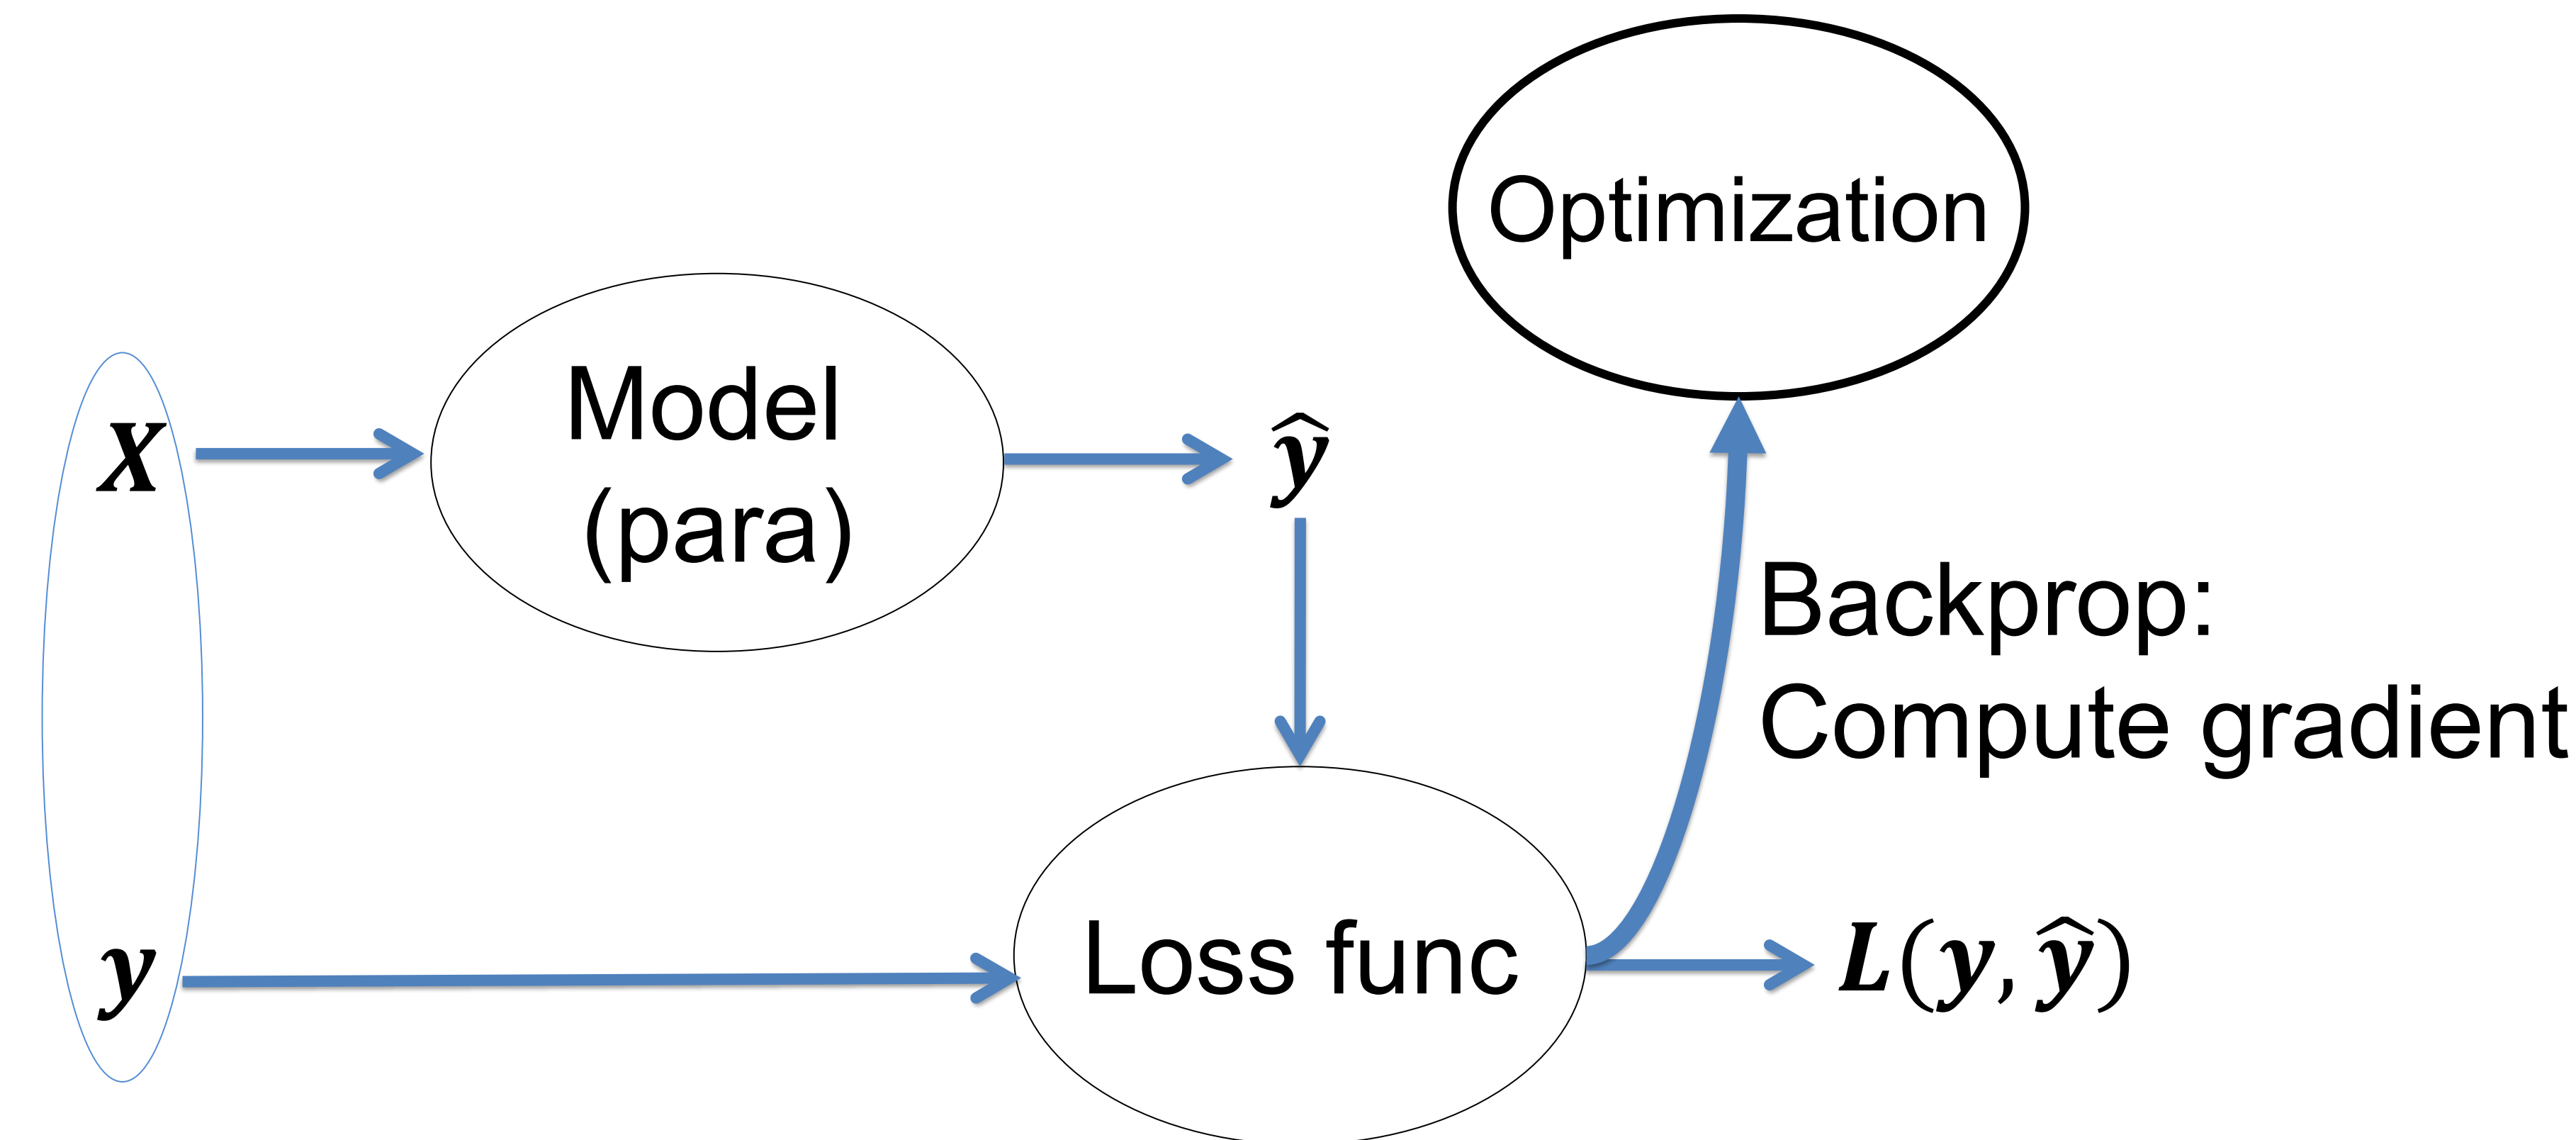


# How do we determine the model parameters?




$W^{[1]}, b^{[1]}$   
 $W^{[2]}, b^{[2]}$   
 $W^{[3]}, b^{[3]}$

Model : map input to output

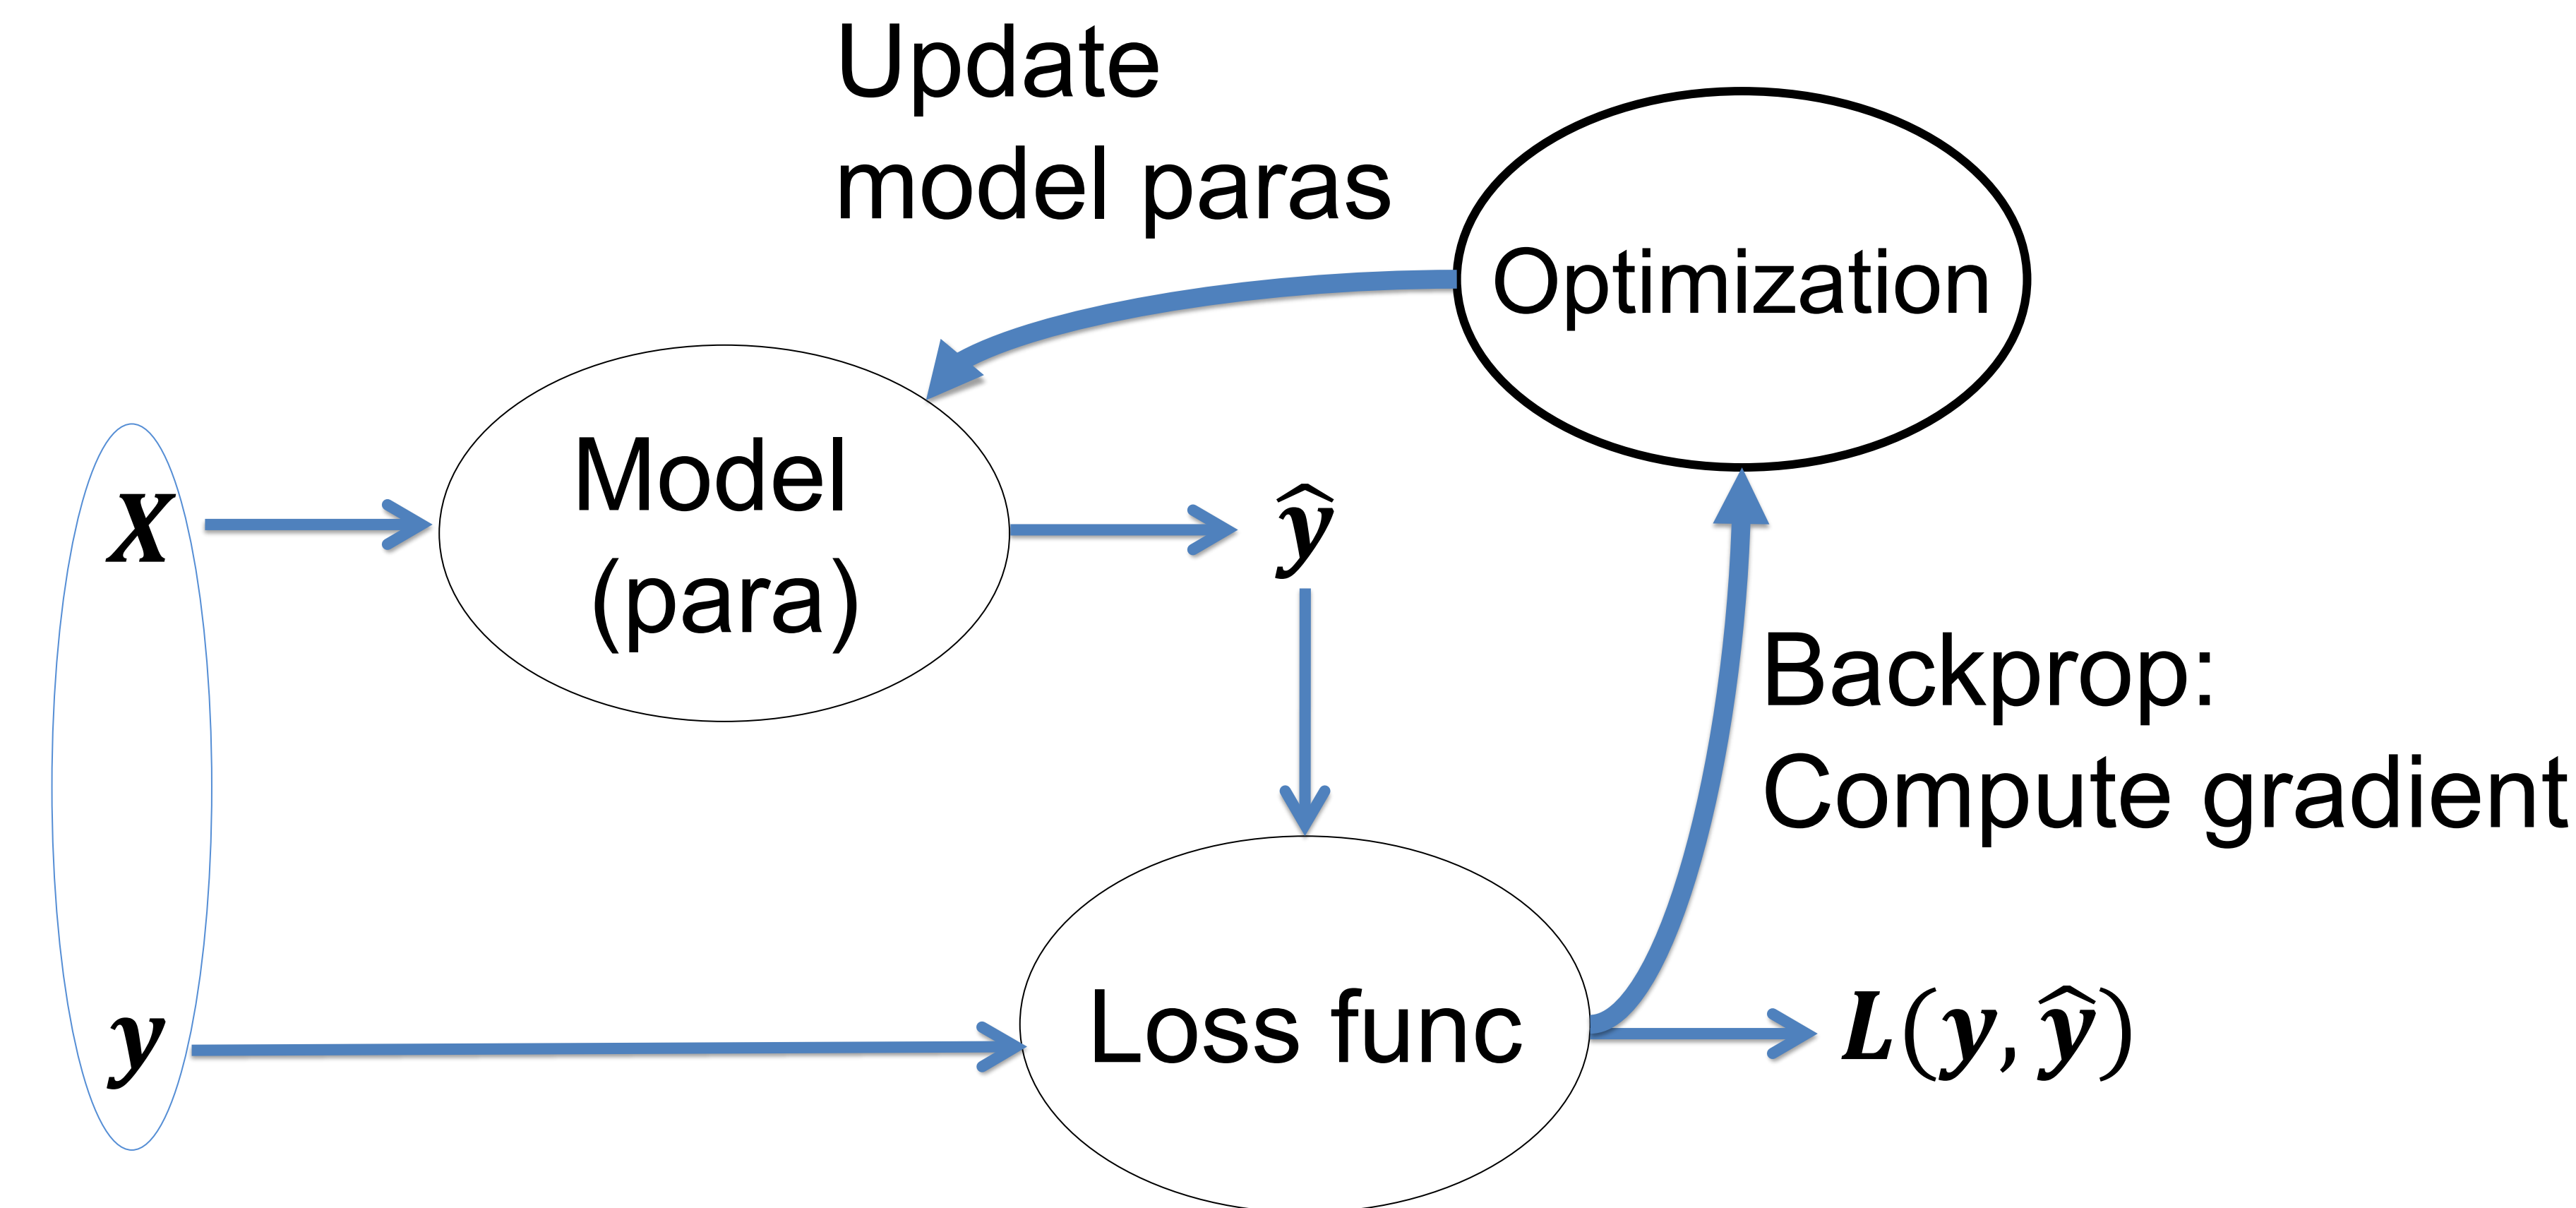


# How do we determine the model parameters?




$W^{[1]}, b^{[1]}$   
 $W^{[2]}, b^{[2]}$   
 $W^{[3]}, b^{[3]}$

Model : map input to output



# Gradient descent



To find minimal point of a function

- Start from an initial point
- Follow the negative gradient direction



# Gradient of a function

Gradient of a function indicates the steepest direction to **increase** this function in a neighborhood.  
**Negative** gradient of a function indicates the steepest direction to **decrease** this function in a neighborhood.


$f$  is a scalar function.

In 1-dimension, the derivative of a function:

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

$$\nabla f(\mathbf{p}) = \begin{bmatrix} \frac{\partial f}{\partial x_0} \\ \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_{N-1}} \end{bmatrix}$$

- $\mathbf{p} = [x_0, x_1, x_2, \dots, x_{N-1}]^T$  is a point in N-dimensional space
- $\nabla f(\mathbf{p})$  is the gradient at point  $\mathbf{p}$ .
- Its direction points to the steepest slope to **increase** the function.
- The steepness of the slope at that point is given by the  $|\nabla f(\mathbf{p})|$



[https://en.wikipedia.org/wiki/Gradient#/media/File:Gradient\\_of\\_a\\_Function.tif](https://en.wikipedia.org/wiki/Gradient#/media/File:Gradient_of_a_Function.tif)

# Gradient descent

Gradient direction  
is  
the steepest direction to **increase** this function

We want to minimize the loss  $L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{W}, \mathbf{b}))$  by  
adjusting the model parameters


For every layer  $l$  :

$$\mathbf{W}^{[l]} = \mathbf{W}^{[l]} - \alpha \frac{\partial L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{W}, \mathbf{b}))}{\partial \mathbf{W}^{[l]}}$$

$$\mathbf{b}^{[l]} = \mathbf{b}^{[l]} - \alpha \frac{\partial L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{W}, \mathbf{b}))}{\partial \mathbf{b}^{[l]}}$$

$\alpha$  : **learning rate**

Guarantee to find a local minima





# Gradient Descent (GD) over a set of samples

We have  $M$  samples

Empirical loss on measured data =  $\frac{1}{M} \sum_{i=0}^{M-1} L(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}(\mathbf{W}, \mathbf{b}))$

$$\frac{\partial L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{W}, \mathbf{b}))}{\partial \mathbf{W}^{[l]}} = \frac{1}{M} \sum_{i=0}^{M-1} \frac{\partial L(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}(\mathbf{W}, \mathbf{b}))}{\partial \mathbf{W}^{[l]}}$$

Find best parameters to minimize the mean loss across all training samples

$$\frac{\partial L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{W}, \mathbf{b}))}{\partial \mathbf{b}^{[l]}} = \frac{1}{M} \sum_{i=0}^{M-1} \frac{\partial L(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}(\mathbf{W}, \mathbf{b}))}{\partial \mathbf{b}^{[l]}}$$

# Gradient descent

Initialize weights and bias

for iter in range(t):

Evaluate loss function (forward pass) over all samples

$$L = \frac{1}{M} \sum_{i=0}^{M-1} L(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}(\mathbf{W}, \mathbf{b}))$$

Compute gradient

$$\frac{\partial L}{\partial \mathbf{W}^{[l]}}, \frac{\partial L}{\partial \mathbf{b}^{[l]}}$$

Update parameter

$$\begin{aligned}\mathbf{W}^{[l]} &= \mathbf{W}^{[l]} - \alpha \frac{\partial L}{\partial \mathbf{W}^{[l]}} \\ \mathbf{b}^{[l]} &= \mathbf{b}^{[l]} - \alpha \frac{\partial L}{\partial \mathbf{b}^{[l]}}\end{aligned}$$

# Stochastic Gradient Descent (SGD)

GD performs one parameter update step after going through entire training set

→ Slow if M is large

→ Gradient update may lack “exploration”

Initialize weights and bias

Random shuffle dataset

for epoch in range(E):

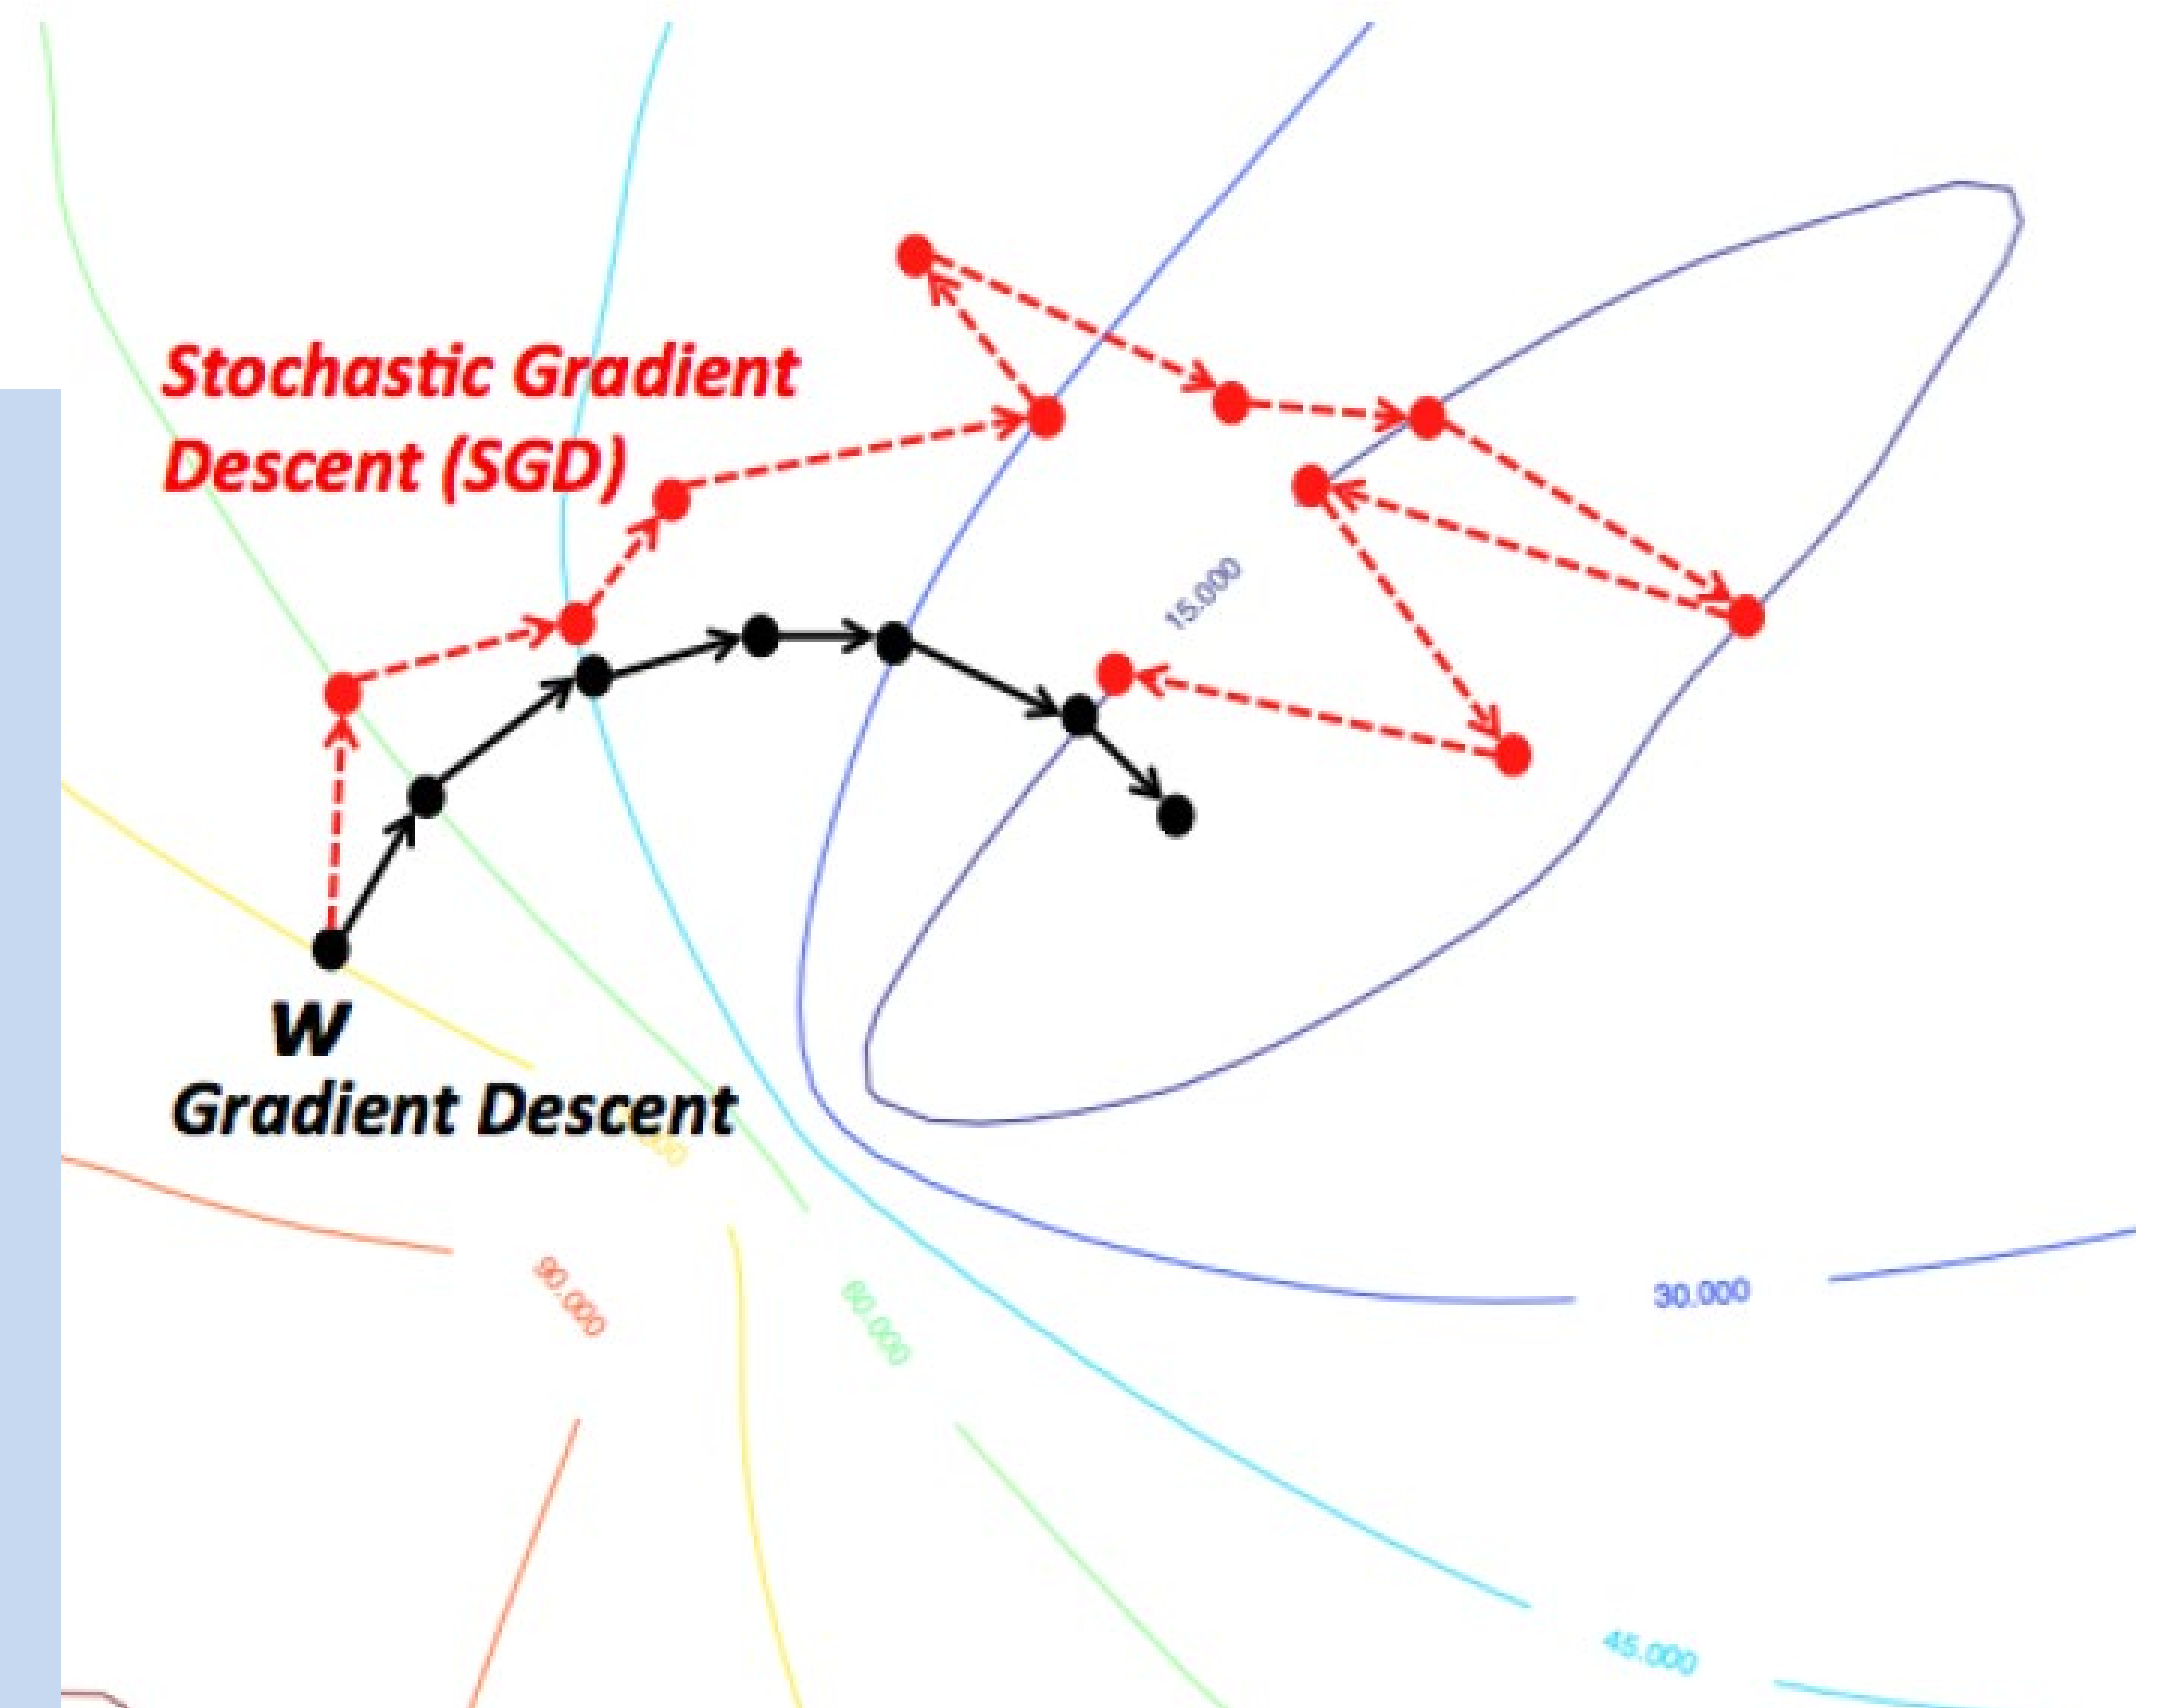
**select one sample with index  $i$**

Evaluate loss function (forward pass) at this sample

$$L = L(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}(\mathbf{W}, \mathbf{b}))$$

Compute gradient  $\frac{\partial L}{\partial \mathbf{W}^{[l]}}, \frac{\partial L}{\partial \mathbf{b}^{[l]}}$

Update parameter  $\mathbf{W}^{[l]} = \mathbf{W}^{[l]} - \alpha \frac{\partial L}{\partial \mathbf{W}^{[l]}}$ ,  $\mathbf{b}^{[l]} = \mathbf{b}^{[l]} - \alpha \frac{\partial L}{\partial \mathbf{b}^{[l]}}$



<https://data-notes.co/the-10-deep-learning-methods-ai-practitioners-need-to-apply-885259f402c1>

# Mini-Batch SGD

SGD can have too much “noise” during convergence  
Not fully utilize the computing hardware

Initialize weights and bias

Random shuffle dataset

**BatchSize (B) = 32**

for epoch in range(E):

**select #BatchSize samples**

Evaluate loss function (forward pass) at this sample **batch**

$$L = \frac{1}{B} \sum_{i=0}^{B-1} L(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}(\mathbf{W}, \mathbf{b}))$$


Compute gradient  $\frac{\partial L}{\partial \mathbf{W}^{[l]}}, \frac{\partial L}{\partial \mathbf{b}^{[l]}}$

Update parameter  $\mathbf{W}^{[l]} = \mathbf{W}^{[l]} - \alpha \frac{\partial L}{\partial \mathbf{W}^{[l]}}, \mathbf{b}^{[l]} = \mathbf{b}^{[l]} - \alpha \frac{\partial L}{\partial \mathbf{b}^{[l]}}$

- Batch size often is limited by the GPU RAM
- Different way to select a batch, e.g. sequential, random, fixed step size etc.
- Every iteration can have a different B



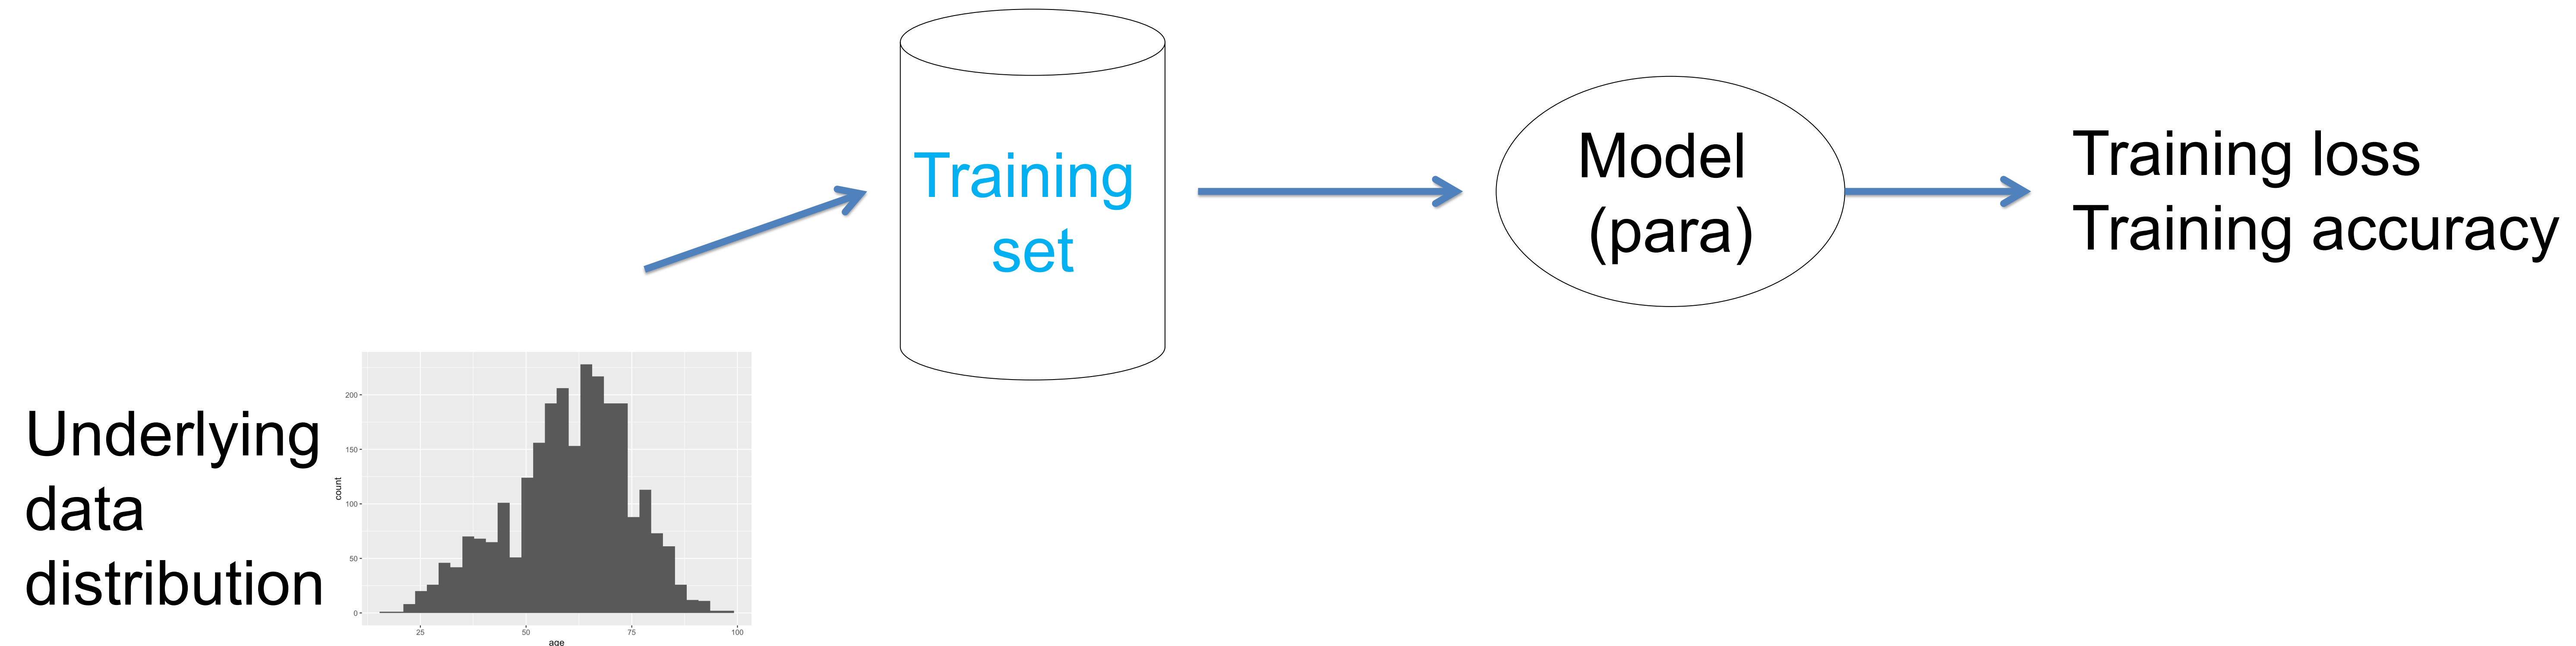
# Larger batch size, higher learning rate



- Larger BatchSize, better estimation of gradient
- Larger BatchSize, less exploration
- High learning rate for small BatchSize can lead to failed convergence
- Overall, set BatchSize large, subject to the GPU RAM limit
- More on how to find good learning rate ...



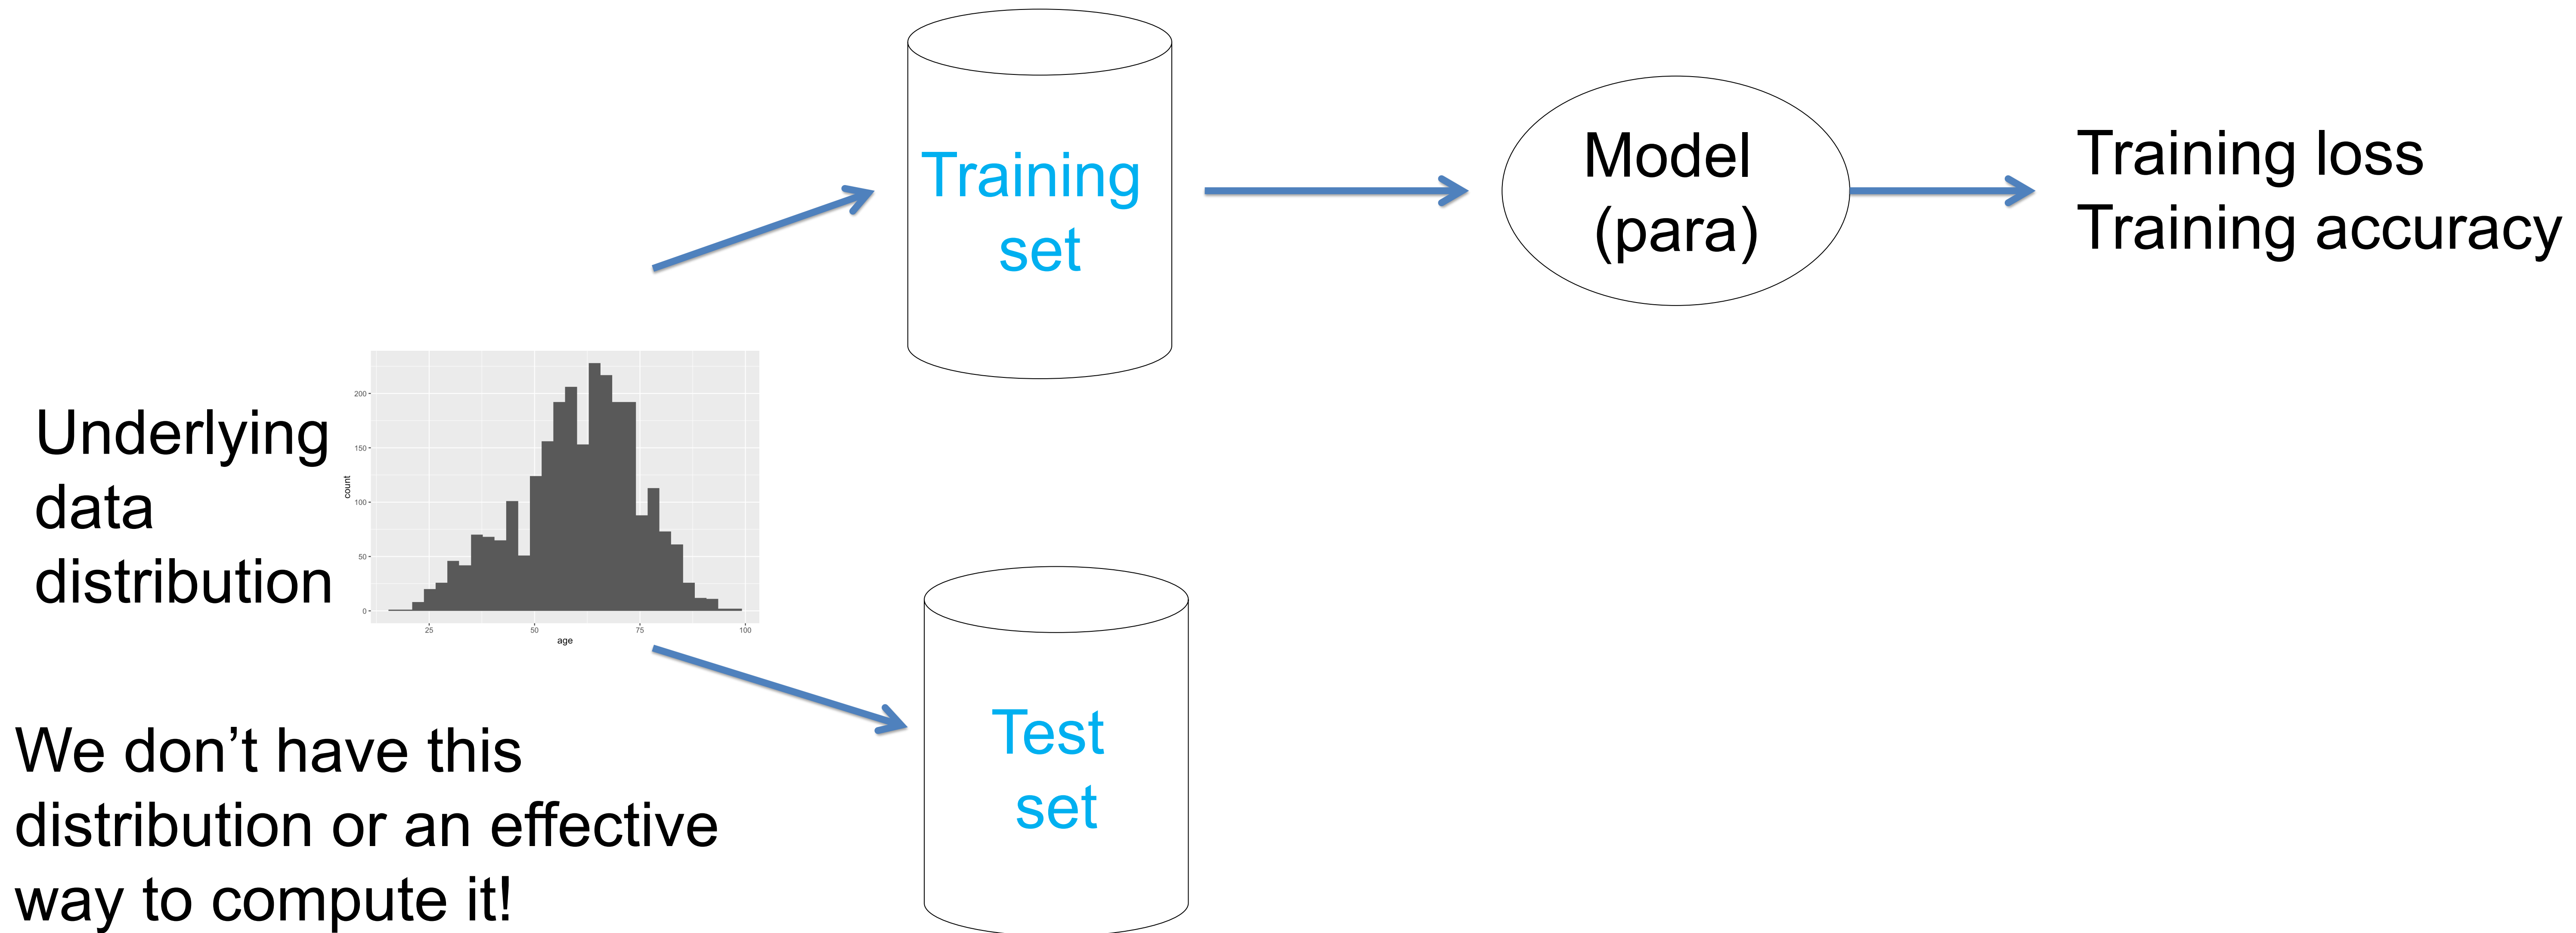
# Model Generalization



We don't have this distribution or an effective way to compute it!

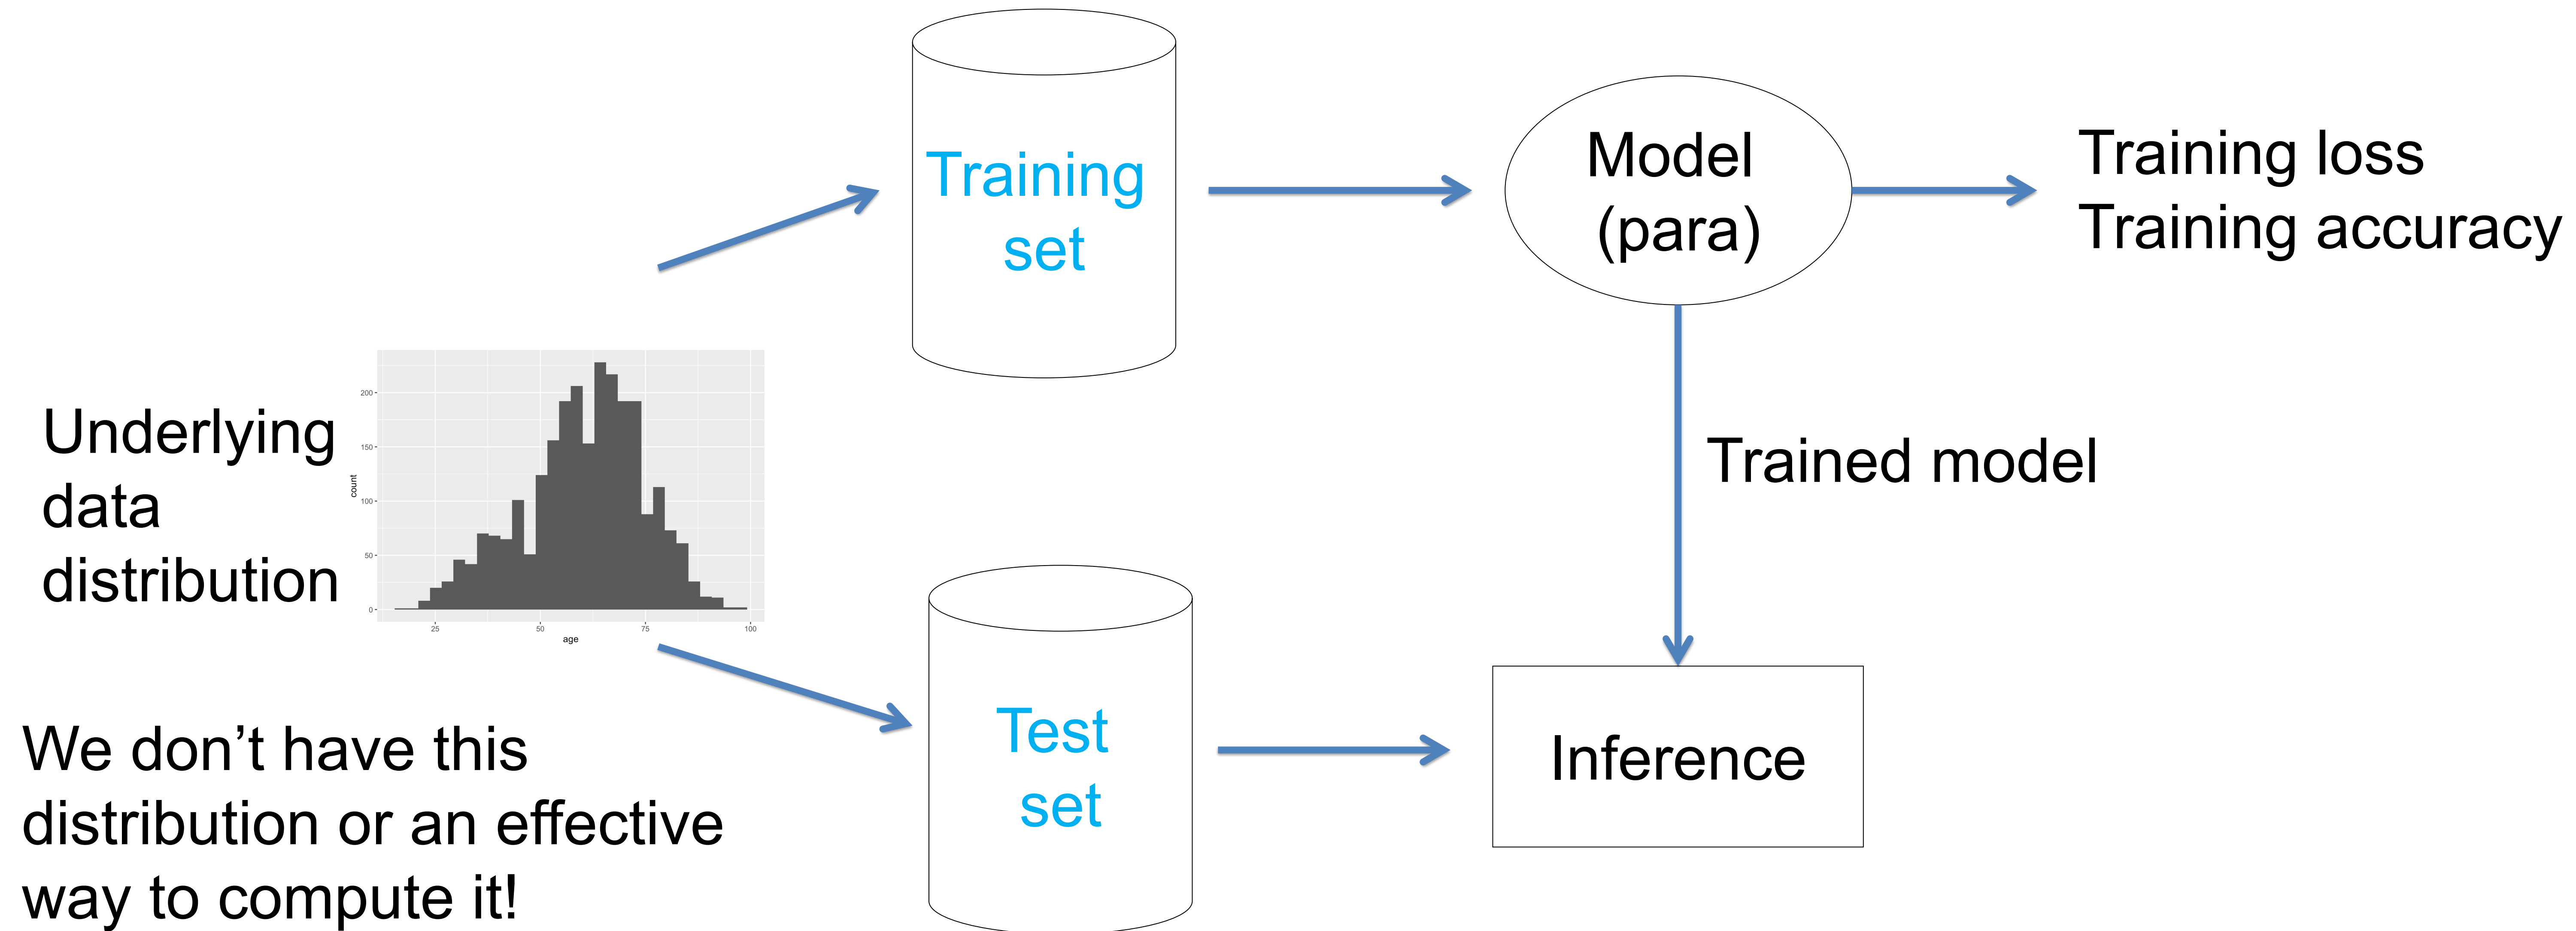
- Training and test sets are sampled from the same distribution, i.i.d. (independent and identically distributed)
- We care test/generalization performance; that is, the performance of training model on a **new** dataset

# Model Generalization



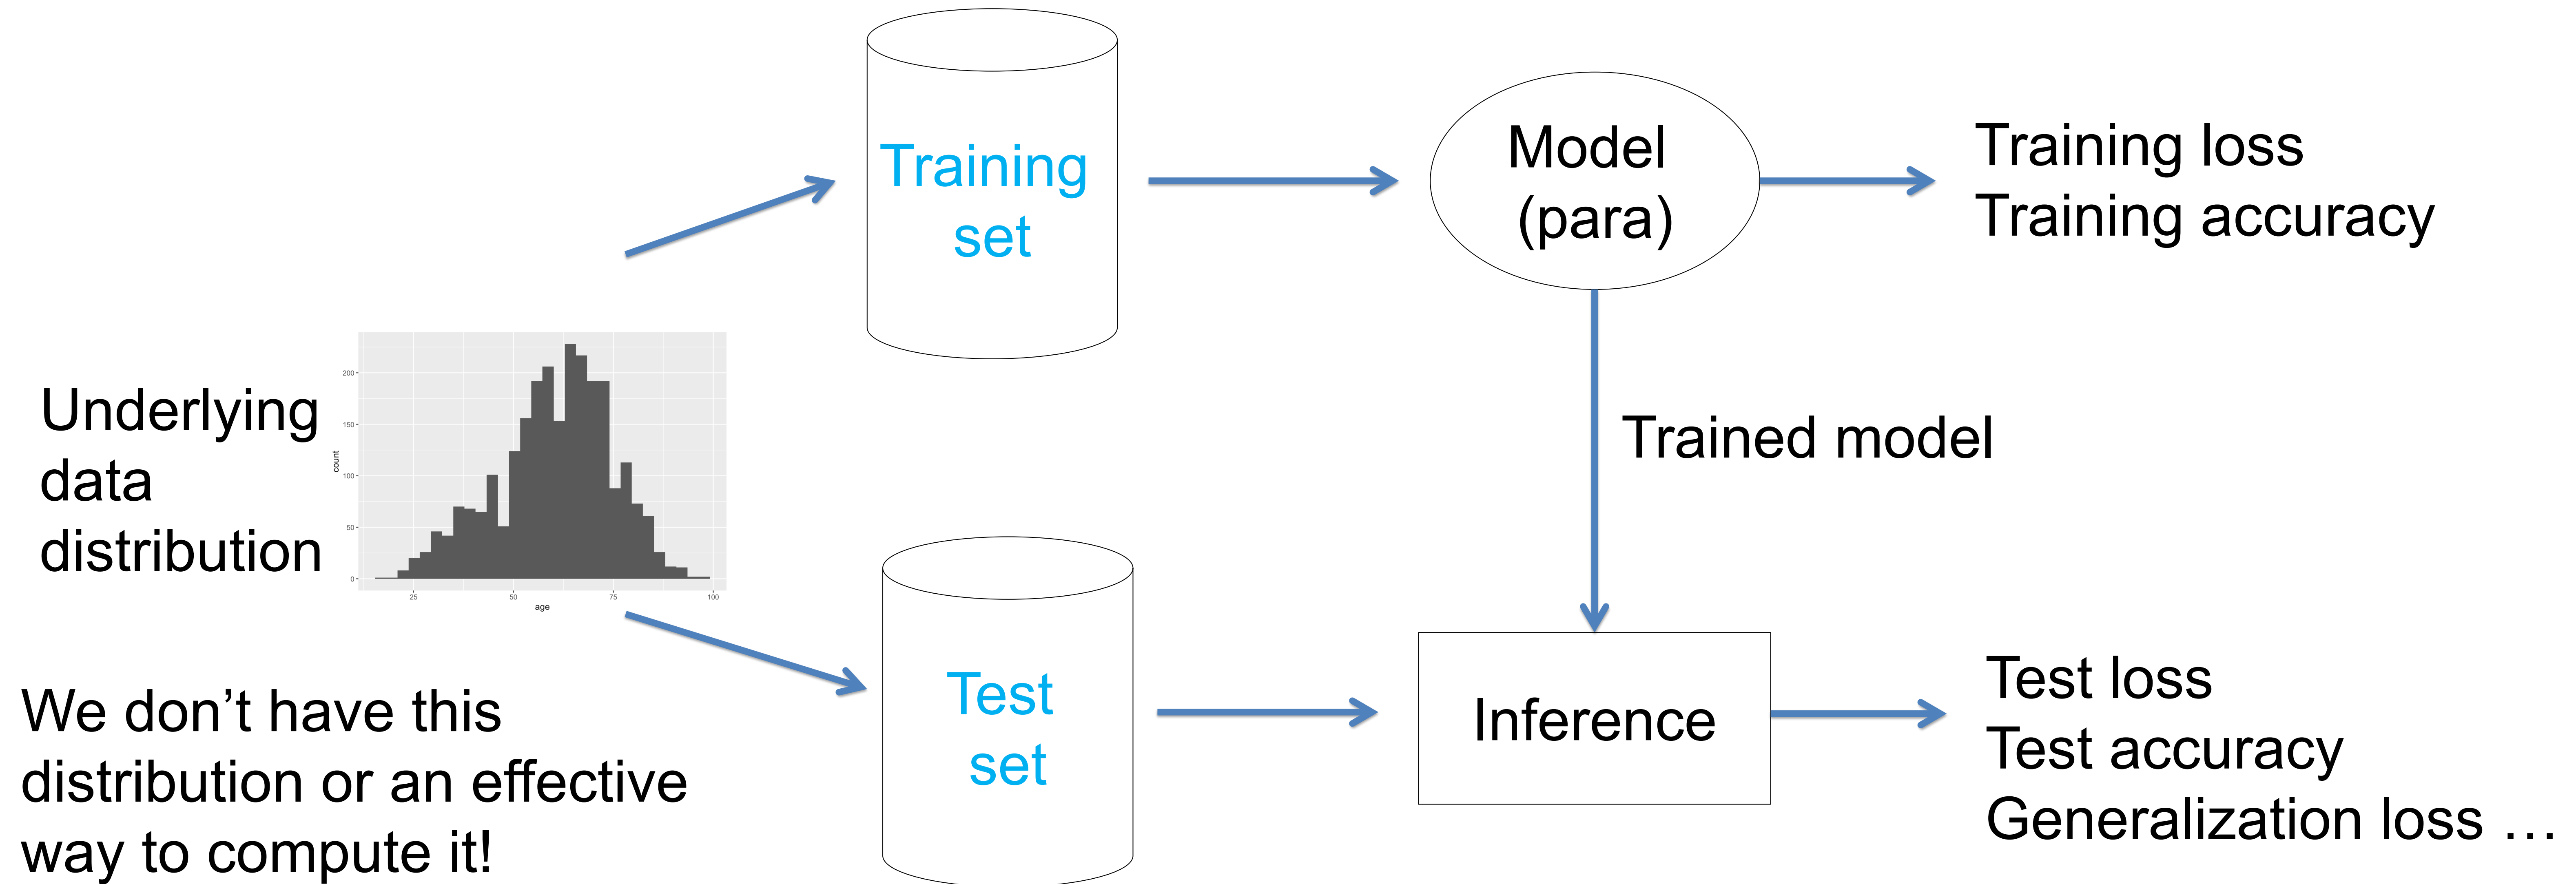
- Training and test sets are sampled from the same distribution, i.i.d. (independent and identically distributed)
- We care test/generalization performance; that is, the performance of training model on a **new** dataset

# Model Generalization



- Training and test sets are sampled from the same distribution, i.i.d. (independent and identically distributed)
- We care test/generalization performance; that is, the performance of training model on a **new** dataset

# Model Generalization



- Training and test sets are sampled from the same distribution, i.i.d. (independent and identically distributed)
- We care test/generalization performance; that is, the performance of training model on a **new** dataset

# Model error

Given the underlying function  $f(\mathbf{x})$ , the model  $M(\mathbf{x}; \mathbf{D})$  to approximate  $f(\mathbf{x})$

$\mathbf{D}$  is a data set sampled from function  $f(\mathbf{x})$ .  $\mathbf{D} = [(\mathbf{x}_0, \mathbf{y}_0), (\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_{N-1}, \mathbf{y}_{N-1})]$

Every sample  $(\mathbf{x}_i, \mathbf{y}_i)$  is contaminated by noise:  $\mathbf{y}_i = f(\mathbf{x}) + \epsilon$ ,  $\epsilon$  is the random noise

We want to know the expected error of model, given the dataset  $\mathbf{D}$ :

$$E_{\mathbf{D}}[(y - M(\mathbf{x}; \mathbf{D}))^2]$$

This error consists of three parts:

$$E_{\mathbf{D}} \left[ (y - M(\mathbf{x}; \mathbf{D}))^2 \right] = \{E_{\mathbf{D}}[M(\mathbf{x}; \mathbf{D})] - f(\mathbf{x})\}^2 + E_{\mathbf{D}} \left[ (E_{\mathbf{D}'}(M(\mathbf{x}; \mathbf{D}')) - M(\mathbf{x}; \mathbf{D}))^2 \right] + \sigma^2$$



# Model error

Model prediction error consists of three parts:

$$E_D \left[ (y - M(\mathbf{x}; \mathbf{D}))^2 \right] = \{E_D[M(\mathbf{x}; \mathbf{D})] - f(\mathbf{x})\}^2 + E_D \left[ (E_D(M(\mathbf{x}; \mathbf{D})) - M(\mathbf{x}; \mathbf{D}))^2 \right] + \sigma^2$$

$Bias(\mathbf{M}, \mathbf{D}) = E_D[M(\mathbf{x}; \mathbf{D})] - f(\mathbf{x})$  This is the **Bias**, for the difference between the mean model performance and ground-truth

$E_D[M(\mathbf{x}; \mathbf{D})]$  is the expected model performance over all possible datasets <- the best model we can get

# Model error

Model prediction error consists of three parts:

$$E_D \left[ (y - M(\mathbf{x}; \mathbf{D}))^2 \right] = \{E_D[M(\mathbf{x}; \mathbf{D})] - f(\mathbf{x})\}^2 + E_D \left[ (E_{D'}(M(\mathbf{x}; \mathbf{D}')) - M(\mathbf{x}; \mathbf{D}))^2 \right] + \sigma^2$$

$Var(M, D) = E_D \left[ [E_{D'}(M(\mathbf{x}; \mathbf{D}')) - M(\mathbf{x}; \mathbf{D})]^2 \right]$  This is the **Variance**, measuring model performance fluctuation due to **different datasets**.

Measure how much the model prediction can change, after trained with different training sets

# Model error

Model prediction error consists of three parts:


$$E_D \left[ (y - M(\mathbf{x}; \mathbf{D}))^2 \right] = \{E_D[M(\mathbf{x}; \mathbf{D})] - f(\mathbf{x})\}^2 + E_D[(E_{D'}(M(\mathbf{x}; \mathbf{D}')) - M(\mathbf{x}; \mathbf{D}))^2] + \sigma^2$$

$\sigma^2$  Bayes error, irreducible error

lowest possible error rate for any classifier

If one would know exactly what process/distribution generates the data, one still cannot achieve 100% accuracy, due to randomness

If the test error has a lower-bound, there will be a Bias-Variance trade-off.




Distribution of two classes can overlap

# Underfitting and overfitting

**Polynomial fit degree 1**

Training error: 0.4

Generalization error: 0.42




Underfit

Model cannot represent the data distribution

**Polynomial fit degree 4**

Training error: 0.14

Generalization error: 0.17




Good fit

Model well represents the data distribution, and does not capture data noise

**Polynomial fit degree 20**

Training error: 0.07

Generalization error: 2000

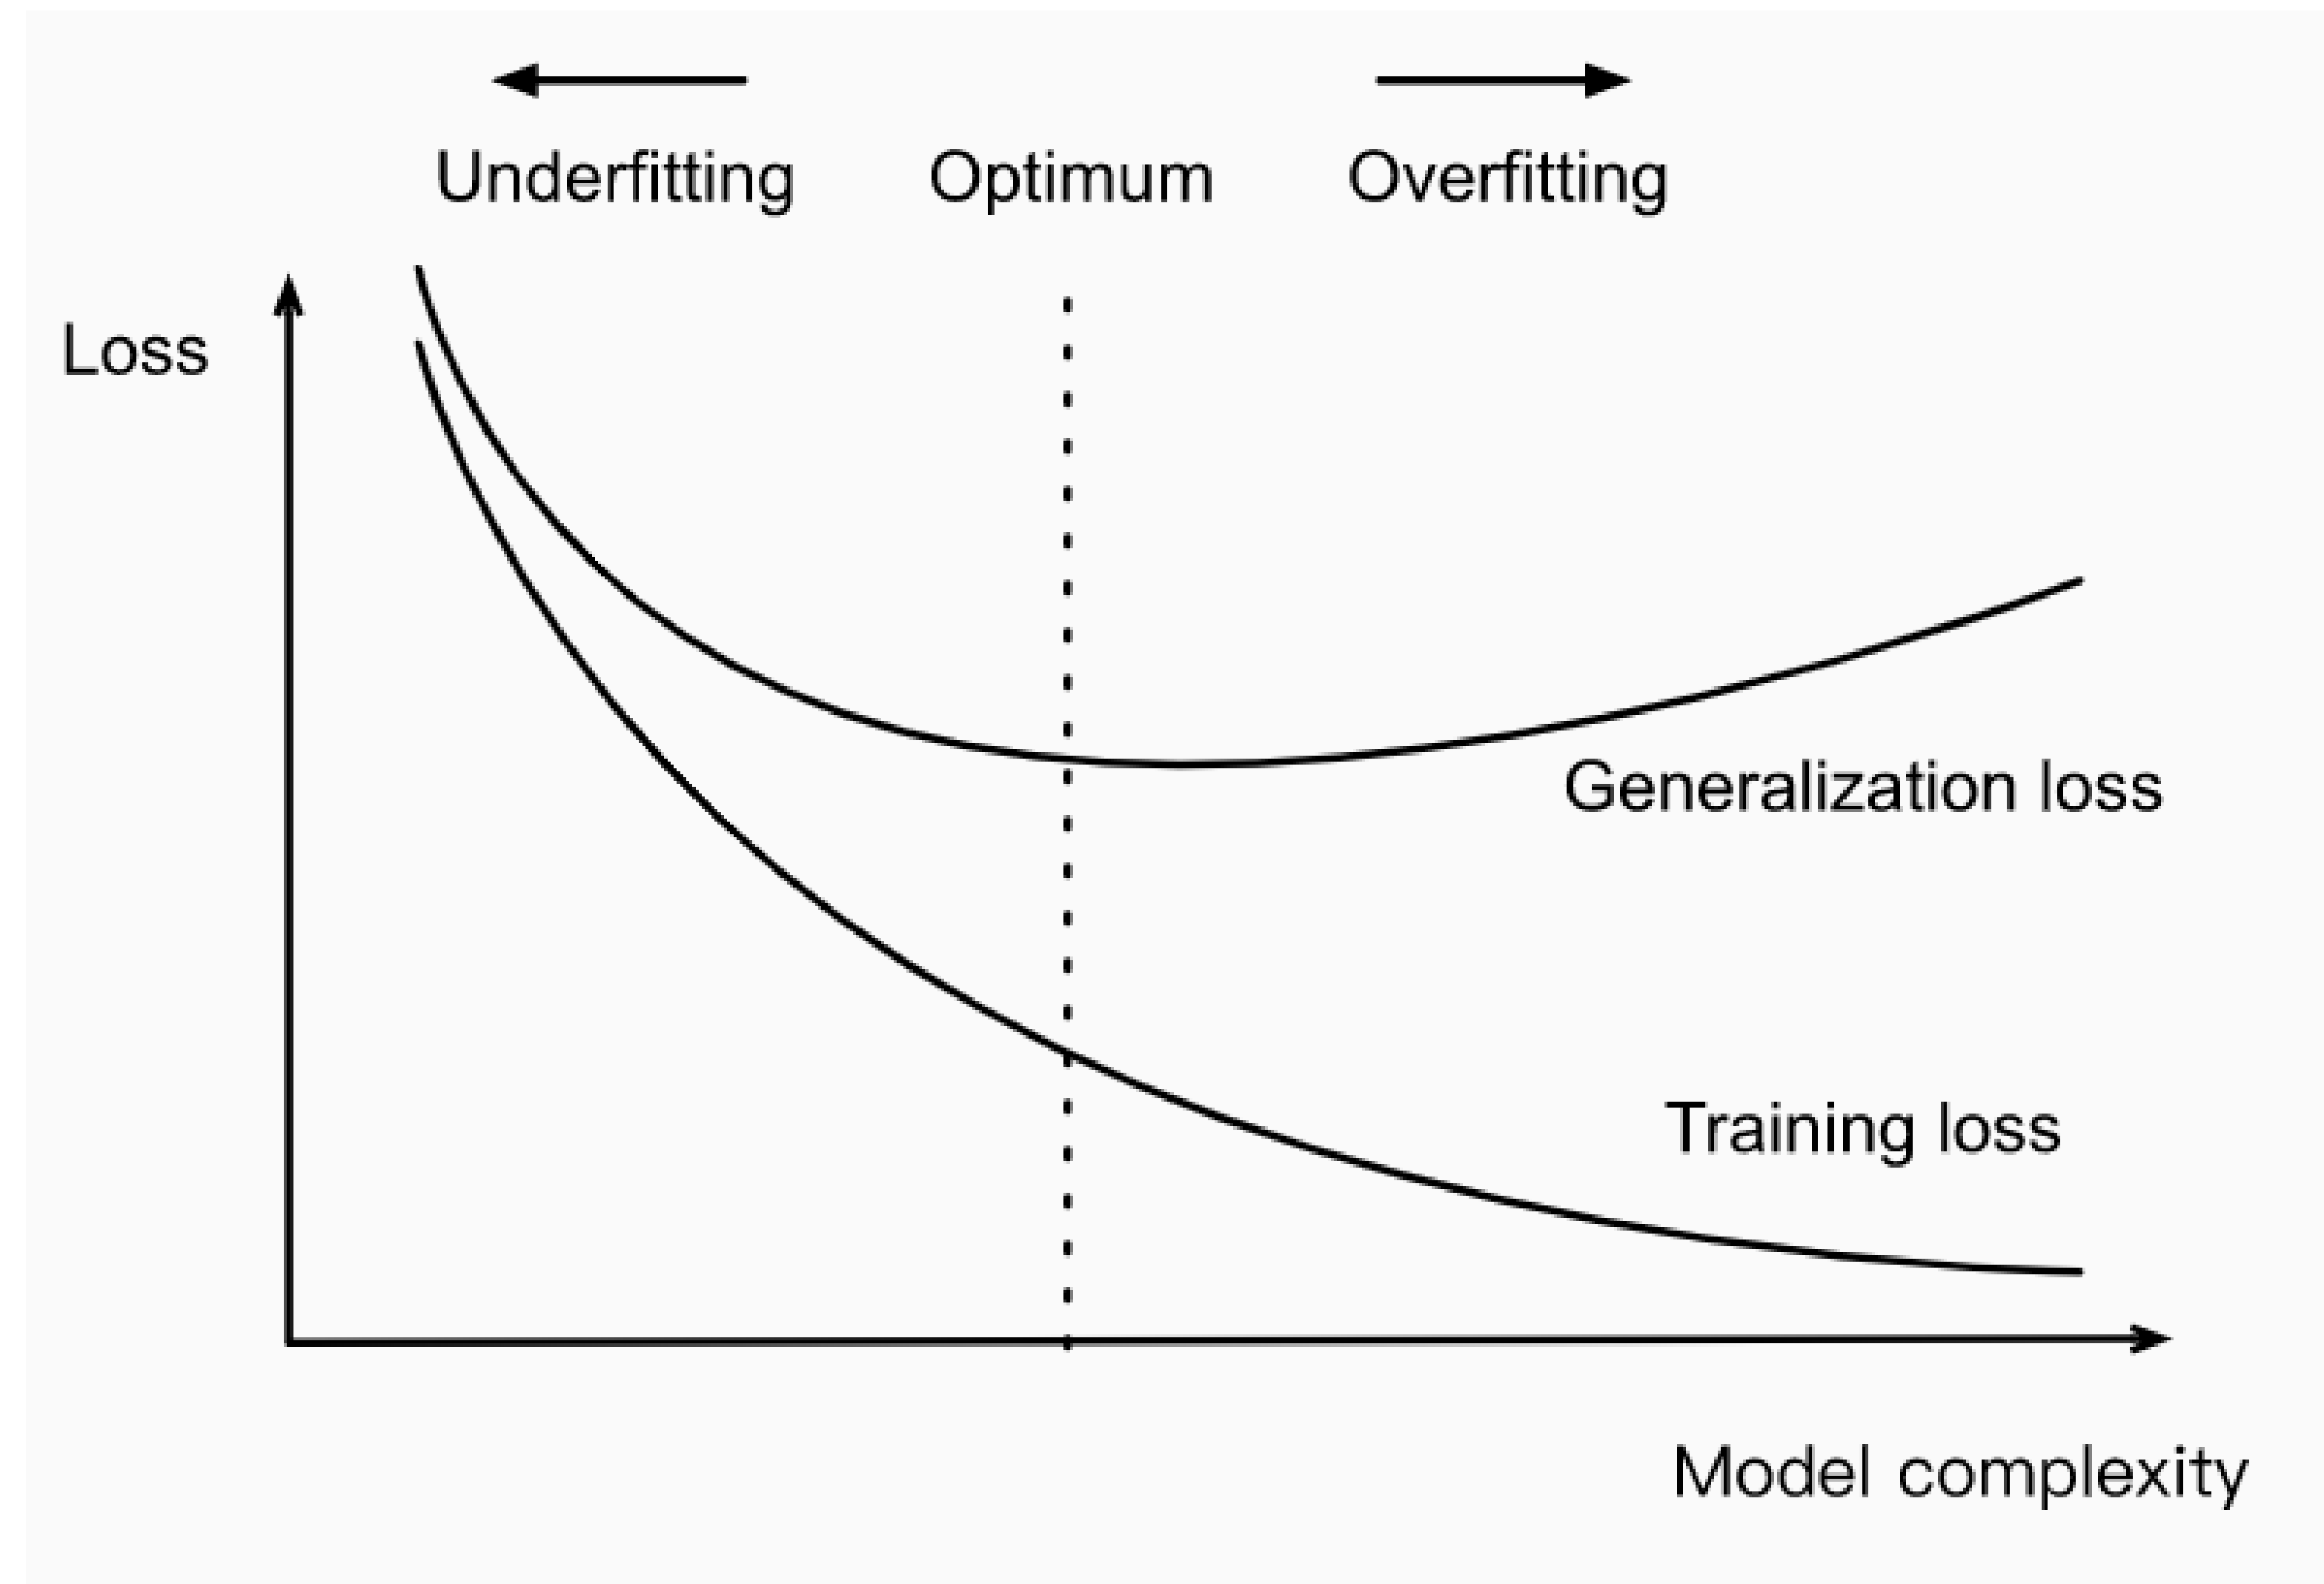


Overfit

Model is so flexible that it captures unwanted fluctuation due to noise

<https://ascpt.onlinelibrary.wiley.com/doi/10.1002/cpt.1796>


# Underfitting and overfitting



- Increasing model capacity/complexity can lead to overfitting
- When applying trained model to a new dataset, e.g. test set, model performance can decrease, as a result of overfitting, indicated by the high **generalization loss**
- Model can also underfitting the data, indicated by the high **training loss**



# In Deep Learning set up





- We train on Training set and use training accuracy to estimate model performance  $\leftarrow$  **Bias**.
- We apply the trained model on Testing set. Performance difference between training accuracy and test accuracy gives an estimation of **Variance**.

# Bias and variance

Bayes error is the best possible error rate if we knew true data distribution. It will not be zero :

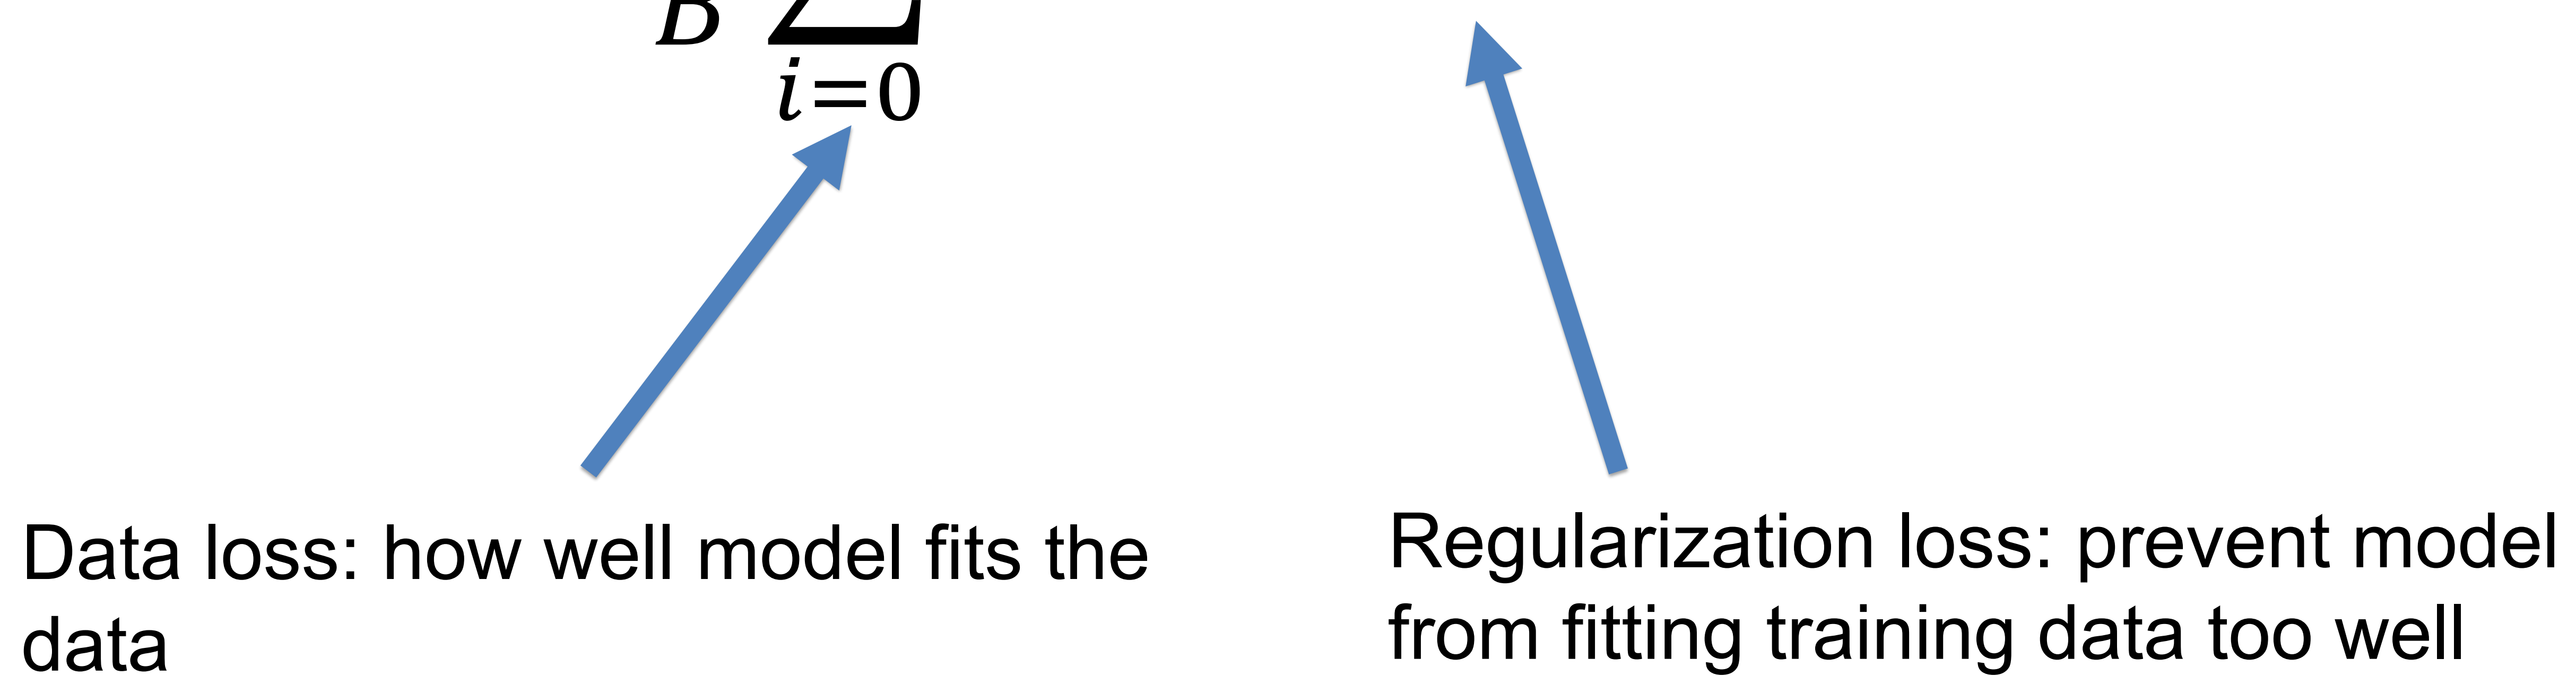
- Overlap between class distribution
- Noise



# Regularization to control model complexity

Deep learning model is very powerful, that it may overfit training data → degraded generalization

IDEA: change loss function to control model complexity

$$L = \frac{1}{B} \sum_{i=0}^{B-1} L^{(i)} + \lambda R(\mathbf{W})$$
The diagram shows the equation  $L = \frac{1}{B} \sum_{i=0}^{B-1} L^{(i)} + \lambda R(\mathbf{W})$ . A blue arrow points from the text 'Data loss: how well model fits the data' to the summation term  $\frac{1}{B} \sum_{i=0}^{B-1} L^{(i)}$ . Another blue arrow points from the text 'Regularization loss: prevent model from fitting training data too well' to the regularization term  $\lambda R(\mathbf{W})$ .

Data loss: how well model fits the data

Regularization loss: prevent model from fitting training data too well

# L2 Regularization, weight decay

$$L = \frac{1}{B} \sum_{i=0}^{B-1} L^{(i)} + \frac{\lambda}{2} \|\mathbf{W}\|_2^2$$

$$\|\mathbf{W}\|_2^2 = \sum_{k=0}^{P-1} w_k^2$$

For all parameters in the model, flatten them and computing the element-wise L2 norm

$$\frac{\partial L}{\partial w_k} = \frac{1}{B} \sum_{i=0}^{B-1} \frac{\partial L^{(i)}}{\partial w_k} + \lambda w_k$$
$$w_k = w_k - \alpha \left( \frac{1}{B} \sum_{i=0}^{B-1} \frac{\partial L^{(i)}}{\partial w_k} \right) - \alpha \lambda w_k$$

weight decay

# L1 Regularization

$$L = \frac{1}{B} \sum_{i=0}^{B-1} L^{(i)} + \frac{\lambda}{2} \|\mathbf{W}\|_1$$

$$\|\mathbf{W}\|_1 = \sum_{k=0}^{P-1} |w_k|$$

For all parameters in the model, flatten them and computing the element-wise absolute value

$$\frac{\partial L}{\partial w_k} = \frac{1}{B} \sum_{i=0}^{B-1} \frac{\partial L^{(i)}}{\partial w_k} + \lambda \text{sign}(w_k) \quad \text{sign}(w_k) = \begin{cases} 1, w_k > 0 \\ 0, w_k == 0 \\ -1, w_k < 0 \end{cases}$$



# Drop Out




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

- Jointly train many smaller network, randomly selected
- Implicitly combine exponentially many different neural network architectures efficiently

- Randomly drop out a set of neurons during training phase, with a probability of  $1-p$  (e.g.  $p=0.5$ )
- It means to randomly select different rows in the  $W$  matrix, for every batch, every epoch
- During the test time, use all neurons, but scale the score by  $p$
- Or, in training time, scale the score by  $1/p \leftarrow$  inverse dropout
- Often used with linear layer, not for convolution
- Require more epochs to train

Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, Ruslan Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. JMLR, 15(56):1929–1958, 2014.

# Other operations with regularization effects

Regularization will:

- Increase training error
- Decrease testing error
- Introduce new hyper-parameters
- Often requires experiments

More on this topic in later  
lecturers

Other operations to improve test error:

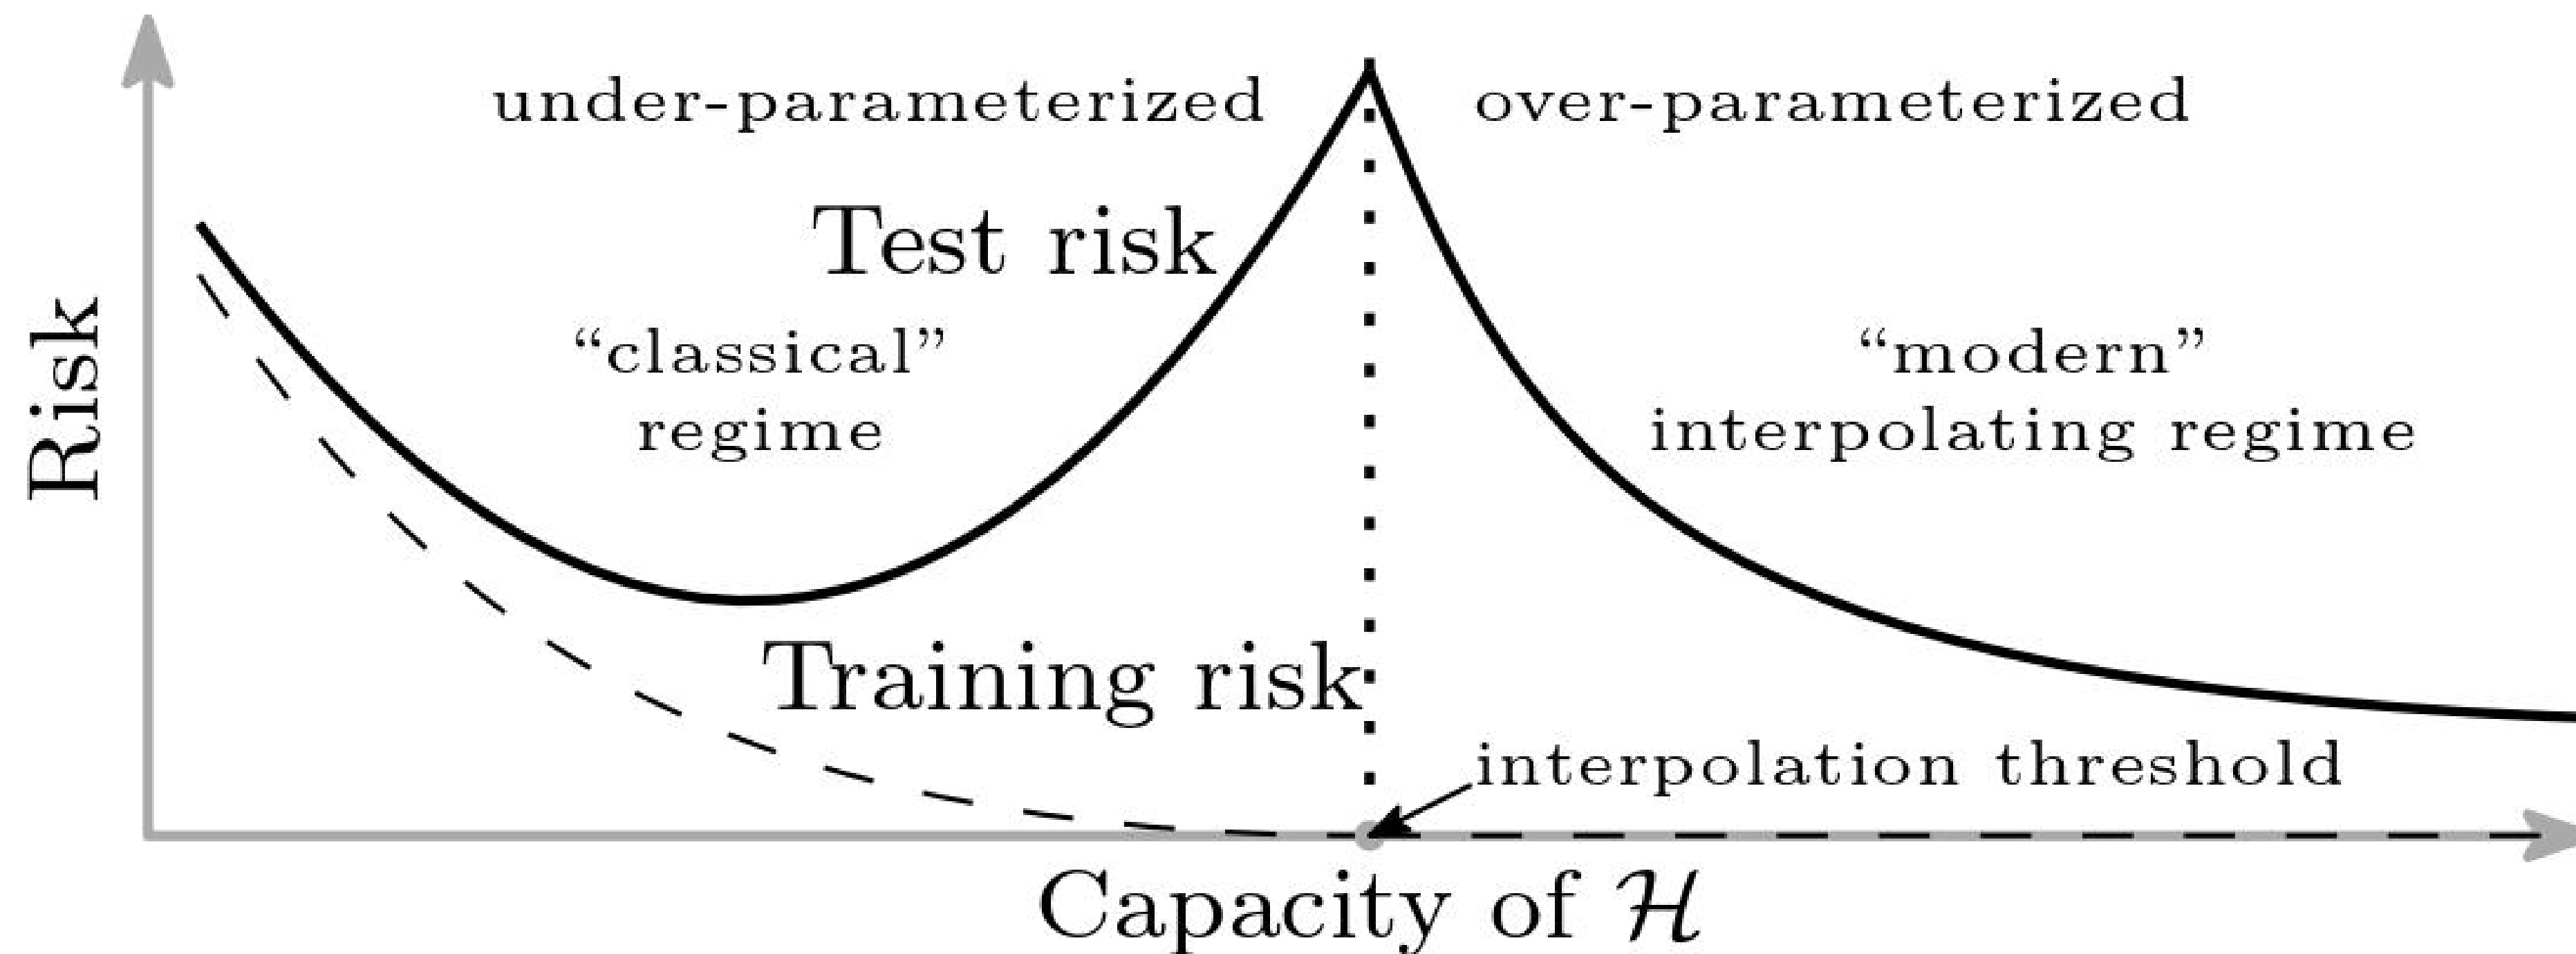
- Data augmentation
- Drop connection, random connection
- Batch/Layer/Channel normalization
- Early stopping

...





# Double-descent of test error in deep models

$$\text{Test error} = \text{Bias}^2 + \text{Variance} + \text{Bayes error}$$



- Total test error can decrease with more complex model and longer training (#epochs)
- Still under research, but for deep model, Bias and variance may not act as a trade-off



<https://arxiv.org/pdf/1912.08286.pdf>



<https://arxiv.org/pdf/1912.02292.pdf>



# Double-descent of test error in deep models

Test error = Bias<sup>2</sup> + Variance + Bayes error

$$\text{Var}(\mathbf{M}, \mathbf{D}) = E_{\mathbf{D}} \left[ (E_{\mathbf{D}'} (M(\mathbf{x}; \mathbf{D}')) - M(\mathbf{x}; \mathbf{D}))^2 \right]$$



- More complex model → reduce variance due to sampling the training set  $\mathbf{D}$ ; may help reduce variance due to optimization
- Longer training → reduce variance due to optimization

<https://arxiv.org/pdf/1912.08286.pdf>  
<https://arxiv.org/pdf/1912.02292.pdf>

