

Deep Learning Crash Course



www.deeplearningcrashcourse.org

Hui Xue

Fall 2021



Outline

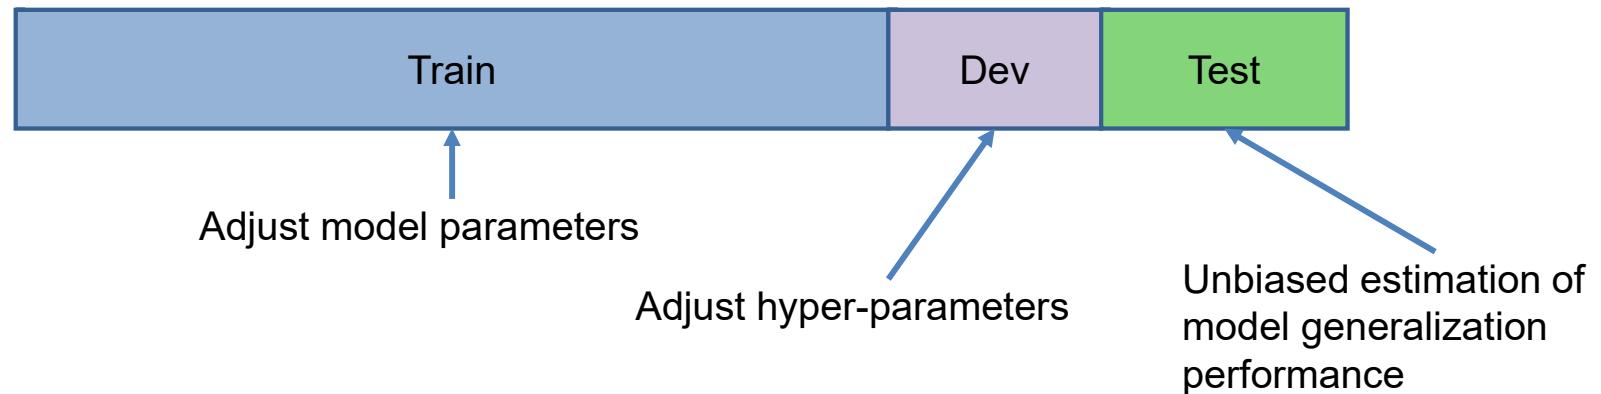
- Set up the training, cont.
- Data pre-processing
- Model weight initialization
- Tips for model training

Model error and training setup

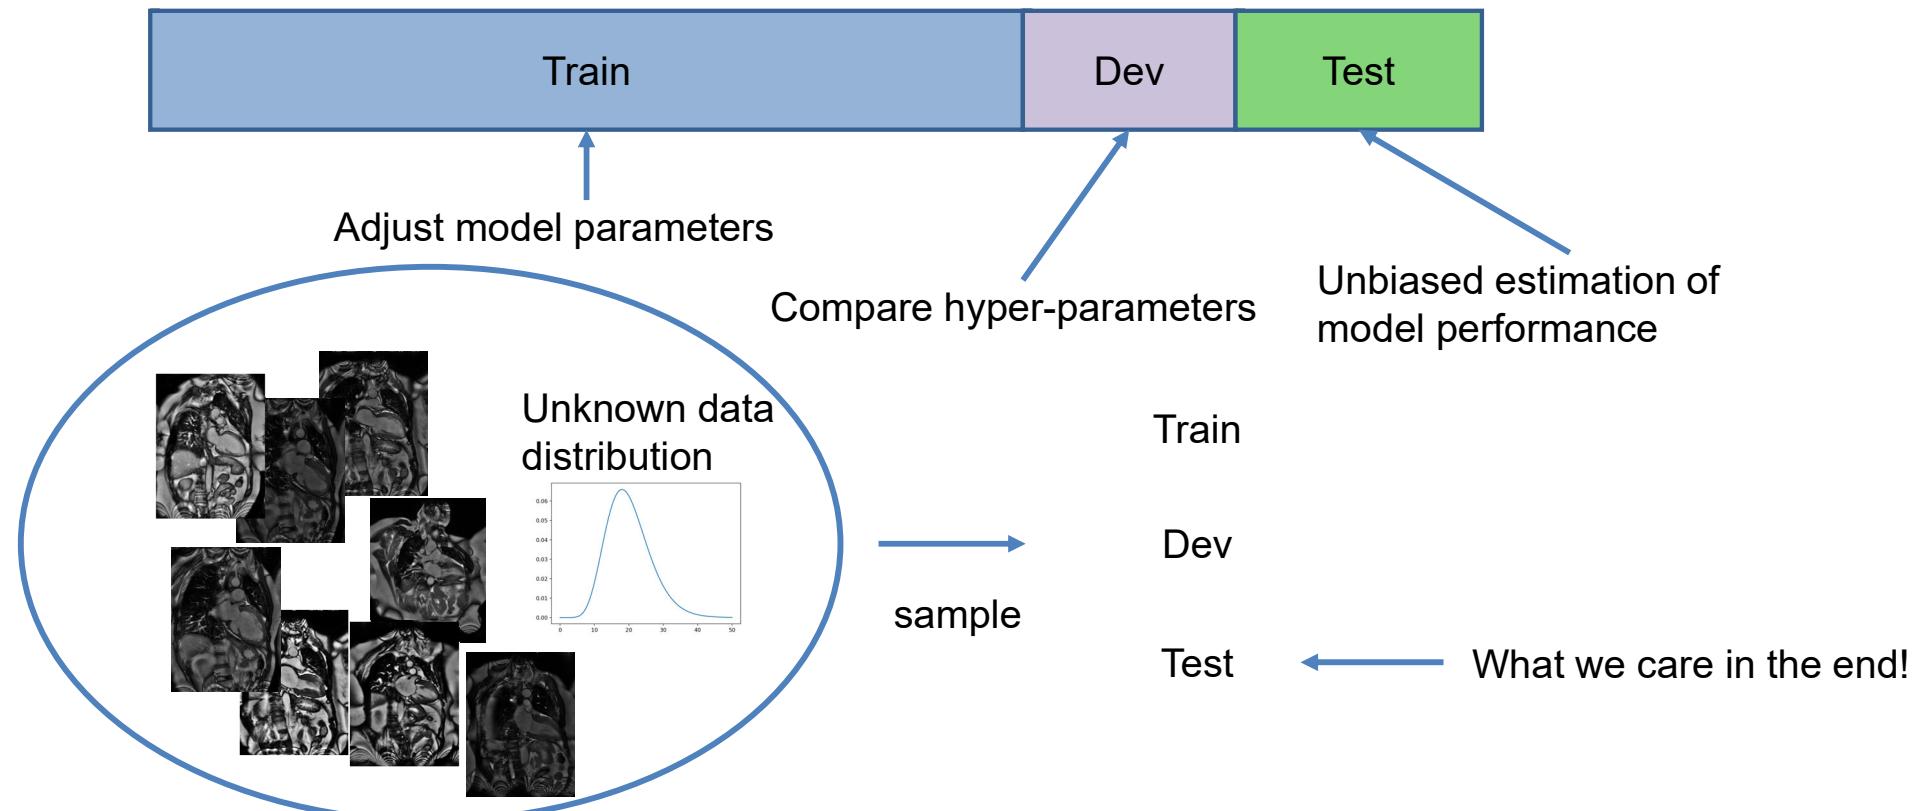
Given a data set, model error is:

$$E_{\mathbf{D}} \left[(y - M(\mathbf{x}; \mathbf{D}))^2 \right] = \underbrace{\{E_{\mathbf{D}}[M(\mathbf{x}; \mathbf{D})] - f(\mathbf{x})\}^2}_{\text{Bias}} + \underbrace{E_{\mathbf{D}}[(E_{\mathbf{D}}(M(\mathbf{x}; \mathbf{D})) - M(\mathbf{x}; \mathbf{D}))^2]}_{\text{Variance}} + \sigma^2$$

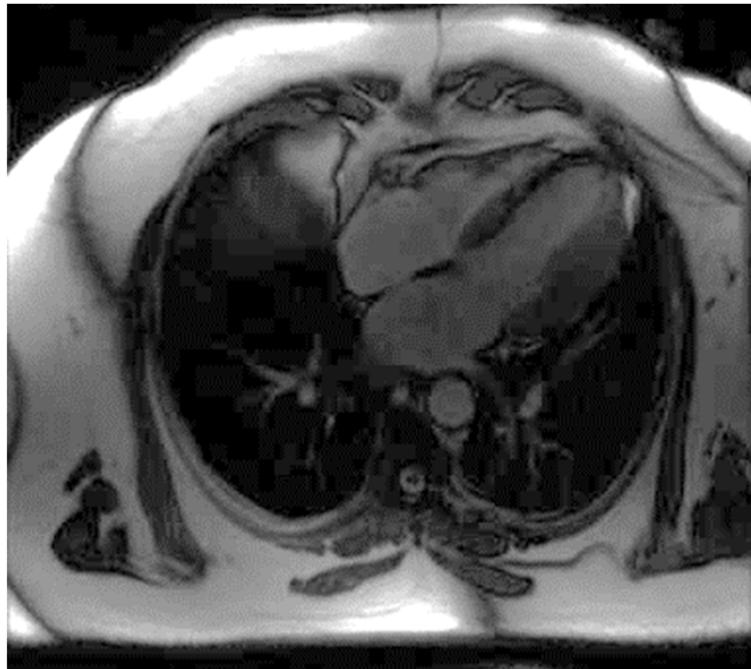
Bayes error



Ideally, Train, Dev, Test sets are from the same data distribution



Test and Train sets can be from different distribution



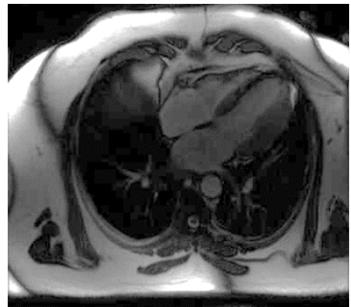
10K, MRI images



1K, CT images

[https://www.clinicalradiologyonline.net/article/S0009-9260\(17\)30058-2/fulltext](https://www.clinicalradiologyonline.net/article/S0009-9260(17)30058-2/fulltext)

Test and Train sets can be from different distribution



10K, MRI images

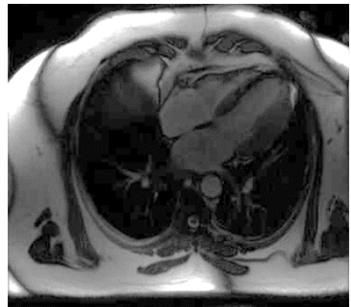


1K, CT images

Goal: Detect whether heart is imaged in MR images

	Option1	Option2	Option3	Option4
Train	8K MRI	8K MRI+1K CT	8K MRI	8K MRI
Dev	1K MRI	1K MRI	1K MRI+1K CT	1K MRI
Test	1K MRI	1K MRI	1K MRI	1K MRI +1K CT

Test and Train sets can be from different distribution



10K, MRI images



1K, CT images

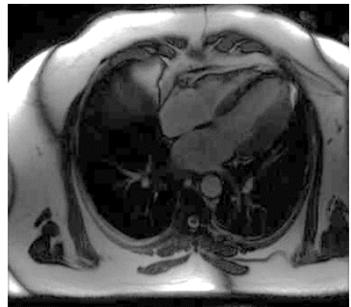
Sample the Dev and Test sets from the same data distribution for the deployment time



Dev and Test sets define the target of the model

https://www.google.com/url?sa=i&url=https%3A%2F%2Ftowardsdatascience.com%2Ftradeoff-bias-or-variance-1409eec38caf&psig=AOvVaw0FzJR_8JxUkPjAiJQaxZY&ust=1620153718709000&source=images&cd=vfe&ved=0CAI0jRxqFwoTCMCroImVnACFQAAAAAdAAAAABAD

Test and Train sets can be from different distribution



10K, MRI images

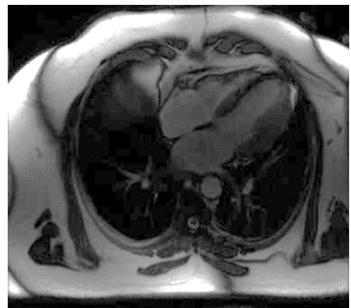


1K, CT images

Goal: Detect whether heart is imaged in **CT** images

	Option1	Option2	Option3
Train	500 CT	10K MRI	10K MRI+500 CT
Dev	250 CT	500 CT	250 CT
Test	250 CT	500 CT	250 CT

Dev and Train set are from the different distribution

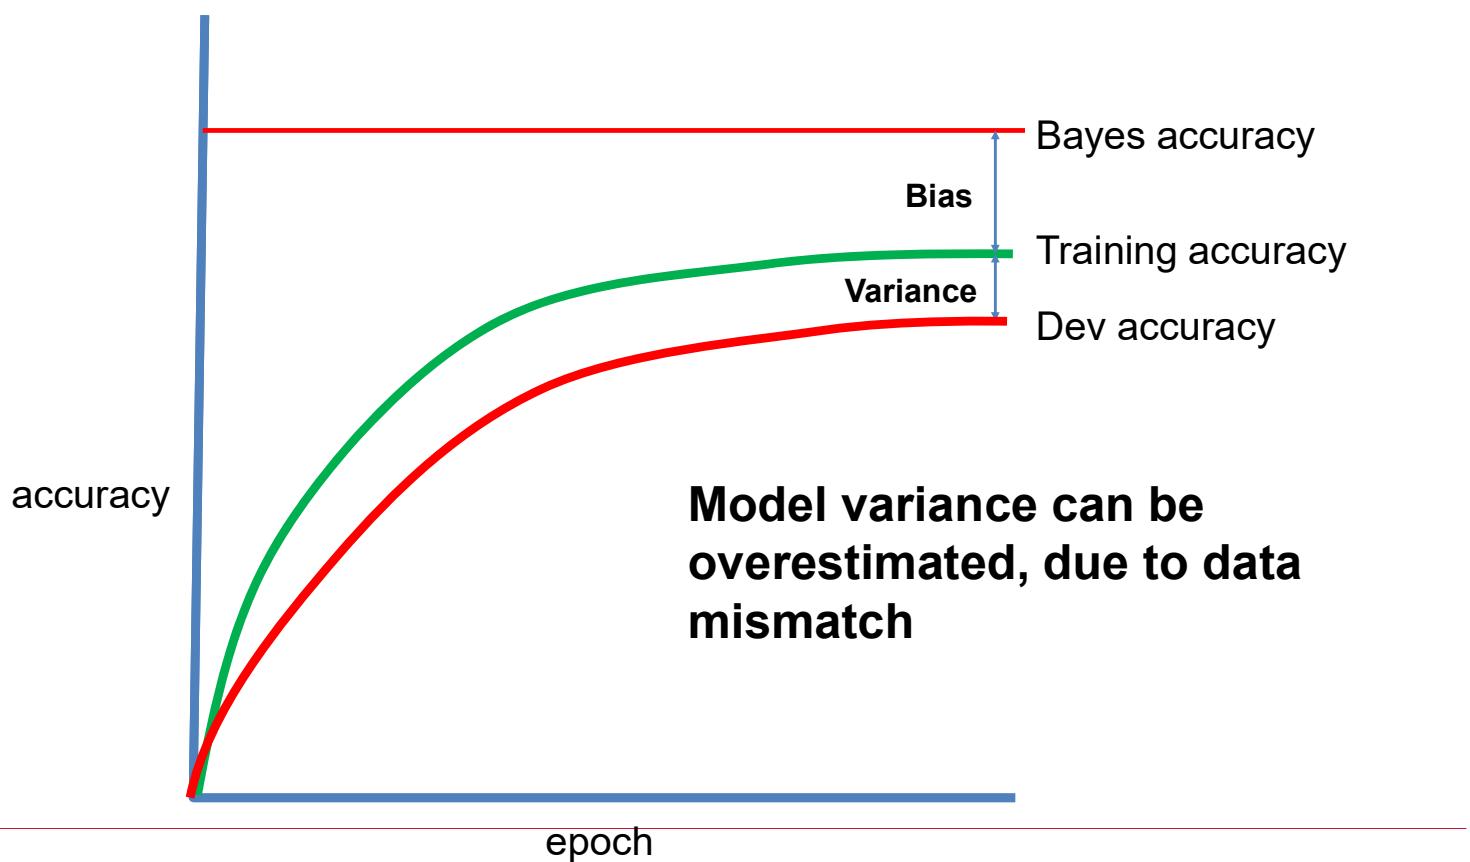


10K, MRI images

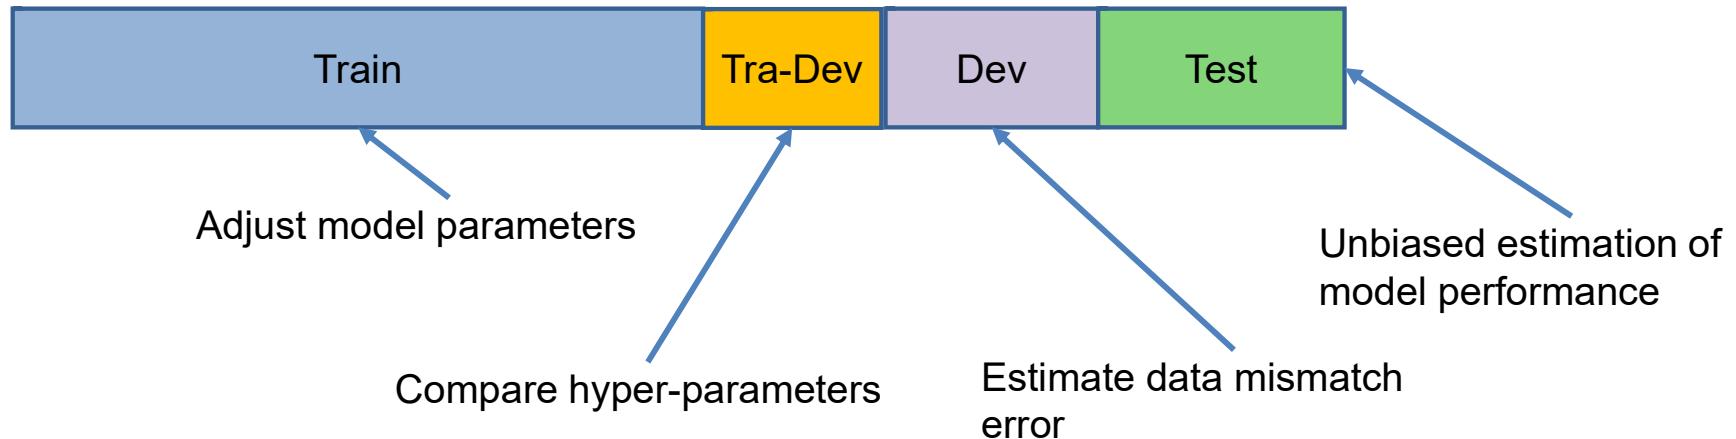


1K, CT images

Goal: Detect whether heart is imaged in **CT** images

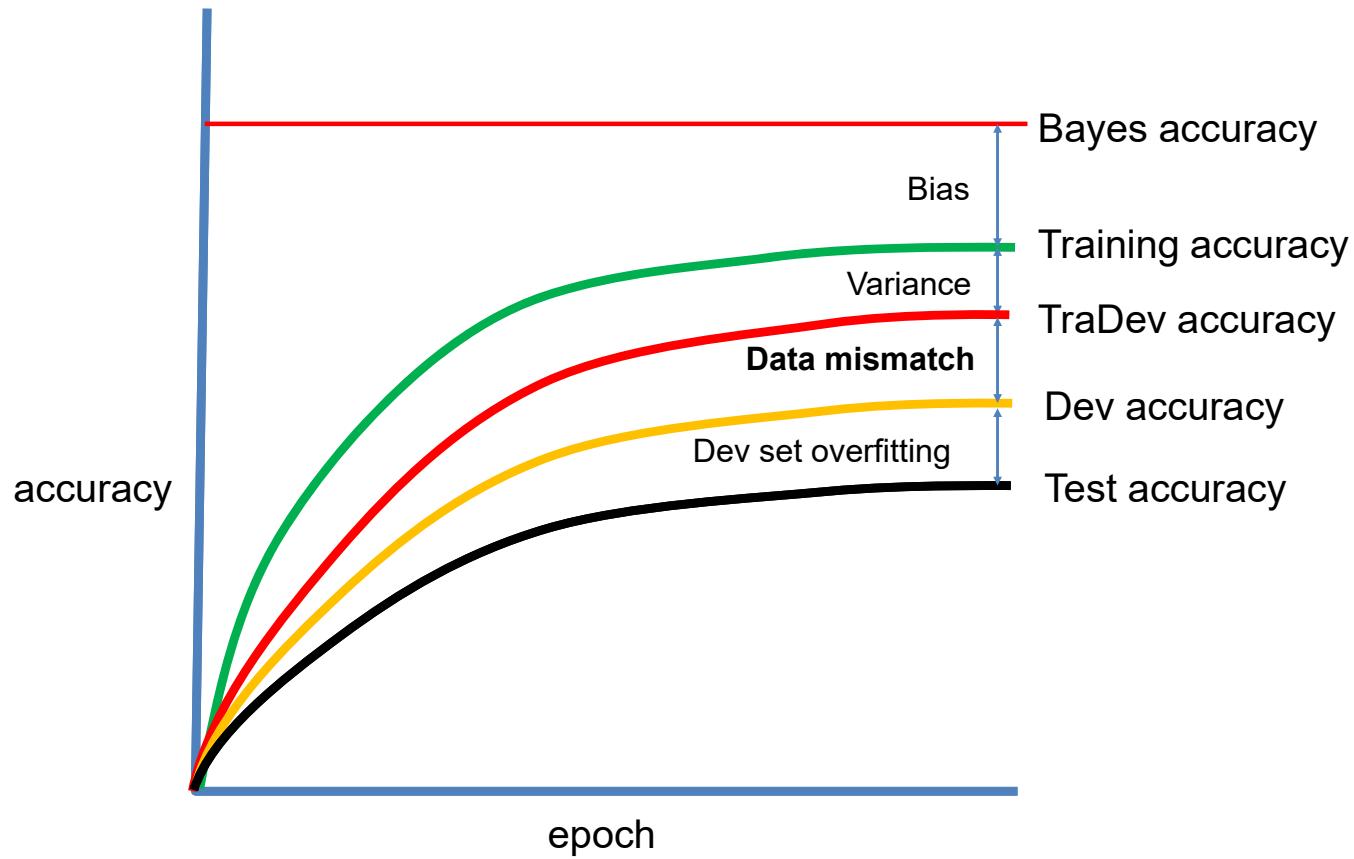


Dev and Train set are from the different distribution



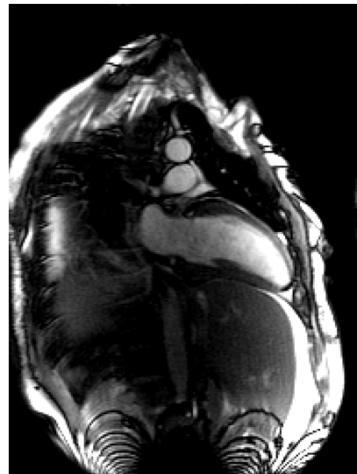
- Split the train set further to get a Tra-Dev set
- Train and Tra-Dev set have the same data distribution
- Used to estimate model variance

Train, Tra-Dev, Dev and Test sets



Spot problems : Error analysis

Error analysis : exam the incorrect/less-perfect samples in Dev set to spot the problems



Typical sample
in Train set

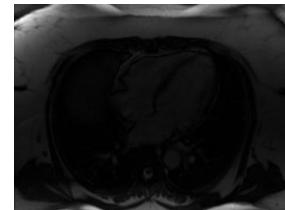
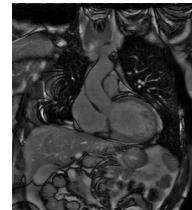


Image is too dark 5%



New image view 11%

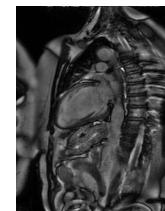
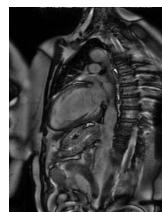
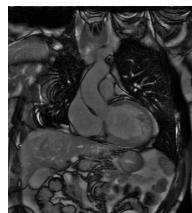
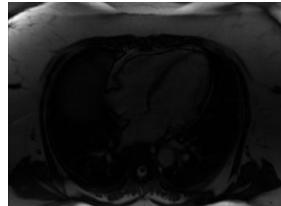


Image is flipped and
with different
contrast 14%

Spot problems : Error analysis

Error analysis : exam the incorrect/less-perfect samples in Dev set to spot the problems



Findings	Error rate	Remedies
Image is too dark	5%	Add pre-processing
New image view	11%	Acquire more data
Image is flipped and with different contrast	14%	Add data augmentation

Iterate error analysis with training

Training Accu	67%
Dev Accu	65%

Correct for first findings: add data pre-processing, data augmentation, add new data

Training Accu	87%
Dev Accu	77%

Increase model depth + L2 Reg

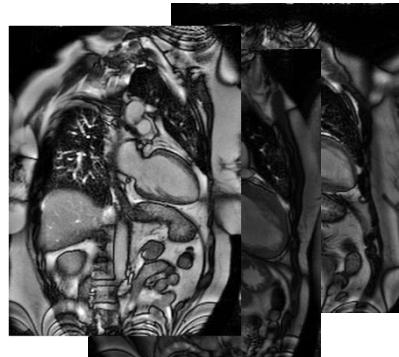
Training Accu	95%
Dev Accu	94%

Second error analysis
....

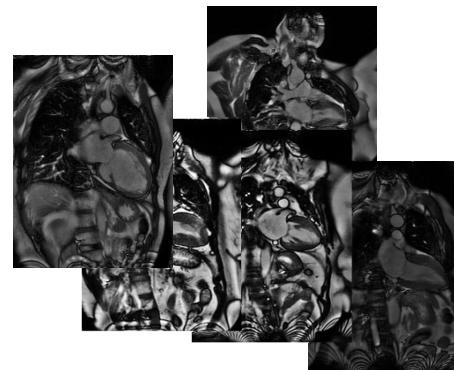
- Human-level-performance: ~ 100% detection rate
- Iterative process
- The key for deep learning is to decide what to do at each step

Until meet the performance requirement

Speedup the error analysis

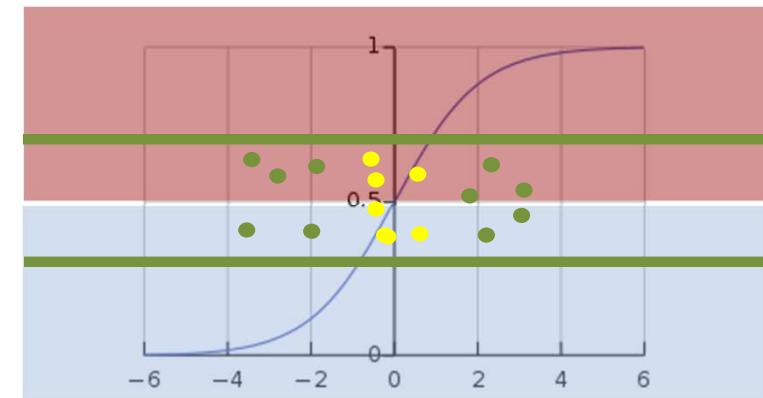


Eyeball set: a portion of dev set to check



Blackbox set: a portion of dev set to apply model, but not exam for error

Active selection : use uncertainty

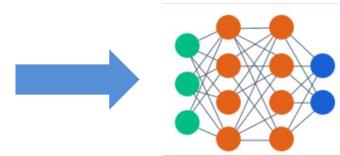


- Use the model output to pick samples to check
- Consider to use different metrics
- Combine domain knowledge, e.g. signal-noise-ratio in different imaging protocols, image appearances in different diseases (such as hypertrophic cardiomyopathy)
- Automate this process and optimize tools

Error analysis of a complex system

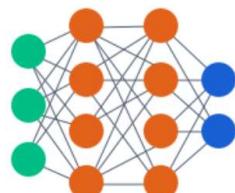


<https://ecg.utah.edu/lesson/3>



R-R interval, location of S-T wave

Model 1



Disease classification

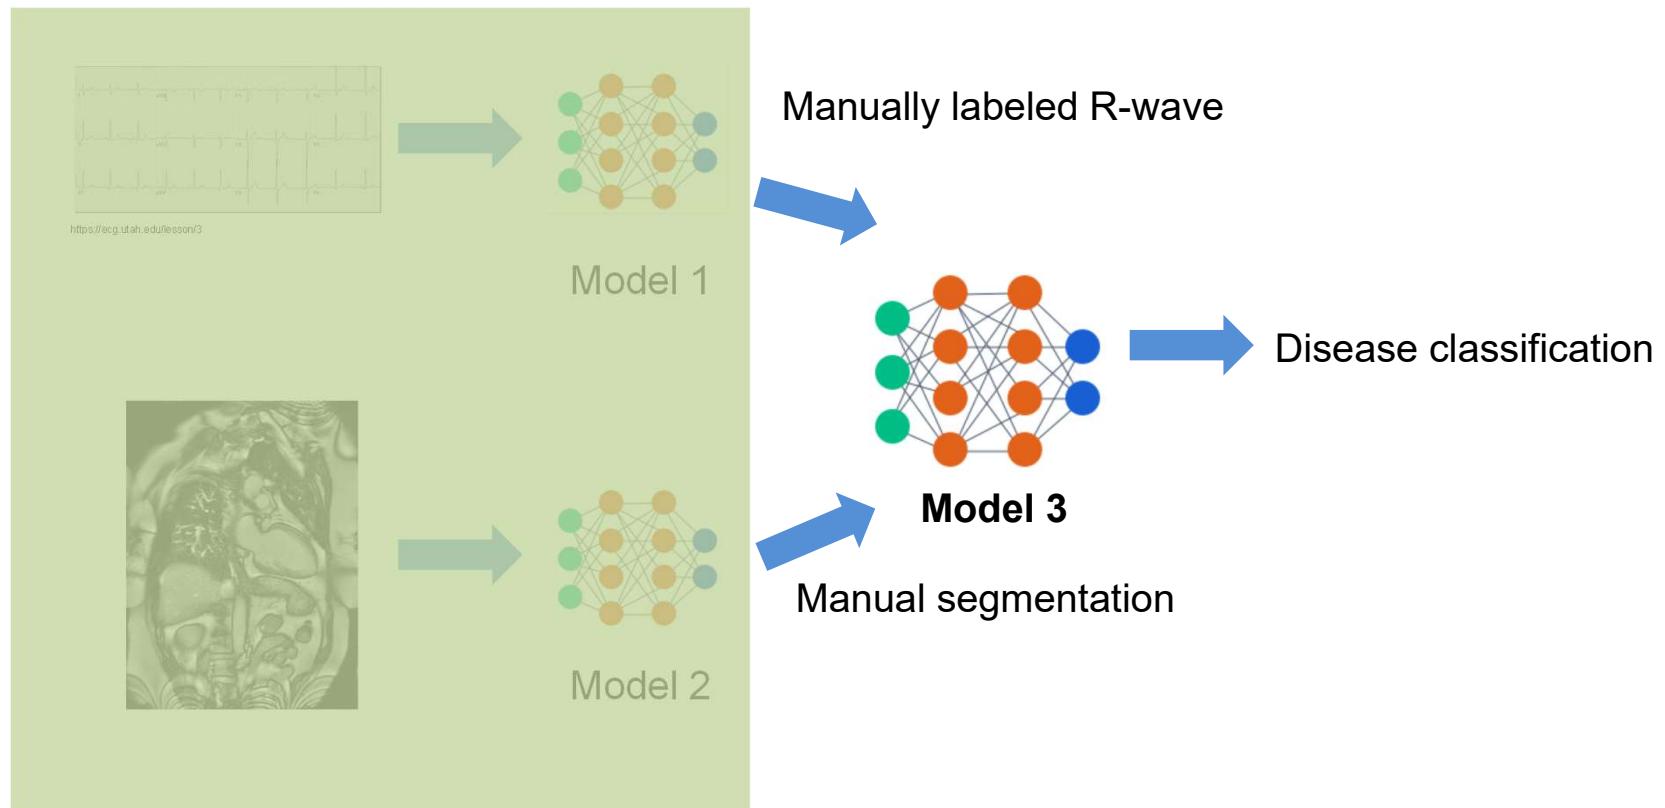


Model 2

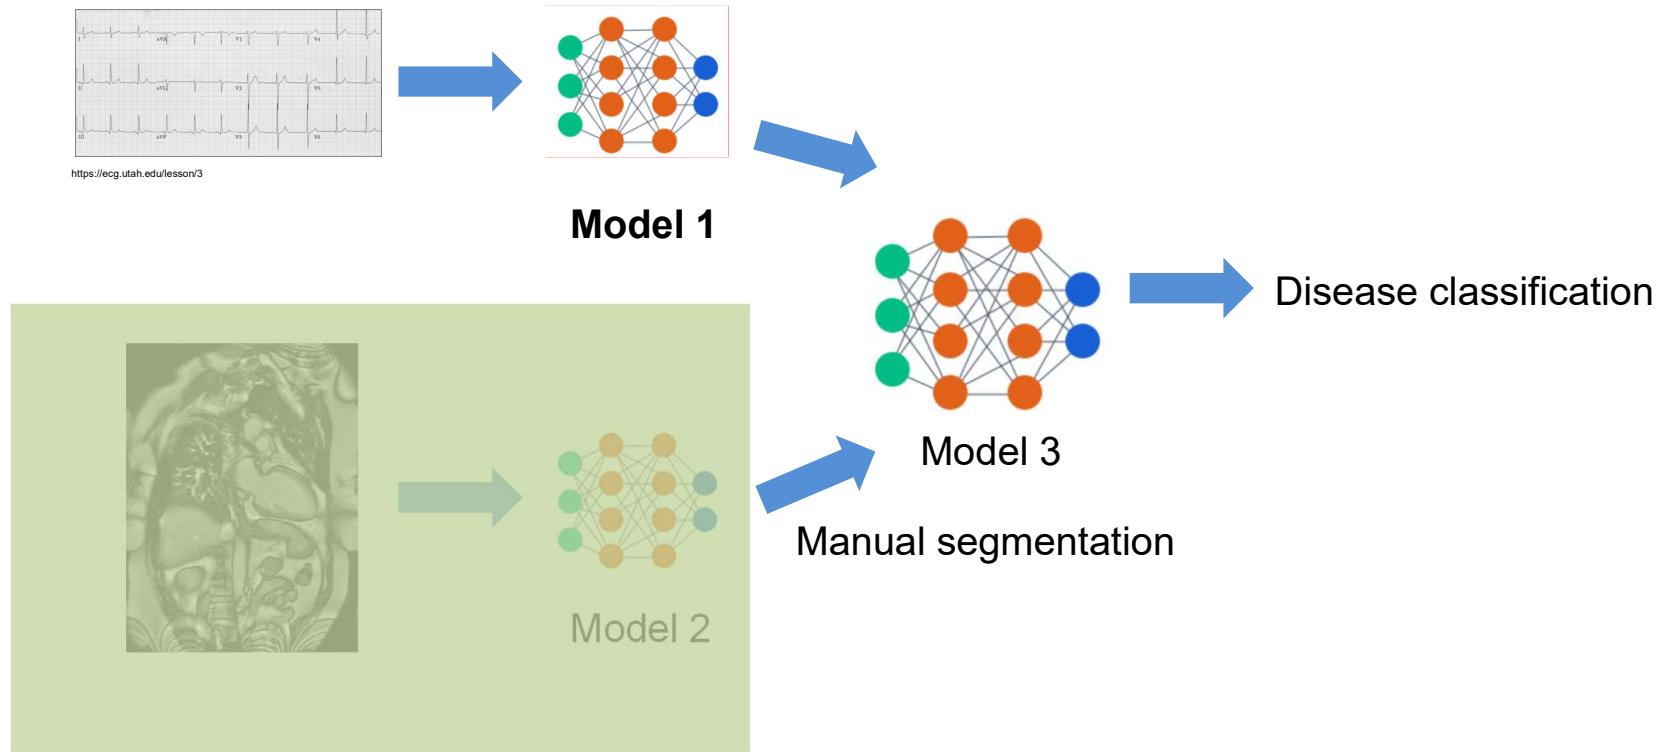
Segmentation of heart

Model 3

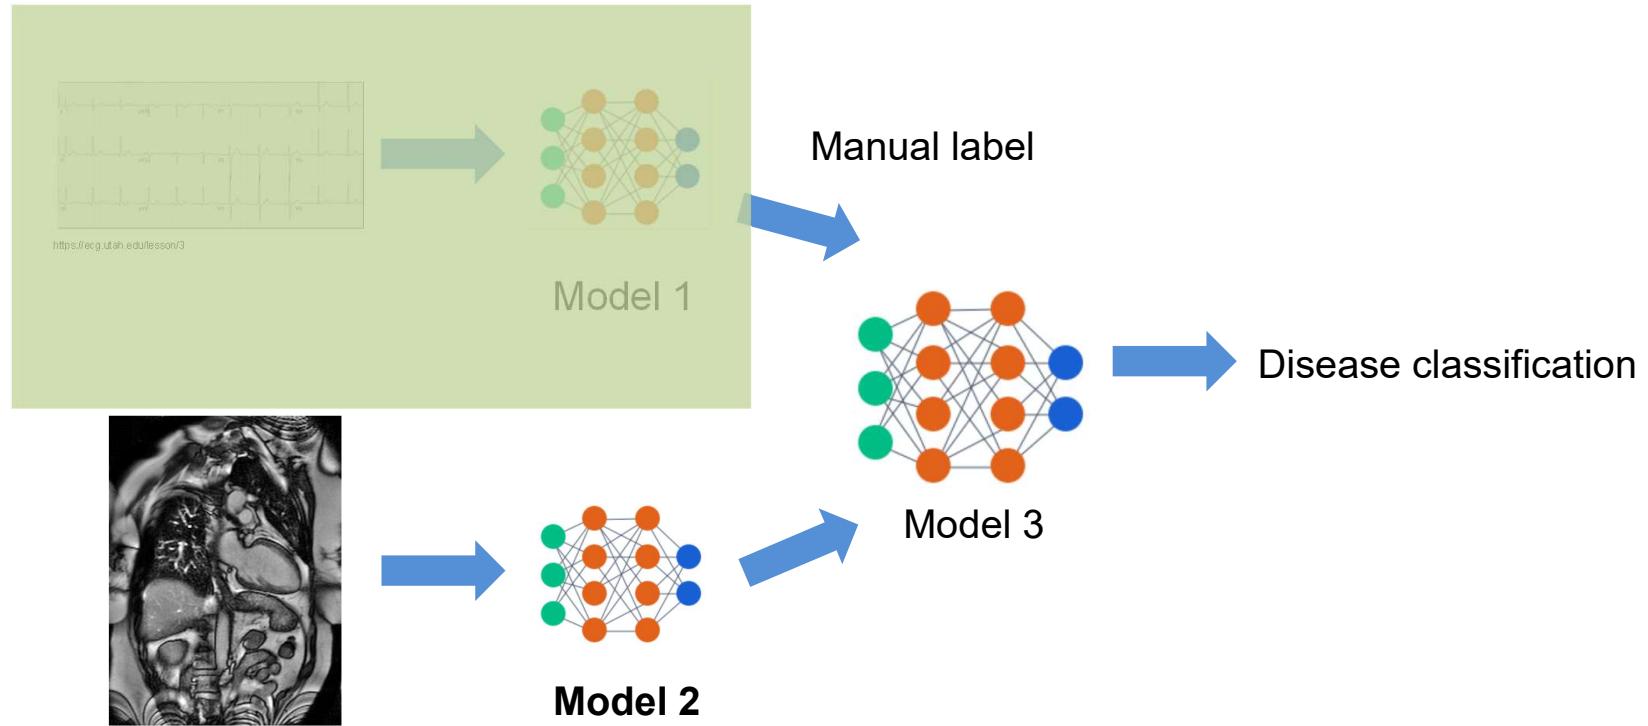
Error analysis of a complex system



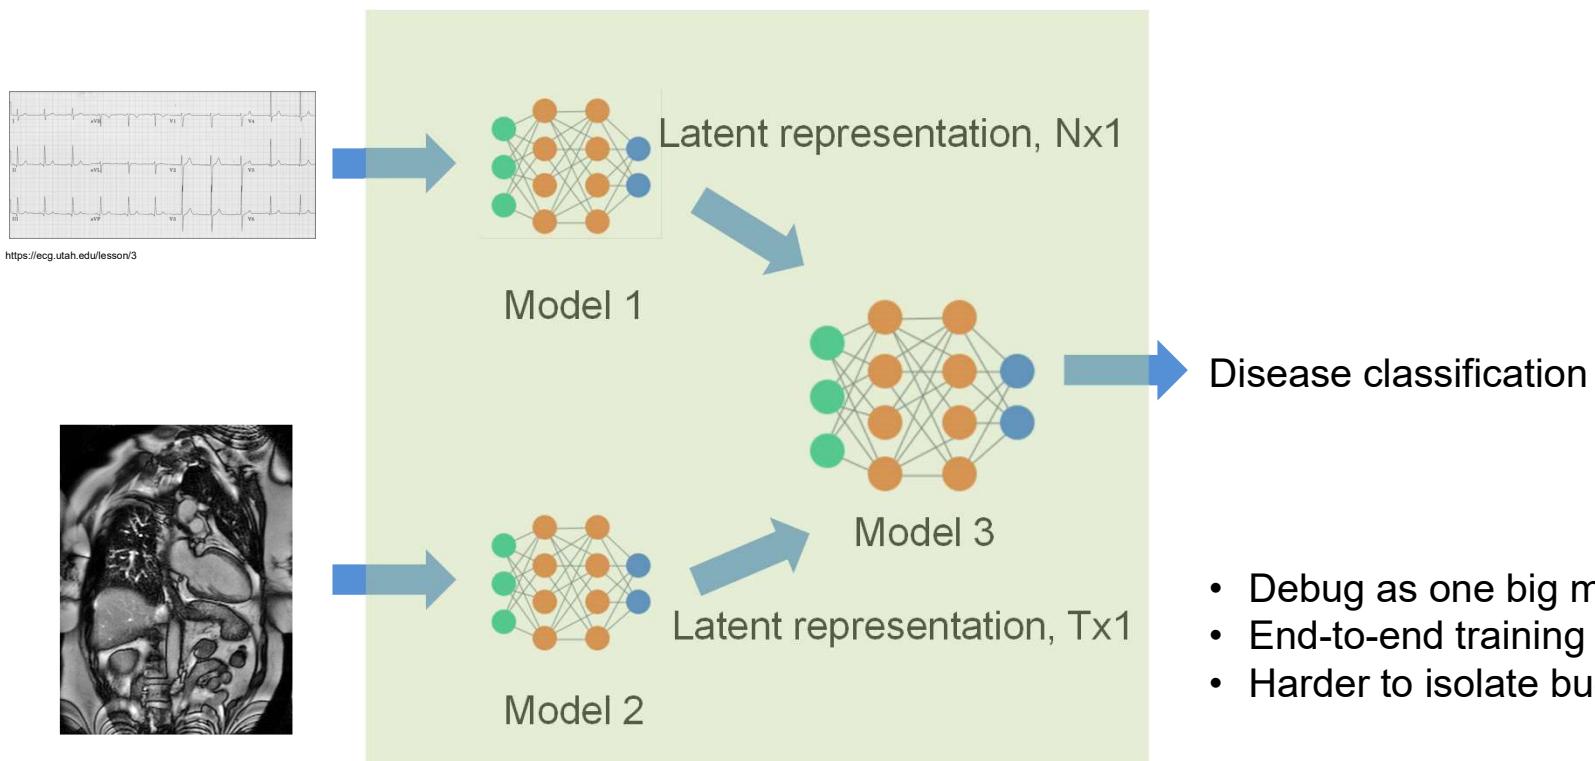
Error analysis of a complex system



Error analysis of a complex system



Error analysis of a complex system



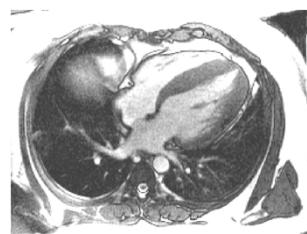
- Debug as one big model
- End-to-end training
- Harder to isolate bugs

Outline

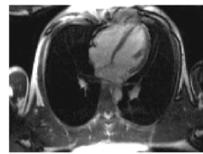
- Set up the training, cont.
- **Data pre-processing**
- Model weight initialization
- Tips for model training

Data pre-processing to reduce variation

Standardizing the input data and reducing the variation can **significantly** help training



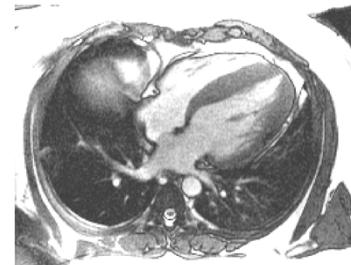
High spatial resolution



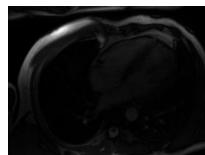
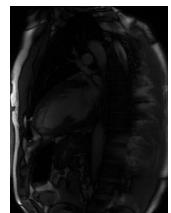
Low spatial resolution



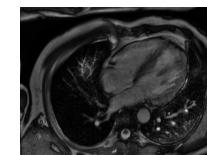
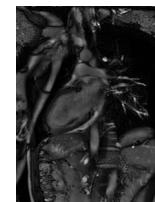
Resample to
fixed resolution



Fixed high resolution of 1mm²



Adjust image
windowing



Data pre-processing to understand

Understand the “edge” samples, exclude the corrupted samples and correct wrong labels



Noisy with artifacts



Signal loss due to devices

MNIST	CIFAR-10	CIFAR-100	Caltech-256	ImageNet	QuickDraw
					
given: 5 corrected: 3	given: cat corrected: frog	given: lobster corrected: crab	given: ewer corrected: teapot	given: white stork corrected: black stork	given: tiger corrected: eye

An average of 3.4% errors across the 10 benchmark ML datasets

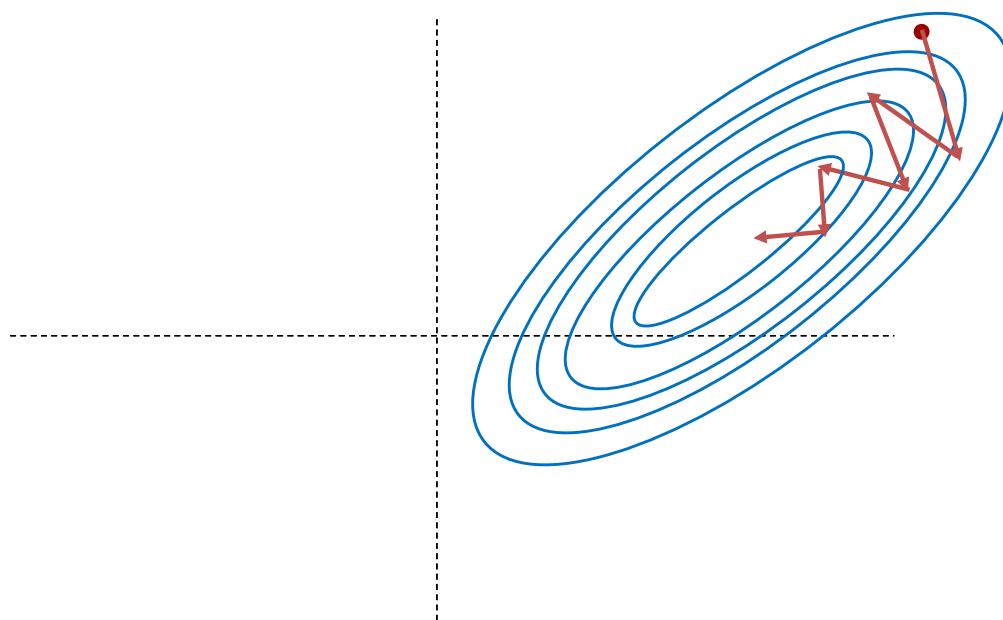
Has impact on evaluating different models

<https://labelerrors.com/>

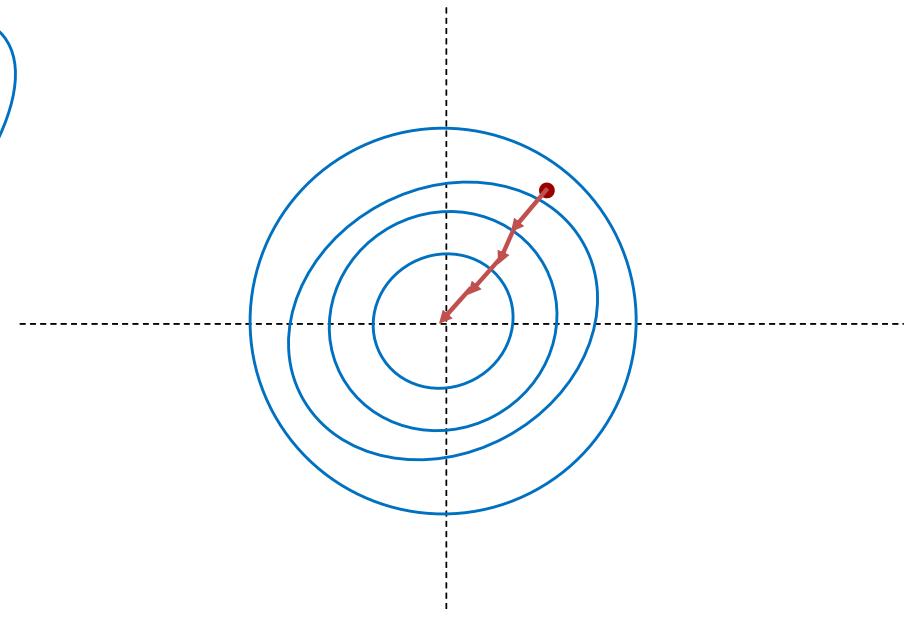
<https://l7.curtisnortcutt.com/confident-learning>

Data pre-processing

Normalize different dimension of input data, help optimization



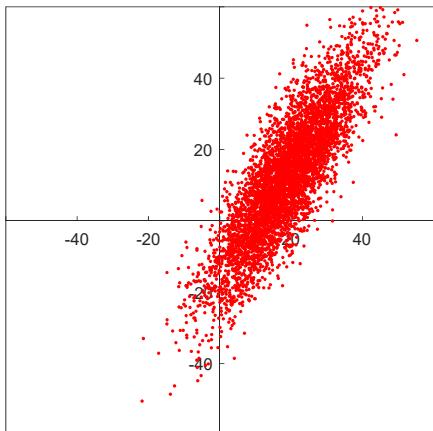
Features with varying magnitude



After scaling and centering

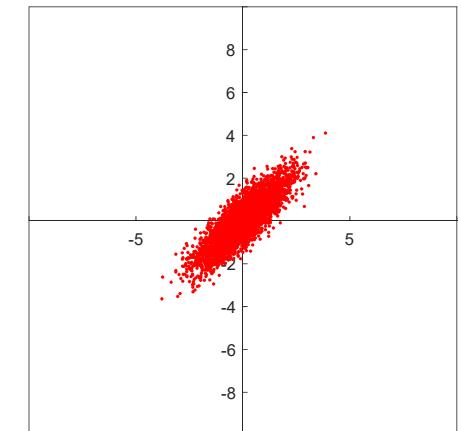
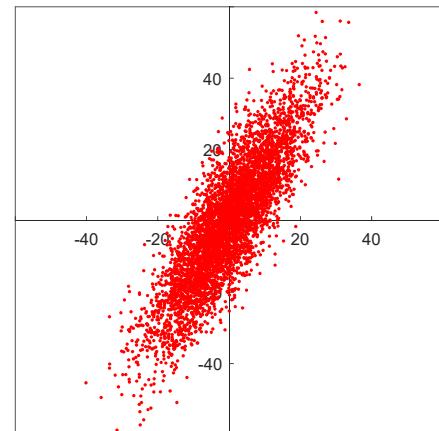
Scale and centering

Normalize different dimension of input data, help optimization



$$X \in \mathbb{R}^{N \times 2}$$

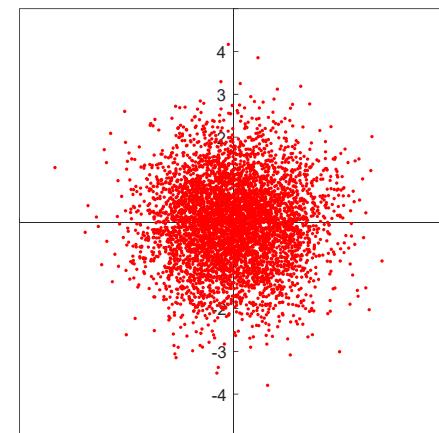
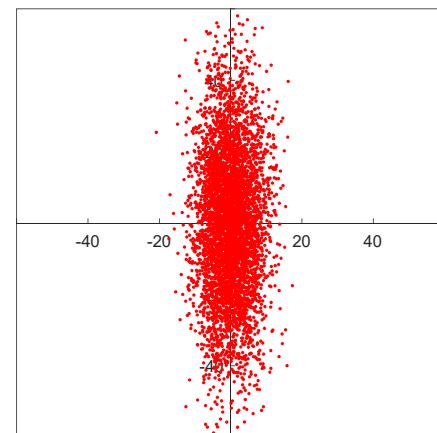
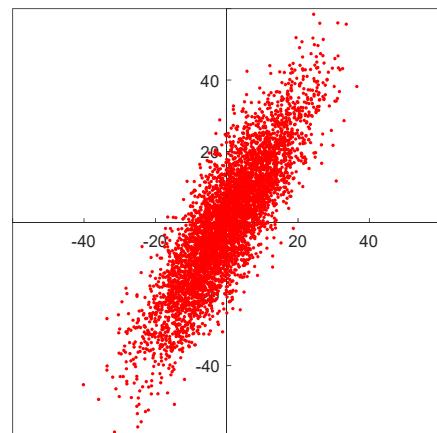
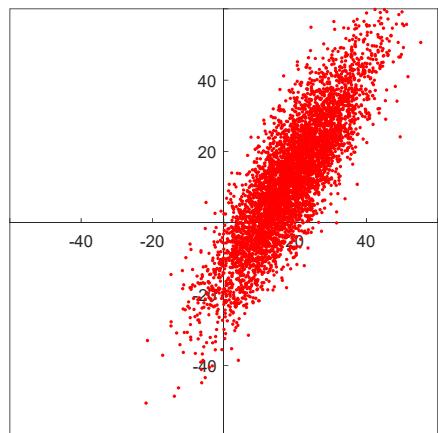
$$X = X - np.mean(X, axis=0)$$



$$X = \frac{X - np.mean(X, axis=0)}{np.std(X, axis=0)}$$

$$\text{Another approach often used: } X = \frac{X - np.mean(X, axis=0)}{np.amax(X, axis=0) - np.amin(X, axis=0)}$$

PCA: Principal Component Analysis



$$X \in \mathbb{R}^{N \times 2}$$

$$X = X - np.mean(X, axis=0)$$

$$\begin{aligned}[E, V] &= np.linalg.eig(np.dot(X.T, X)/N) \\ X &= np.dot(X, V)\end{aligned}$$

$$X = X / np.sqrt(X \cdot \text{diagonal}())$$

PCA: Principal Component Analysis

$X \in \mathbb{R}^{N \times D}$ If the input dimension D is large, dimension reduction can be performed

```
[E, V] = np.linalg.eig(np.dot(X.T, X)/N)
X = np.dot(X, V)
X = X[:, 0:K]
```

Only keep the top K dimension, $X \in \mathbb{R}^{N \times K}$, can filter out some noise

These steps are also called “Whitening” :

the covariance matrix of X after applying the eigenvector and divided by $\text{sqrt}(\text{eigenvalue})$ becomes identity matrix

Tips for data processing

- For image, often only apply centering and scaling

$$X = \frac{X - np.mean(X, axis = 0)}{np.amax(X, axis = 0) - np.amin(X, axis = 0)}$$

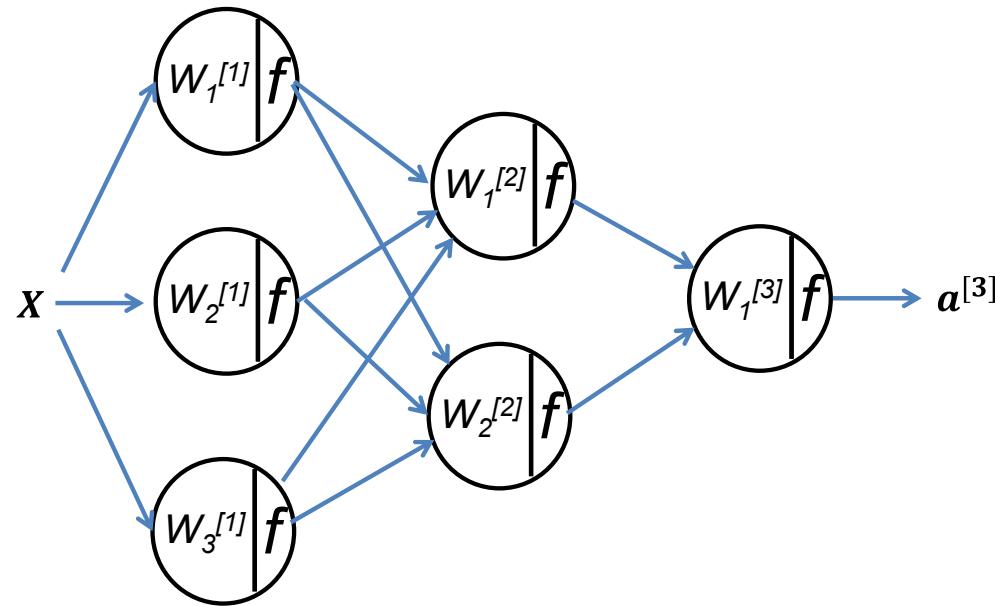
Or simply, $X=X/\text{np.max}(X)$ ← scale image to be [0, 1]

- For time signal or multi-channel signals, PCA can be useful
- Start simple and plot results from every preprocessing steps
- Including the **identical** preprocessing in the regression test
- Make sure the inference data go through the **identical** pre-processing steps
- Computing the pre-processing statistics (e.g. dataset mean, eigen vectors, eigen values ...) on the training set **ONLY** and applying them in inference
- Complicated pre-processing can make model more vulnerable to data distribution shift

Outline

- Set up the training, cont.
- Data pre-processing
- **Model weight initialization**
- Tips for model training

Need a strategy to give weights some initial values



$$\begin{aligned} \mathbf{Z}^{[1]} &= \mathbf{W}^{[1]} \mathbf{X} + \mathbf{b}^{[1]} \\ \mathbf{a}^{[1]} &= f(\mathbf{Z}^{[1]}) \\ \mathbf{Z}^{[2]} &= \mathbf{W}^{[2]} \mathbf{a}^{[1]} + \mathbf{b}^{[2]} \\ \mathbf{a}^{[2]} &= f(\mathbf{Z}^{[2]}) \\ \mathbf{Z}^{[3]} &= \mathbf{W}^{[3]} \mathbf{a}^{[2]} + \mathbf{b}^{[3]} \\ \hat{\mathbf{y}} &= \mathbf{a}^{[3]} = f(\mathbf{Z}^{[3]}) \end{aligned}$$

How do we initialize weights and biases?

Need a strategy to give weights some initial values

How about $W=0$ and $b=0$?

$$\mathbf{z}^{[1]} = \mathbf{W}^{[1]} \mathbf{X} + \mathbf{b}^{[1]}$$

$$\mathbf{a}^{[1]} = f(\mathbf{z}^{[1]})$$

$$\mathbf{z}^{[2]} = \mathbf{W}^{[2]} \mathbf{a}^{[1]} + \mathbf{b}^{[2]}$$

$$\mathbf{a}^{[2]} = f(\mathbf{z}^{[2]})$$

$$\mathbf{z}^{[3]} = \mathbf{W}^{[3]} \mathbf{a}^{[2]} + \mathbf{b}^{[3]}$$

$$\hat{\mathbf{y}} = \mathbf{a}^{[3]} = f(\mathbf{z}^{[3]})$$

All $\mathbf{z}^{[l]}$ will be zero

All $\mathbf{a}^{[l]}$ will be a constant, depending on the non-linear activation function used

When performing backprop, then gradient will be zero ...

So SGD cannot update network parameters ...

Need a strategy to give weights some initial values

How about $W = \sigma \times np.random.rand(input_D, output_D)$, zero-mean, gaussian random values

$$Z^{[1]} = W^{[1]}X + b^{[1]}$$

$W^{[1]}$: 3xN matrix

$$Z_k^{[1]} = \sum_{i=0}^{N-1} W_{k,i}^{[1]} x_i + b^{[1]} \quad \text{The k-th element of score Z}$$

Let's whitening the input X, so it has same variance for all features.

$$\text{Var}(Z_k^{[1]}) = \sum_{i=0}^{N-1} \text{Var}((W_{k,i}^{[1]})\text{Var}(x_i) + \text{Var}(b^{[1]}) = N\sigma^2 \text{Var}(x)$$

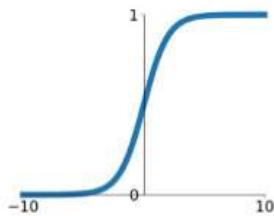
Since all x_i and $W_{k,i}^{[1]}$ are independent upon initialization, with zero mean

$$\text{Var}(wx) = E(w^2)E(x^2) - E(w)^2 E(x)^2 = E(w^2)E(x^2) = \text{Var}(w)\text{Var}(x) \quad \text{If w and x are independent and have zero mean}$$

For deep network, activation function can be saturated

Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



Layer 1
mean=-0.00
std=0.49

Layer 2
mean=0.00
std=0.29

Layer 3
mean=0.00
std=0.18

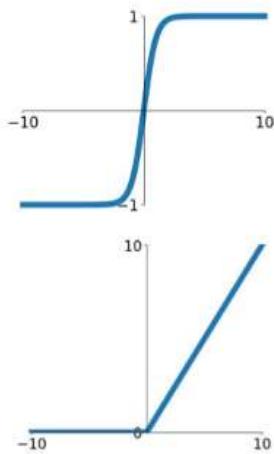
Layer 4
mean=-0.00
std=0.11

Layer 5
mean=-0.00
std=0.07

Layer 6
mean=0.00
std=0.05

tanh

$$\tanh(x)$$



Layer 1
mean=0.00
std=0.87

Layer 2
mean=-0.00
std=0.85

Layer 3
mean=0.00
std=0.85

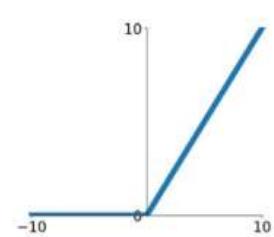
Layer 4
mean=-0.00
std=0.85

Layer 5
mean=0.00
std=0.85

Layer 6
mean=-0.00
std=0.85

ReLU

$$\max(0, x)$$



$\sigma=0.01$

$\sigma=0.05$

tanh activation was used

Figure credit: Stanford CS231N notes

Xavier Initialization

$$Var(\mathbf{Z}_k^{[1]}) = N\sigma^2 Var(x)$$

For a deep network, the non-linear activation can be saturated for native initialization

Make the learning harder

Sensitive to hyperparameter configuration

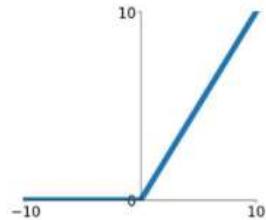
To keep the output variance unchanged:

$$W = np.random.rand(output_D, input_D) / np.sqrt(input_D) \quad \sigma = \frac{1}{\sqrt{N}}$$
$$Var(\mathbf{Z}_k^{[1]}) = \sum_{i=0}^{N-1} Var((\frac{1}{\sqrt{N}} W_{k,i}^{[1]}) Var(x_i) + Var(\mathbf{b}^{[1]}) = Var(x)$$

Xavier Glorot and Yoshua Bengio, "Understanding the difficulty of training deep feedforward neural networks", AISTAT 2010

He initialization: Xavier Initialization for ReLU activation

ReLU
 $\max(0, x)$



$$W = np.random.rand(output_D, input_D) \times \frac{2}{np.sqrt(input_D)}$$

$$\sigma = \frac{2}{\sqrt{N}}$$

To compensate the ReLU function for the negative inputs

$$Var(wx) = E(w^2)E(x^2) - E(w)^2 E(x)^2 = E(w^2)E(x^2) = Var(w)E(x^2) = \frac{1}{2}Var(w)Var(x)$$

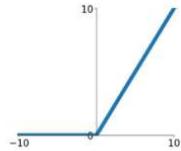
x is the response after RELU, which has positive mean

- Use Xavier or He initialization
- Bias is initialized to zero; randomness in weights often is sufficient to drive the training
- With Batch Normalization, network training is much less sensitive to parameter initialization

He et al, "Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification", ICCV 2015

He initialization: Xavier Initialization for ReLU activation

ReLU
 $\max(0, x)$



y is the score before ReLU:

$$x = \max(0, y)$$

y has zero-mean and symmetric distribution,
if the network is in good condition before this layer

$$\text{Var}(wx) = E(w^2)E(x^2) - E(w)^2 E(x)^2 = E(w^2)E(x^2) = \text{Var}(w)\text{Var}(x) = \frac{1}{2}\text{Var}(w)\text{Var}(x)$$

x is the activation after ReLU, which has positive mean

$$E[x^2] = \int_{-\infty}^{+\infty} \max(0, y)^2 p(y) dy$$

where the part $y < 0$ does not contribute to the Integral

$$= \int_0^{+\infty} y^2 p(y) dy$$

which we can write as half the integral over the entire real domain (y^2 is symmetric around 0 and $p(y)$ is assumed to be symmetric around 0):

$$= \frac{1}{2} \int_{-\infty}^{+\infty} y^2 p(y) dy$$

now subtracting zero in the square we get:

$$= \frac{1}{2} \int_{-\infty}^{+\infty} (y - E[y])^2 p(y) dy$$

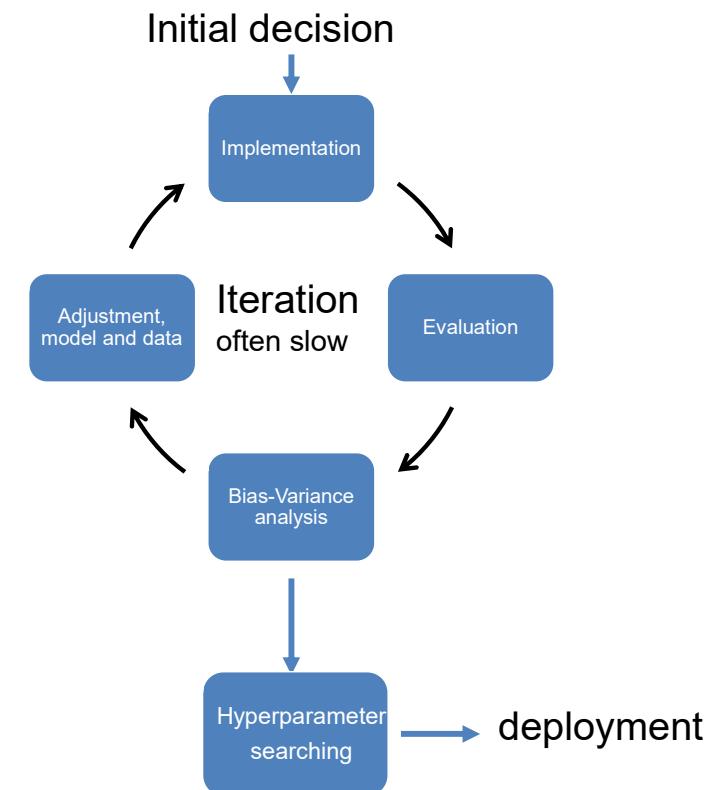
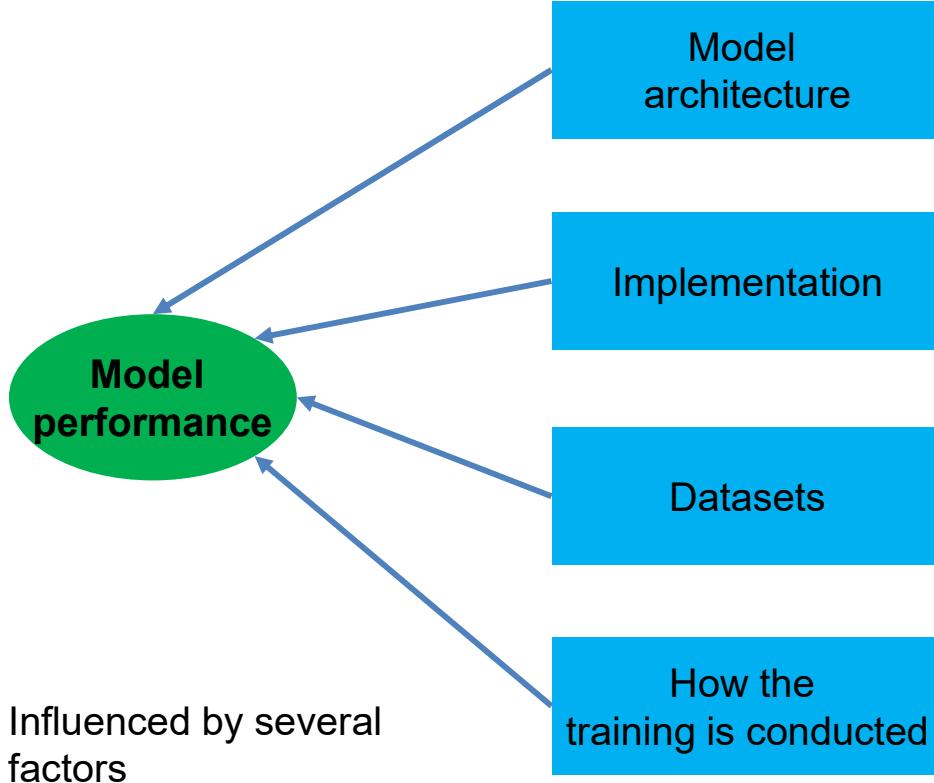
which is

$$= \frac{1}{2} E[(y - E[y])^2] = \frac{1}{2} \text{Var}[y]$$

Outline

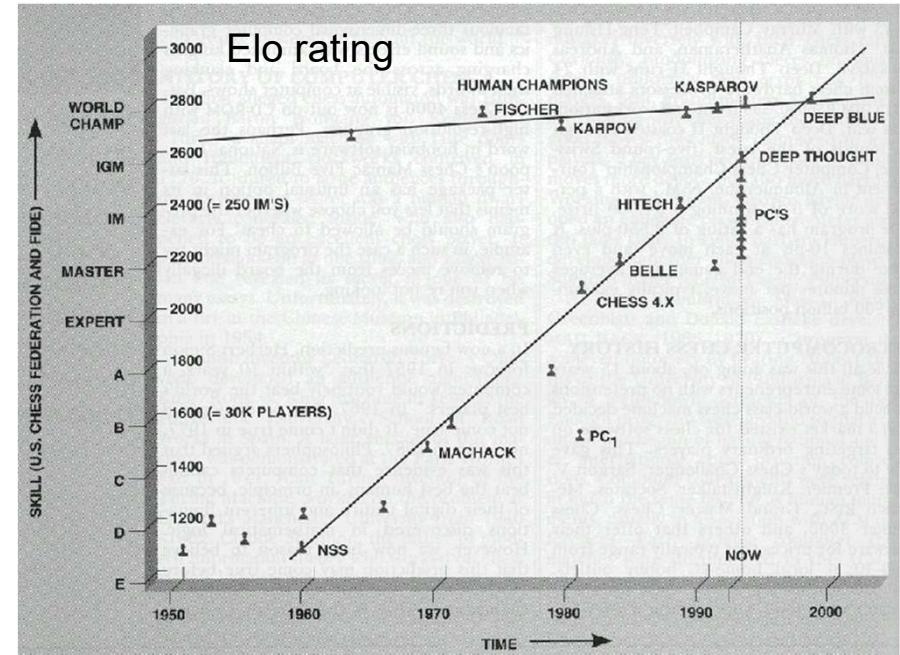
- Set up the training, cont.
- Data pre-processing
- Model weight initialization
- **Tips for model training**

Debug deep neural network model can be difficult



Before you started ...

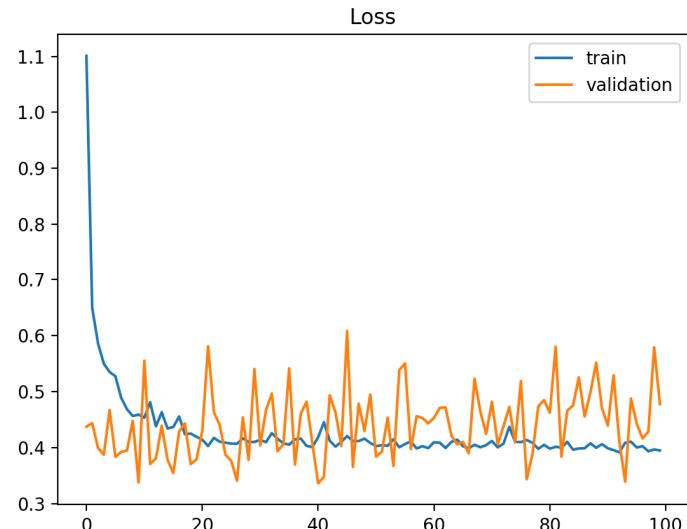
- Have an estimation of Bayes accuracy
 - Human level performance as a surrogate
 - Expert performance – experienced operators vs. less experienced
 - Voting strategy
 - Baseline and results from publications
 - Whether it has related competition



<https://www.drdobbs.com/parallel/computer-chess-the-drosophila-of-ai/184405171>

Before you started ...

- Have the Dev and Test set established



- Dev and test sets should be large enough to give a reliable estimation of model performance
- If the total amount of data is small, get more data and use cross-validation

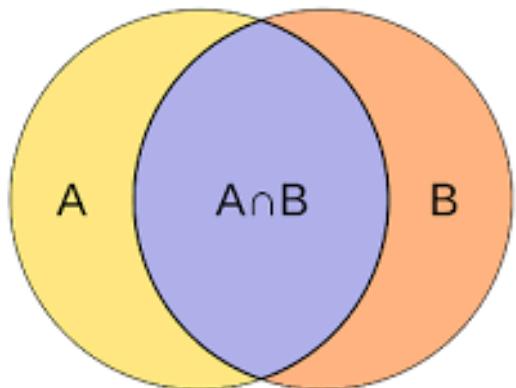
<https://machinelearningmastery.com/learning-curves-for-diagnosing-machine-learning-model-performance/>

Before you started ...

- Have the evaluation metric

Accuracy, Dice ratio, L2 distance, business specific metrics?

Pick the key metric



$$\text{Dice coefficient}(A, B) = \frac{2 \times |A \cap B|}{|A| + |B|}$$

		True Class	
		Positive	Negative
Predicted Class	Positive	TP	FP
	Negative	FN	TN

$$\text{Accuracy} = (TP+TN)/(TP+TN+FP+FN)$$

$$\text{Precision} = TP/(TP+FP)$$

$$\text{Sensitivity} = TP/(TP+FN)$$

$$\text{Specificity} = TN/(TN+FP)$$

$$\text{F1 score} = 2 \times \text{Pre} \times \text{Sen} / (\text{Precision} + \text{Sensitivity})$$

<https://towardsdatascience.com/confusion-matrix-for-your-multi-class-machine-learning-model-ff9aa3bf7826>

Before you started ...

- **Have the expectation**

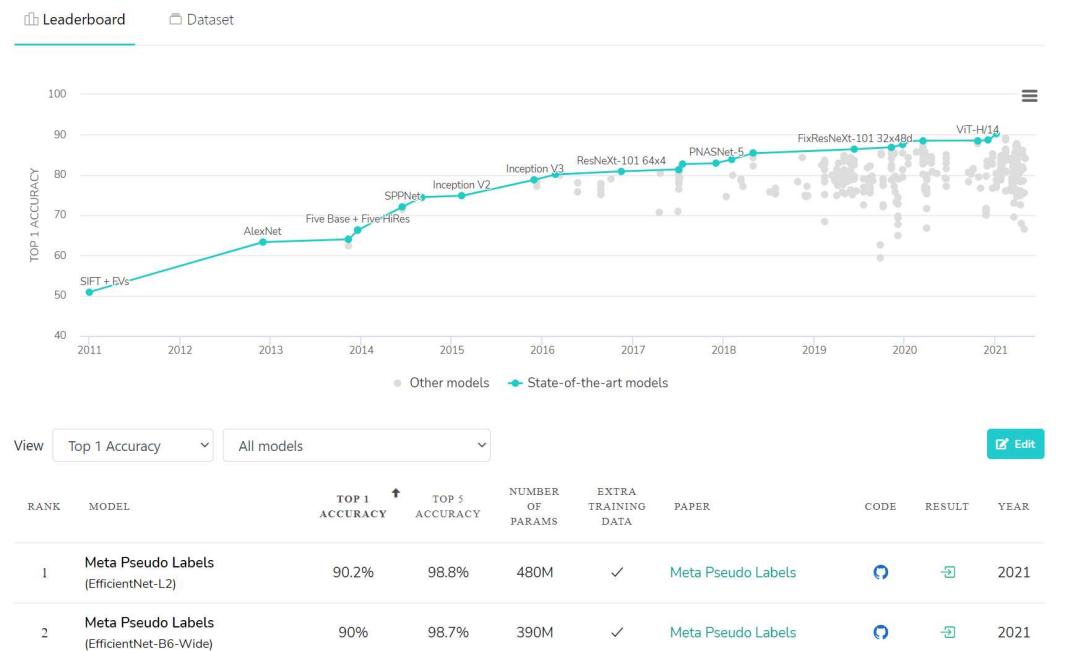
Know when to stop
Be the state-of-the-art?
Meet the business needs?

Consider the resource limitation

Data and computing

90% and 90.2% may be the same for your deployment

Image Classification on ImageNet



<https://paperswithcode.com/sota/image-classification-on-imagenet>

Model architecture

- Start from standard choices and make changes if needed
- Look for available resources
- Look for related publications
- Read the code if possible and do not take it blindly

Tasks	First try
Image classification	Standard ResNet, e.g. implementation in Pytorch repo
Segmentation	Unet architecture
Language, NLP	Pre-trained GPT family models
Time series signal	Bidirectional RNN with LSTM, transformer
Tabular data	Dense MLP, DropOut, BatchNorm

Implementation

- Pay attention to output layers

With the same base network, different functionalities can be achieved by changing the output layers:

Pooling and dense linear for classification

CONV layer for segmentation

Dense linear for regression

- Pay attention to loss function – need logits or need probabilities?

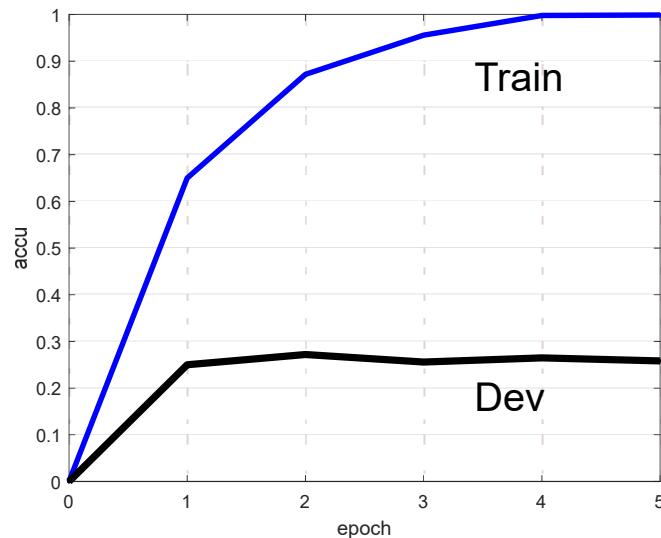
different implementation may pack extra computation into loss function

```
>>> m = nn.Sigmoid()
>>> loss = nn.BCELoss()
>>> input = torch.randn(3, requires_grad=True)
>>> target = torch.empty(3).random_(2)
>>> output = loss(m(input), target)
```

```
>>> m = nn.LogSoftmax(dim=1)
>>> loss = nn.NLLLoss()
>>> # input is of size N x C = 3 x 5
>>> input = torch.randn(3, 5, requires_grad=True)
>>> # each element in target has to have 0 <= value < C
>>> target = torch.tensor([1, 0, 4])
>>> output = loss(m(input), target)
```

Overfit a small dataset, e.g. one mini-batch

- For classification, accuracy should quickly reach 100%
- For segmentation, Dice ratio should quickly increase to ~1.0



Very useful to find implementation bugs:

Loss goes up

learning rate too high, error in output layer for class order...

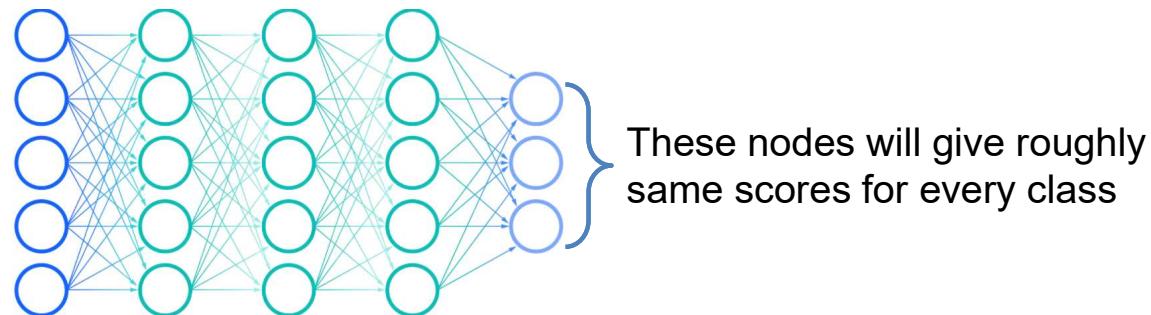
Loss is NaN

Overflow or underflow or divided-by-zero

Accuracy plateaus

learning rate too low, bad initialization, regularization too high, data/label mismatch ...

Check the initial loss function, if making sense



- Random initialized
- All output classes are treated equally
- For the binary classification, initial loss will be close to $-\log(0.5)$
- For the multi-class classification, initial loss will be close to $-\log(1/C)$, C is the number of total class

First turn off regularization/drop out/Batch Norm

$$L = \frac{1}{B} \sum_{i=0}^{B-1} L^{(i)} + \frac{\lambda}{2} \|\mathbf{W}\|_2^2$$

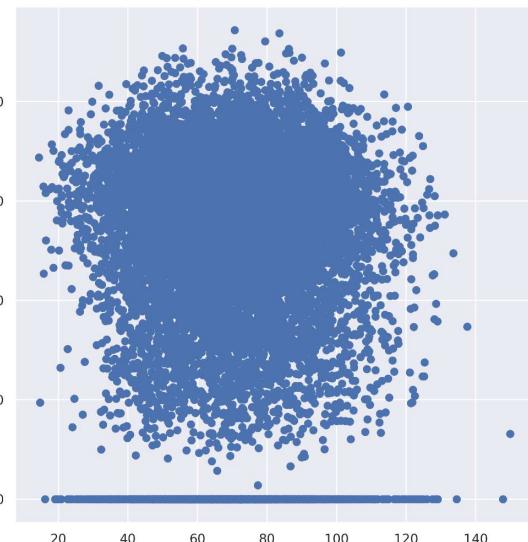
- Set $\lambda = 0$ to begin with
- Make sure regularization loss not dominate total loss
- Turn off operations to introduce randomness to training, e.g. drop out
Add controls to allow turn them on/off easier
- Consider to turn off batch norm at first trials, since BN layer can hide problem such as NaN in computation, explosion/vanishing gradients, and make debug harder
- With BatchNorm, make sure Batch size large enough (>16)

Check **every** steps, starting from pre-processing

Errors can happen in pre-processing, such as
forget to change label after processing images (e.g. image scaling for segmentation)

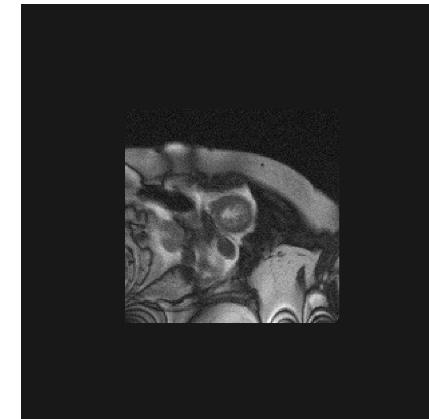
Errors can happen in datasets:
forget to shuffle the samples, some samples are corrupted ...

Errors can happen when loading data:
image/label mismatch

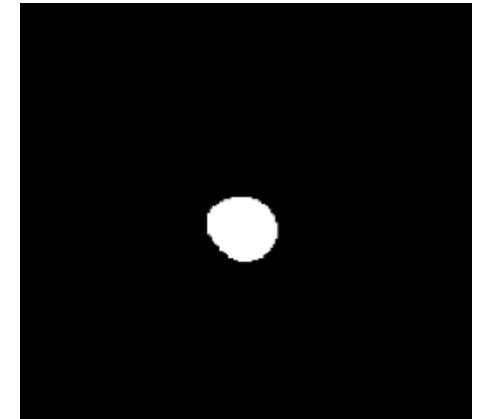


L2 norm of all
samples

Some images
are blanket!

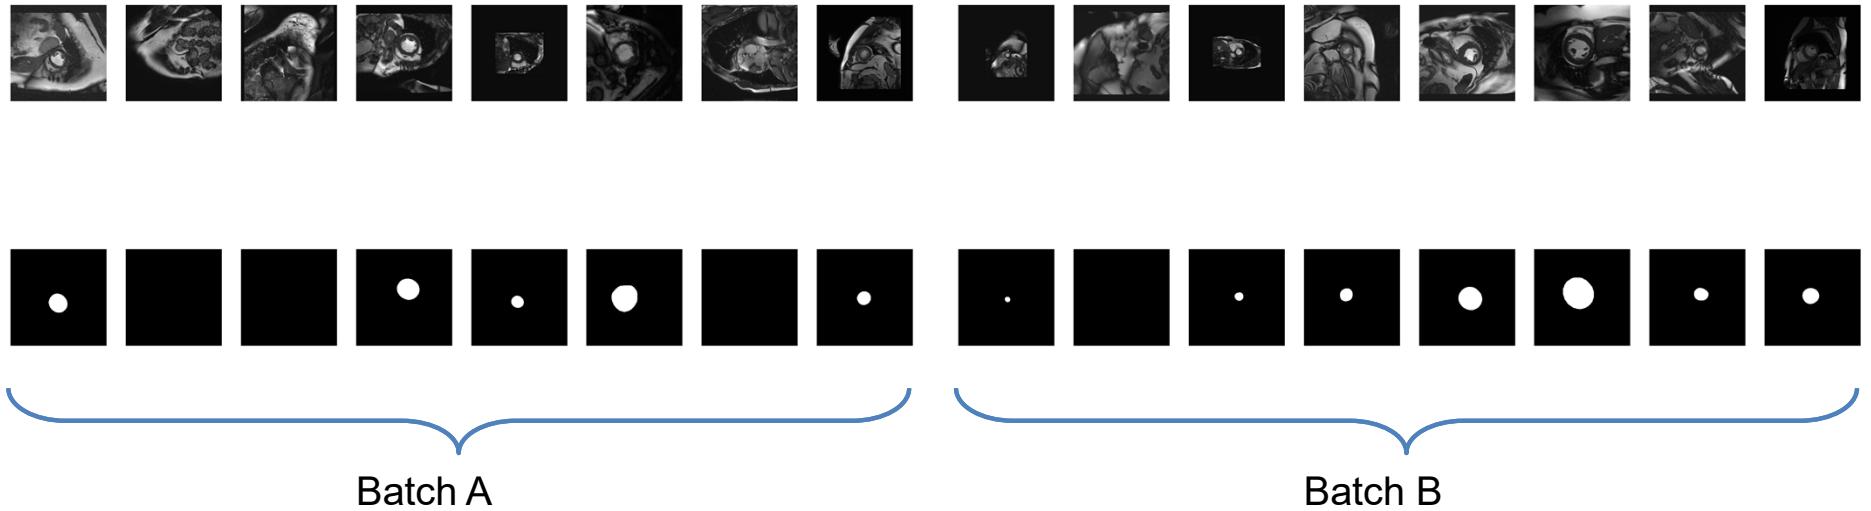


Scale the image and forget to scale mask



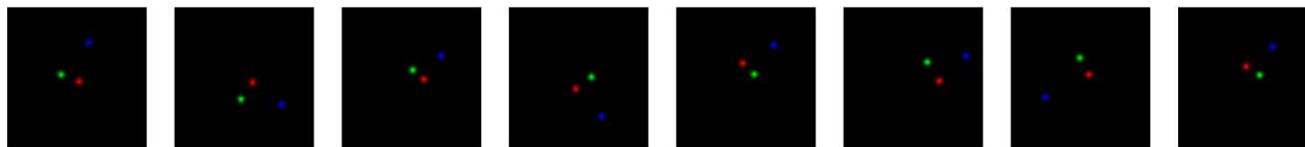
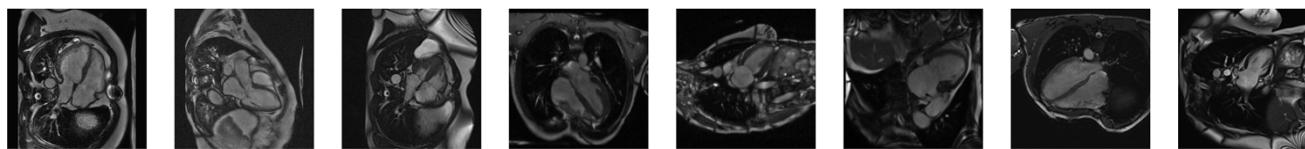
Check immediate inputs to your model training code

Print/plot/save a few mini-batch and check them before training

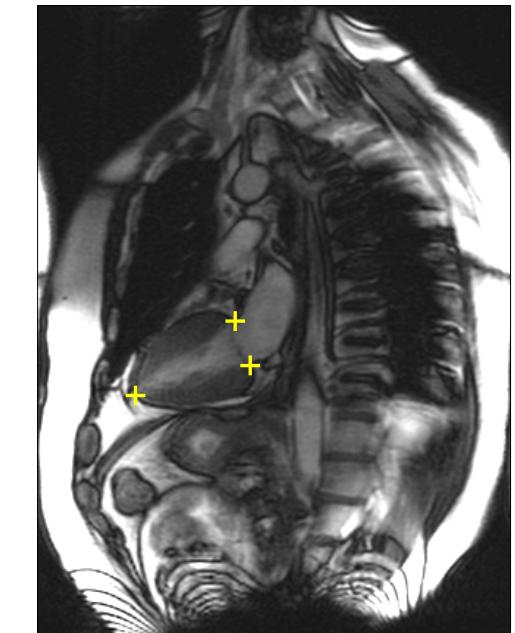


Check immediate outputs of your model before any post-processing

```
scores = model(images)
```



Model outputs are spatial probability maps



Goal : detect key landmarks

Postprocessing: extract landmark location

Add Debug mode and verbose mode, add working_dir

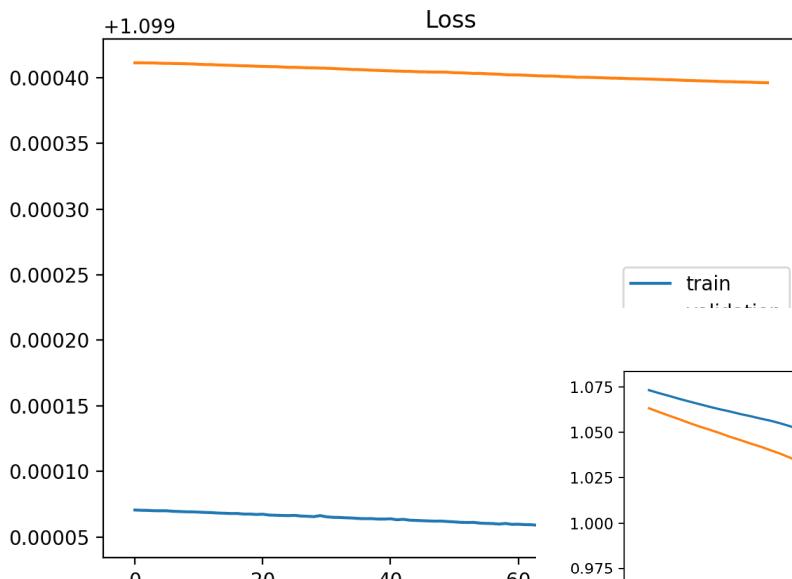
```
if(self.working_dir is not None):
    self.save_batch(self.working_dir, images, batch_no)

scores = model(images)

if(self.working_dir is not None):
    self.save_output(self.working_dir, scores, batch_no)
```

- Easily enable or disable
- Fast, not slow down training if disabled
- Save intermediate results with its index/name/time stamp/data role ...
- Keep these code lines in your repo

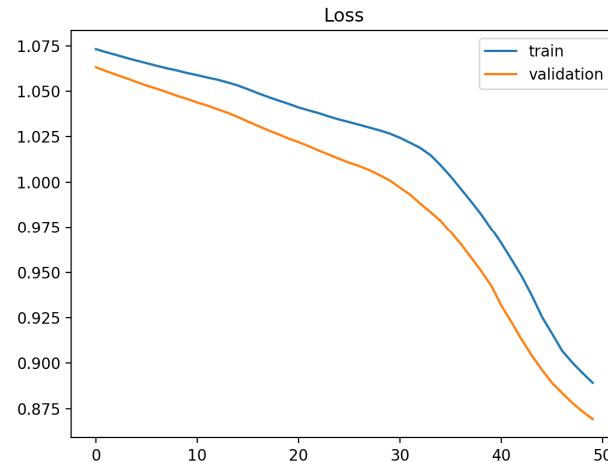
Monitor the learning curves



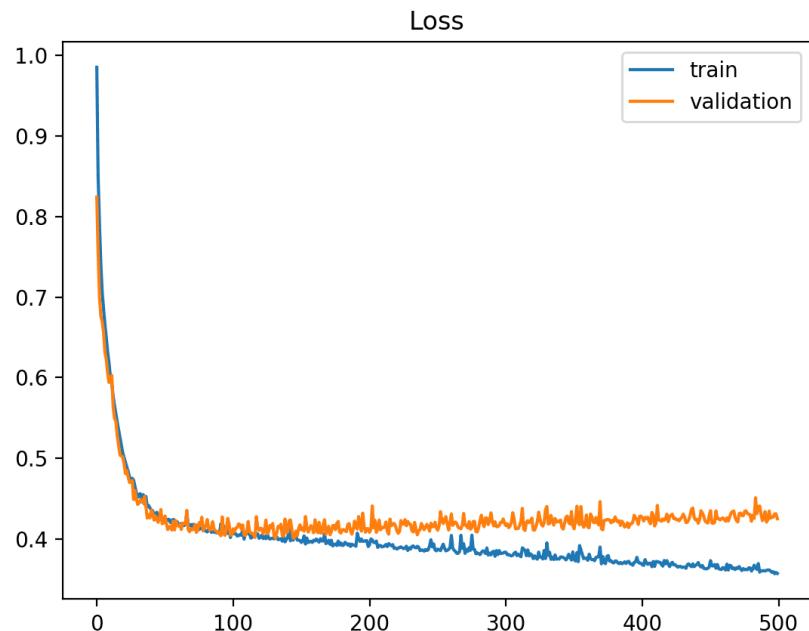
Loss is not decreasing quickly

learning rate too low
model underfitting

try to fix the problem by fitting to a tiny dataset



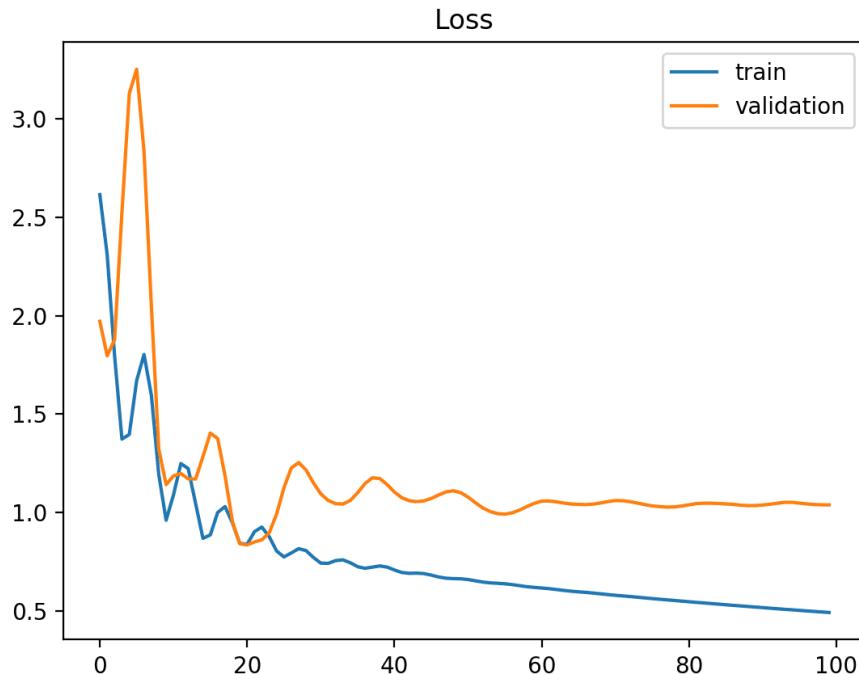
Monitor the learning curves



Loss goes up for dev set and goes down for training set

- model overfitting
- add regularization
- reduce model capacity
- add more training data

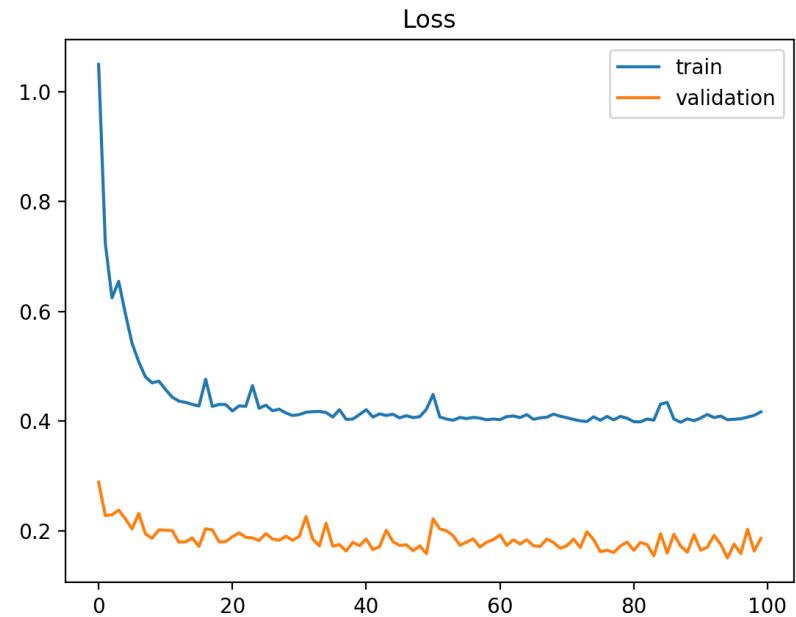
Monitor the learning curves



Loss shows strong variance in dev set

model overfitting
training set too small
Add more sample to training set

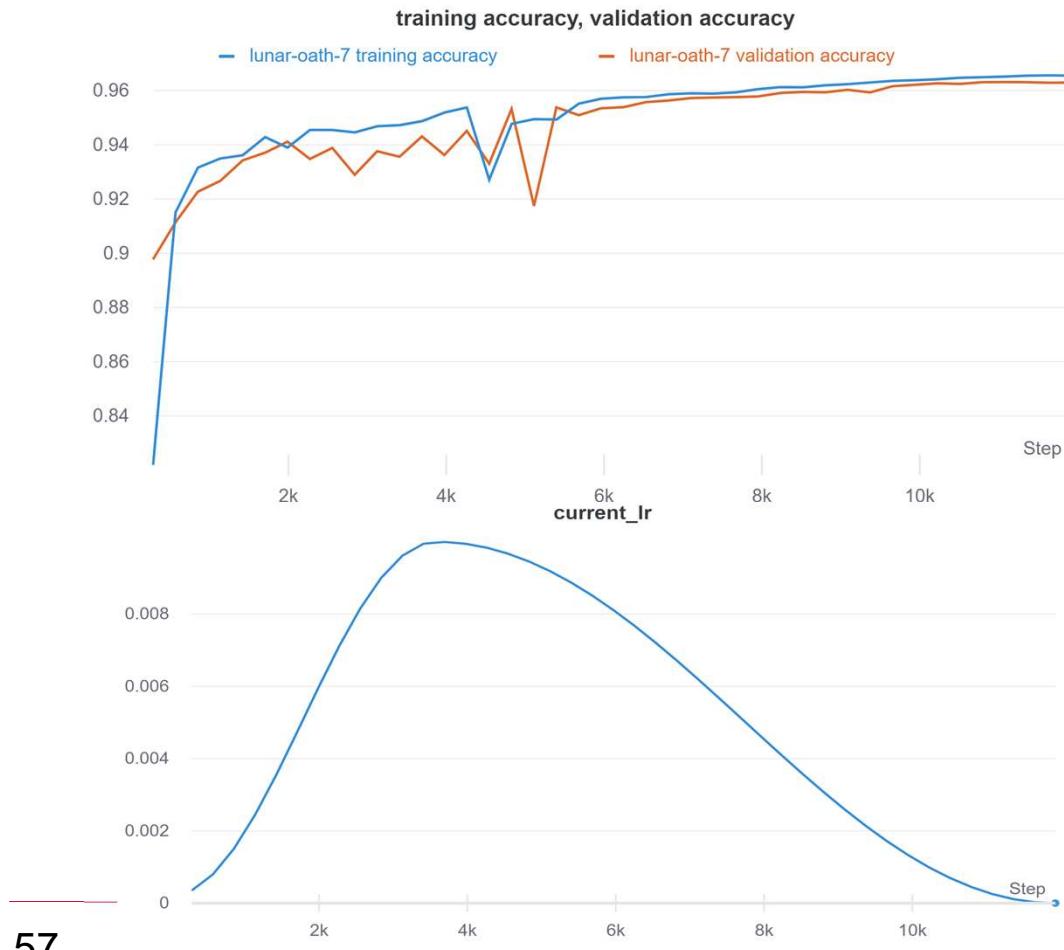
Monitor the learning curves



Loss is lower for dev set than train set

dev set is easier than train set
too much data augmentation
dev/tra set distribution mismatch
exam dev set

Monitor the learning curves



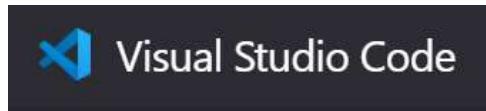
Good training

Accuracy is high, but still increasing
Consider to train longer

Note with one-cycle scheduler, at high LR, the accuracy can dip

Use Debugger

Tensor shape mismatch
Loss function exceptions
Nan or inf error
All other exceptions ...



or, pdb or ipdb for CLI
experiences of debugging

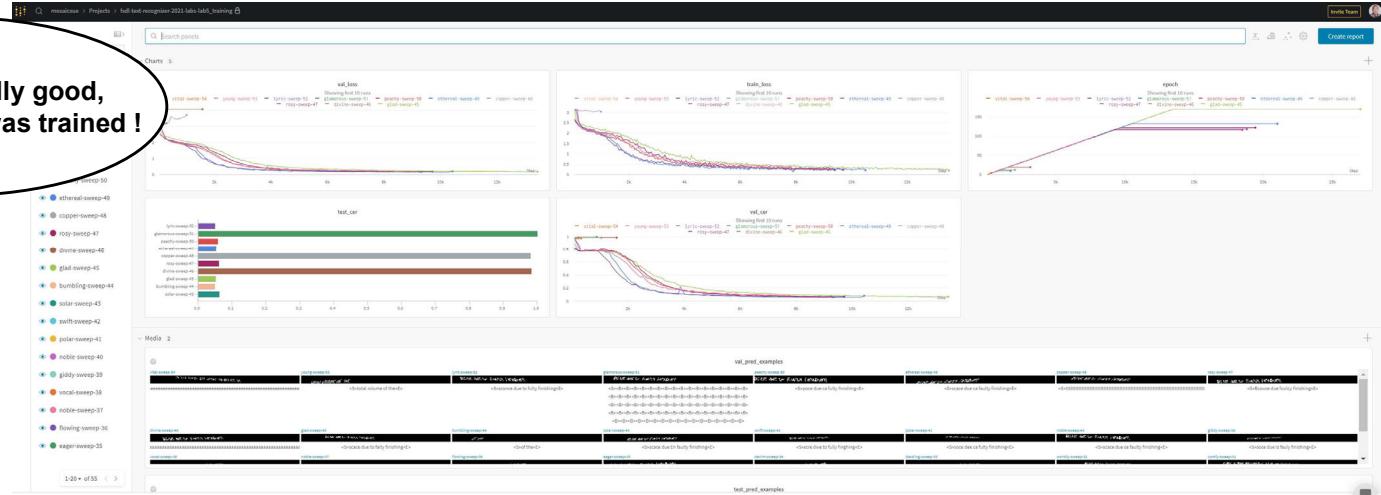
The screenshot shows a debugger interface with several panes:

- File Explorer:** Shows project structure with files like `run_training.py`, `networks_savn.py`, and `perf_aif_lv_detection.py`.
- Code Editor:** Displays the `run_training.py` file with line numbers and some highlighted code.
- Variables:** Shows variables `tra`, `val`, `train_idxs`, `val_idxs`, `test_idxs`, `dataset`, `loader_for_train`, `loader_for_val`, `iter_train`, `images`, `masks`, `labels`, `dms`, and `names`.
- Output:** Shows command-line output including file paths and dimensions for images and masks.
- Terminal:** Shows the command `python run_training.py` being run.

On the right side, there are additional panes for `PROBLEMS`, `OUTPUT`, and `TERMINAL`. A sidebar on the far right shows a list of files and a search bar.

• set breakpoints
• check variable values
• interactive debugging

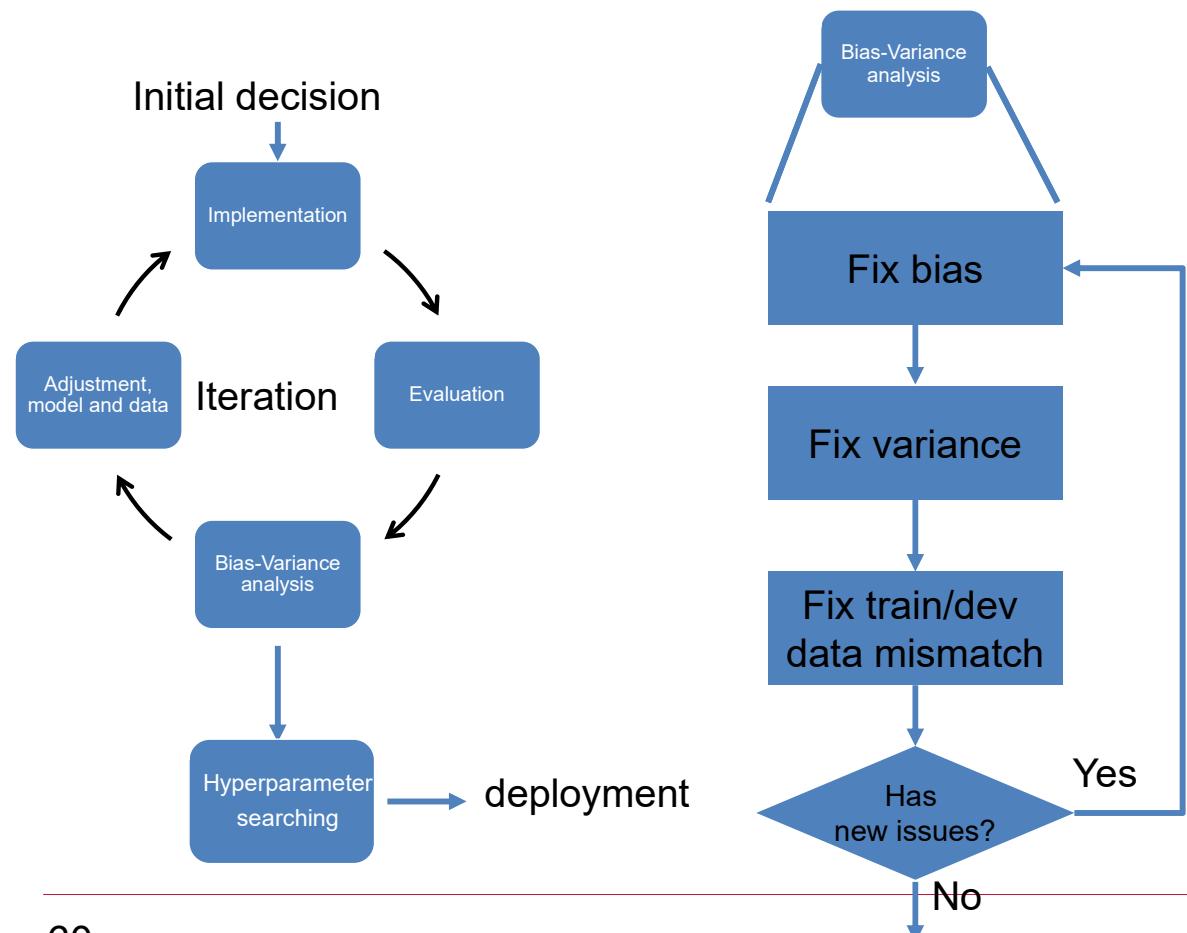
Use Experimental management tools



- Record your training and parameters
- Integrated with python training code
- Support hyperparameter searching
- Easy to compare different experiments

Many iterations may be needed to reach required performance

Do not use hyperparameter searching as the first remedy



- Start hyperparameter searching after fixing obvious bias/variance problems
- Often large performance gain is from error analysis and collecting specific new data for problems found
- Fix bias/variance problems can change model and which hyperparameters to tune

